

GPUs, Wombs, and Money Printers

Scenarios and Responses
for Human Society in the Age of AI

Tony Seah

v1.4.0 August 2026

Contents

Preface: How to Read This Book	1
Volume I · The Extrapolation	3
I. Introduction: From Cat-vs-Dog to Passing the Bar—A Decade Apart	3
II. The Collapse of Intelligence Premium: When Human Brains Aren’t Scarce	4
2.1 The Human-Intelligence Displacement Spiral	4
2.2 Friction to Zero & the Collapse of Middlemen	5
2.3 Ghost GDP & the Composition Fallacy	6
2.4 Autonomous Driving: Hand-Driving Will Look Like Horseback	9
2.5 Humanoid Robots: When AI Grows Limbs	11
2.6 The 35-Year Crisis Flips: AI Favors Seniors	13
2.7 The Degree Devaluation Cascade	13
2.8 Three Misaligned Timelines: Unemployment, Fiscal Pressure, Longevity	16
III. Drag Forces: What Could Slow AI	17
3.1 Physical: Power, Chips, Atoms	17
3.2 Trust: Hallucinations, Last Mile, ROI Hangover	17
3.3 Institutional: Laws, Labor, Legal Vacuum	18
3.4 Fast vs Slow: Differential Diffusion Rates	18
3.5 Calm Before the Storm: Why 2026 Jobs Look Fine	19
3.5.1 The Lights-Out Office: AI at the Centre, Humans at the Edge	21
3.5.2 Three Gaps: Asking the Question From the Other Side	24
3.6 Dashboard: Six Indicators, Three Thresholds	26
3.7 Brake Conditions and Scorecard	26
IV. Geographic Differentiation	27
4.1 U.S.: Eye of the Storm	28
4.2 China: Different Script, Similar End	28
4.3 India: Digital Colonization Cascade	28

4.4 Europe: Regulator’s Dilemma	30
4.5 Developing World: Advantage Erased	30
V. Population Decline: Natural, Not a “Crisis”	30
5.1 Fertility Push Fails	31
5.2 Economics of Marriage/Childbearing Collapse	31
5.2.1 The Unbundling of Marriage: Four Things Were Tied Together, Then Loosened One by One	32
5.2.2 Equality Is a Companion of Development, and So Is the Tension	33
5.2.3 The Chinese Version: Urbanization Made the Structural Mismatch Spatial	34
5.3 AI–Fertility Spiral	36
5.4 Robot Caregiving	37
5.5 Implications for Policymakers: Production Is Not the Problem—Distribution Is	37
VI. Financial Chain Reaction	38
6.1 Mortgages: Good Loans in a Bad World	38
6.2 From Income Repricing to Municipal Finance: Three Different Paths	39
6.3 Dual-Speed Cities	41
6.4 Private Credit Landmine	45
6.5 Policy Toolbox Failure	45
6.6 Crypto, Speculation, “Unemployed Casino”	45
VII. Tragedy of the Commons	46
7.1 Firm-Level Race	46
7.2 Nation-Level Arms Race	46
7.3 Diversionary Conflict Risk	47
Why Global Coordination Is Nearly Impossible	48
7.4 AI-Enabled Weapons of Mass Destruction	49
7.5 AI Deception, Automation Complacency, and the Control Paradox	50
7.6 Mechanism of Becoming a “Transitional Species”	50
VIII. AI Reshapes a Human Life: From Cradle to Grave	50
8.0 Before We Begin (I): LLM and Agent Are Two Different Things	51
Going One Level Higher: Agent Is a “Soul” That Can Be Installed in Any Body	52

8.0 Before We Begin (II): Five Defenses That Are All Collapsing at Once	52
Defense 1: “AI hallucinates”	53
Defense 2: “It hasn’t threatened my job yet”	53
Defense 3: “AI has no feelings, no empathy”	53
Defense 4: “AI has no creativity, it just copies”	53
Defense 5: “AI has no values or sense of purpose” (the deepest one)	54
Summary of the five defenses	54
8.0 Before We Begin (III): Evidence Already Running—OpenClaw	54
8.1 The Chaperoned-by-AI Generation	56
8.2 The Meaning Crisis: When Work No Longer Defines “Who I Am”	56
8.3 AI Companions and the End of Intimacy	57
Health Divergence: Rich Live to 150, Poor Stay at 75	58
8.5 Digital Afterlife: When Death Is No Longer the End	58
8.6 The Death of Consumption and the Collapse of Cultural Consensus	59
8.7 The Creativity Paradox: Process Is the Value	59
IX. Far Future (2035–2125)	60
9.1 Species Split: Augmented vs Natural	60
9.2 Longevity Paradox: Who Pays the 55-Year Retirement?	62
9.3 Immortality Second-Order Effects	62
9.4 Space Commercialization: One Person’s Empire	62
9.5 AI as an Interstellar Species	64
Not Just Humans: The Twilight of Most Carbon-Based Life	66
Energy efficiency: carbon life’s fatal weakness	66
Other advantages matched one by one	66
Who gets hit first? Not all carbon life equally	67
The womb’s endgame: two paths converge	67
Garden or wilderness?	67
Post-Scarcity: When Material Abundance Actually Arrives	68
X. One Person’s Life, Eight Billion Fates	68
Xiao Ming’s Life: Born in 2026	69
Eight Billion Fates: Seven Phase Transitions	70

The Canary, Reflexivity, and Your Choice	71
Volume II · The Response	74
XI. The Optimist’s Algorithm: Turning Worst Cases into a Task List	74
11.1 Three Premises Taken for Granted—All Three Are Questionable	74
11.2 Policy Levers: Three Lines, Three Horizons	76
11.3 The Test: When to Change Gear	76
11.4 Optimism Is Not Naivety	76
11.5 For Managers and Policymakers: The One-Line Version	77
XII. The Transition: Why the Road Between Here and There Is the Most Dangerous Part	77
12.1 The Endgame May Be Good	77
12.2 But the Transition Will Kill People	78
12.3 Transitions in History	78
XIII. Mental Resilience: The Foundation for Everything Else	79
Why Psychology Comes First	79
13.1 The Psychological Structure of the AI Shock: Crushed From the Top Down	79
13.2 A Lifetime of Accumulation, Wiped Out: How Deep That Pain Goes	80
13.3 The Special Risk Facing Men	81
13.4 Create Your Own ”Luck” – Evidence From a Decade of Scientific Research	81
13.5 The Anatomy of Anxiety: A Simple but Useful Formula	83
Reducing Uncertainty	83
Reducing Importance – The More Subversive Move	84
13.6 Concrete Mental Survival Strategies	85
13.7 If You Are Already in Crisis	87
13.8 For the People Around You: Someone You Know May Be Struggling	87
XIV. Face Reality: HALO — What Is Losing Value, What Is Gaining Value	87
14.1 HALO: Wall Street’s Panic Signal	88
14.2 The Death of SaaS: When Code Becomes a Consumable	88
14.3 What This Means for You Personally	89
14.4 Money: What Counts as a Truly Safe Asset	89

14.4.1 Asset Tiers Under the HALO Framework	89
14.4.2 Cash Is Optionality, Not Waste	90
14.4.3 An Honest Word About "Passive Income"	90
14.5 Skills: What Counts as a "HALO-Type Skill"	91
14.5.1 HALO-Type Skills vs. Light Skills	91
14.5.2 "Learning AI" Is Not the Answer	91
14.5.3 Skill Strategies Have Expiration Dates Too	91
XV. Time and Circumstance: Different Windows, Different Groups	92
15.1 Phase 1: Now to 2030 – "AI Augments Humans"	92
15.2 Phase 2: 2030-2040 – "AI Replaces Humans"	93
15.3 Phase 3: After 2040 – "The Post-Human Era"	93
15.4 The Real Situation Facing Different Groups	93
15.4.1 Those With Capital and Skills (Top 10%)	93
15.4.2 Average White-Collar Workers (Middle 60%)	93
15.4.3 Manual Laborers and Low-Income Groups (Bottom 30%)	94
XVI. Children: Let Go of the Obsession – This May Be the Greatest Liberation	94
16.1 A Realization That Could Halve Parental Anxiety	94
16.2 What Your Child Should Have	94
16.3 What to Let Go Of	95
16.4 A Different Perspective	95
XVII. The Good Things AI Brings: This Chapter May Be Longer Than You Expect	96
17.1 Healthcare: For the First Time, the Poor Can Afford a Good Doctor	96
17.2 Accessibility: The World Is Opening Up for People With Disabilities and the Elderly	97
17.3 Education: The Best Tutor in History, for Free	98
17.4 One Person, One Army: The Barrier to Entrepreneurship and Creation Drops to Zero	98
17.5 Everyday Improvements: Unglamorous but Real Benefits	99
17.6 Shifting Perspective: From Threat to Tool	100
17.7 "Survive the Storm and Better Days Await" – The Power of This Narrative	100
17.8 Your Own Health: Living to 2035 May Be the Best Investment You Ever Make	101
Conclusion · Crossing the Transition	101

Preface: How to Read This Book

This book has two volumes. They do not answer the same question.

Volume I asks: if current trends are not interrupted, what happens? It makes no single-point prophecies; it runs a conditional extrapolation — taking observable data and trend lines already in motion and following each one to its end: if these trends are not interrupted by institutions or by technology itself, where do they take us. Its tone is cold, because an extrapolation owes no one comfort. Volume I gives you a map: where the cliffs are, how the road collapses, how much time is left.

Volume II asks: once you see it, what do you actually have? The same data, turned to a different use — from “what will happen” to “so what should be done.” It works on two levels: what institutions can do (Chapter XI), and what an ordinary person can do while institutions lag (Chapters XII through XVII).

The two volumes speak differently, and that is deliberate.

Volume I rarely addresses you directly, because an extrapolation has to set individual circumstance aside — its credibility comes from accommodating no one. Volume II puts “you” back at the centre, because advice only means something once it lands on a specific person. The difference is measurable: **“you” appears more than four times as densely in Volume II as in Volume I.**

Why the two must sit together.

Volume I alone is a diagnosis with nowhere to put it. Volume II alone is advice cut loose from its grounds — why believe that low leverage matters most? Why believe physical skills have a closing window? Why is psychological resilience not just comfort food? The grounds for every one of those judgments live in Volume I’s hundred-plus data points.

A map and a pair of feet. Neither one gets you down the road alone.

If you are short on time: to understand what is happening, read Volume I; if you already understand and need to know what to do, start at Chapter XI. The load-bearing judgments in Volume II name the trend, the risk chain, or the class of outside evidence they rest on.

Volume I

The Extrapolation


If current trends are not interrupted, what happens. No single-point prophecies — a conditional extrapolation: observable data and trend lines already in motion, each followed to its end.

I. Introduction: From Cat-vs-Dog to Passing the Bar—A Decade Apart

In 2012 AlexNet won the ImageNet contest by “telling cats from dogs” [3]. Most people saw it as an academic breakthrough; few expected an industry-scale revolution to arrive this fast. Ten years later, on Nov 30, 2022, ChatGPT went live—1 million users in five days, 100 million in two months. It writes poems, code, and passes the U.S. bar and medical licensing exams. In one decade, AI leapt from classifying pets to acing the hardest human exams; our adaptation window is likely much shorter than ten years.


Each tech wave is faster: AT&T dominated telephony ~50 years, IBM mainframes ~30, Microsoft PCs ~20, Google the web ~15 [2]. This time the wave isn’t giving humans mechanical arms—it’s giving machines a human brain. The target has moved from hands and feet to cognition, and the clock is ticking.

Promise of this book: not sci-fi prediction, but data-anchored simulation: if today’s trends keep running, what happens in 5, 15, even 100 years across economy, geopolitics, demography, finance, and the species itself? You’ll leave with a set of leading indicators, a timeline, and an action checklist.



Each dominance cycle shrinks—your decision window may be under ten years.

Speed is the surface; the deeper story is which **human capability layer** gets eaten.



From muscles to minds: each wave ate a higher layer of humanity.

The layers: steam engines (coarse muscle) → electrification/assembly (rhythm) → industrial robots (fine repetitive motion) → PCs/Internet (calculation & memory) → deep learning (perception) → LLMs/agents (general cognition) → embodied AI (cognition + action). Each wave is faster and targets a higher layer.

In three short years, LLMs have gone from “chatty search” to autonomous coding/research/decision agents. Citrini Research’s Feb 2026 report “The Consequences of Surplus Intelligence” asks: standing in June 2028, what will we see? [1]

This book pushes further—a 100-year simulation across economic structure, geopolitics, population, finance, and the human species. It’s not a forecast; it’s a chain of “if trends continue” logic.

II. The Collapse of Intelligence Premium: When Human Brains Aren’t Scarce


2.1 The Human-Intelligence Displacement Spiral

For 200 years, economies assumed cognition is scarce. Lawyers earn six figures because few can pass the bar. Same for programmers, analysts, consultants. AI is ending that scarcity.

By late 2025, agentic coding tools jumped a step. A mid-level dev with Claude Code or Codex can rebuild a mid-size SaaS core in weeks. CIOs ask: “Why buy? Let’s build.” Long-tail SaaS (Monday, Zapier, Asana) feel it first; 30% renewal discounts are “good.” Even system-of-record vendors aren’t safe [1].

The deeper flywheel, per Citrini, is the **human-intelligence displacement spiral** [1]:

Stronger AI → layoffs → savings buy more AI → stronger AI → more layoffs




The intelligence displacement spiral—no built-in brake.

Enterprise AI budgets rise while total OpEx falls—AI replaces high-cost labor with low marginal-cost compute. No built-in brake.

Laid-off white-collar workers “downshift.” A \$180k PM drives Uber for \$45k [1]. One person downshifting is fine; hundreds of thousands flood service jobs and push wages down—overqualification becomes a gravitational crush on the bottom. And just as they shift to Uber and delivery, AV and drones are coming for those routes. Exit doors are closing.

2.2 Friction to Zero & the Collapse of Middlemen

AI agents eliminate consumer “laziness.” They comparison-shop, haggle, cancel subs. Any industry monetizing friction or asymmetry—travel OTAs, insurance renewals, financial advisors, real-estate brokers—faces extinction. Realtors’ fees could compress from 2.5–3% to ~1%—an “agent vs agent knife fight” [1].




Intermediary fee compression by industry.

DoorDash’s moat isn’t logistics; it’s “the app on your home screen.” Agents have no home screen—they pick the optimal option each time. Brand loyalty is a human inertia; machines have none. Brand premium is a tax on irrationality. Once agents hold the wallet, decades of brand equity can vanish in years.

2.3 Ghost GDP & the Composition Fallacy


Macro numbers may look rosy—nominal GDP up, productivity up. But Citrini calls it **Ghost GDP**: value that shows in accounts but never circulates, because machines don’t buy things [1].



The “scissors gap” between productivity and nominal wages.


The U.S. is 70% consumption; white-collar drive 75% of discretionary spend [1]. Who buys when they're jobless or paid less?

Classic **composition fallacy**: for any single firm, replacing staff with AI is rational; when all do it, they erase their own customers. The more efficient the firms, the faster the economy dies.



Individually rational $\neq$ collectively rational.

2.4 Autonomous Driving: Hand-Driving Will Look Like Horseback



The autonomous driving race: Tesla, Waymo, and China.

Horses→cars: 21.5 M horses in 1900; by 1912 cars outnumbered horses in NYC. 13,800 carriage makers in 1890 → 90 by 1920 (−99.3%) [83]. A whole profession vanished. AV is replaying this, but the “horse” is human drivers.

By Feb 2026, Tesla FSD logged **8 billion+ miles**, 1 billion in the first 50 days of 2026 [84]. Tesla self-reports one serious crash per 5.1 million miles vs U.S. average ~0.7 million [84]. On Jan 22, 2026 Tesla launched **safety-driver-free robotaxis** in Austin [84].

By 2026 the capability curve and the deployment curve had visibly separated. They must be told apart, because the binding constraint is on the second one.

Capability is compounding. Waymo started earliest: **450k+ weekly** paid rides by end-2025 [85]; by March 2026 it had driven **220.6 million** fully driverless miles with no human behind the wheel, was covering **4 million miles every week** by June 2026, and officially spanned **11 cities and over 1,400 square miles** [129]. **China is catching up hard:** Baidu Apollo Go covered **22+ cities** with 1000+ fully driverless cars and **17M+ rides** by end-2025 [85]; by June 2026 cumulative rides had passed **23 million**, at roughly **one airbag deployment every 14.4 million kilometres** [129]. Pony.ai holds fully driverless permits in all four tier-1 cities and WeRide runs 24/7 driverless in Guangzhou [85]. In H1 2026 Apollo Go launched fully driverless commercial operations in Dubai, obtained the world’s first fully driverless testing permit in a right-hand-drive, left-hand-traffic market (Hong Kong), and began open-road testing in London [129].

Deployment, however, is extremely concentrated. The same Waymo data by city: Phoenix 80.6M, SF

Bay Area 67.1M, Los Angeles 51.8M, Austin 15.8M, Atlanta 5.4M — **the top three total 199.4M, or 90.4% of everything** [129]. Years into commercial operation, **nine out of ten driverless miles ever driven by anyone happened in three metropolitan areas**. Not because the model cannot drive in Nashville — it is the same model. Every new city means redoing safety validation, obtaining regulatory approval, negotiating insurance, securing local political consent, and absorbing the first local accident. **The constraint has moved from the model to everything around the model — and everything around the model does not scale with compute.**

A common misreading is worth clearing up: Waymo and Tesla are two technical routes, but this is not a contest about reliability. Waymo fuses lidar, radar and cameras, leans on HD maps and validates city by city; Tesla is camera-only and bets on generalisation from fleet data. Tesla’s advantage is **per-vehicle cost and potential rollout speed** — a scaling advantage, not a reliability one. As of August 2026 Tesla’s no-in-car-monitor operation spanned six metros but the number of genuinely unsupervised vehicles remained in the dozens, against Waymo’s 220 million miles. **“Waymo has more incidents and recalls, so it is less reliable” is a statistical illusion** — those follow from operating at scale inside a regime where incidents are visible and reportable; fleet sizes two orders of magnitude apart cannot have incident counts compared directly. The routes are also converging: Waymo’s sixth-generation Driver on the Zeekr platform is an explicit cost-reduction move, and Tesla’s Cybercab is purpose-built.


BYD game-changer: 2025 “God’s Eye” ADAS **standard on all models**, even the \$9.5k Seagull [86]. By Nov 2025, **2.3 M** cars shipped with it, generating **150 M km/day** data [86].

Impact: **~100 M** driving jobs at risk globally [83]; U.S. **~5 M**, China **~40 M**, EU **5–7 M**, India **~15 M**. Goldman: U.S. could lose **25k driver jobs/month** in full adoption [87].

- **Urban land:** shared AV can replace ~11 private cars; LA parks/parking are 14% of land.
- **Insurance:** liability shifts from driver to maker/software.
- **Accessibility:** seniors/disabled mobility gains could create 2 M jobs & save \$19 B/yr healthcare [83].

Consensus: L4 at scale between late 2020s–2035; L5 maybe 2045–2060 [83]. Hand-driving in 20 years will feel like commuting on horseback.

2.5 Humanoid Robots: When AI Grows Limbs



Humanoid robots: China's hardware dominance (2025).

2025 saw the first real spike in humanoid robot shipments. ~90% came from Chinese firms [36]: AgiBot shipped 5,168; Unitree 5,500+ [37]. 140+ Chinese companies, 330+ models [36]. Tesla Optimus not yet commercial [37]. China's advantage: full manufacturing stack.

Key bottleneck: **autonomy**. Today's bots rely on scripted moves or teleop. Ken Goldberg notes: we have 100k+ years of language data for LLMs, near-zero physical manipulation data [38]. Also, **battery**: 2–3 hours, far from an 8-hour shift.

By 2026 the spike had more than doubled again—but in three diverging directions.

H1 2026 global humanoid shipments were roughly **19,100 units, up 272% year-on-year**; **Chinese vendors accounted for over 97% of global shipments and China for over 85% of demand** [127]. AgiBot overtook Unitree as the largest vendor (Unitree shipped 5,500 units in all of 2025). On **19 August 2026 Unitree listed on Shanghai's STAR Market, rising as much as 620% intraday and closing up about 460%**, touching a market value of **¥445 billion** against a ¥61 billion IPO valuation [127].

Divergence one: shipments tripled while the most-hyped player has produced zero. As of mid-2026 Tesla's Fremont line had produced **not a single mass-production Optimus**, despite decommissioning the Model S/X lines to make room. In January 2025 Musk said "will we succeed in making several thousand? Yes, I think we will"; by Q1 2026 that became "later this year." The first units go to an internal programme, "Optimus Academy," not to customers [127]. **Shipments up 272% while the loudest player ships nothing means volume is concentrated in the few Chinese firms that can actually build hardware—not that the industry has matured.**


Divergence two: the capability narrative and the technical reality are separating. Toyota Research Institute CEO Gill Pratt, April 2026: “we still don’t have system two”—robots have reactive intuition but not deliberative reasoning, and patching system one into system two keeps failing. He acknowledged “a lot of overpromising” and judged humanoids to be approaching “a peak of inflated expectations” [127]. Rodney Brooks, reviewing his own predictions in July 2026, pushed widely available dexterous hands out to **not-earlier-than 2030, by 2040**, writing that “deployable dexterity will remain pathetic compared to human hands beyond 2036” [127].

Divergence three, the one most worth recording: institutional resistance arrived before technical maturity. In July 2026 roughly **35,000 workers** struck at Hyundai’s Korean assembly plants—the **first auto factory strike explicitly triggered by a plan to put humanoid robots on the line** [127]. The trigger was Hyundai’s plan to deploy over **25,000** Boston Dynamics Atlas units, starting in the United States. Two details matter. **No deployment date had been set for the Korean plants at all**—the workers struck before the robots existed on site. And the union’s core demand was not a ban on automation but: **not a single robot enters a Hyundai workplace without a labour-management agreement first**. Their logic is clear-eyed: once the robots are on the line, the leverage is gone.

This matters beyond a labour dispute. It is the first real-world footnote to §3.5.1—as technical capability and legal accountability come apart, what people fight for is no longer “I can do this job” but **a place for personhood in the decision chain**.

Sim-to-real and rising battery density (8–10%/yr) may yield truly autonomous bots by ~2030. Then the displacement spiral expands from cognition to physical work—closing white-collar “exit ramps.”

2.6 The 35-Year Crisis Flips: AI Favors Seniors



The age crisis reversal: AI favors seniority.

Harvard 2025 study of **62 M workers / 285k firms**: adopting gen-AI cut junior hiring 7.7–10% within six quarters; senior roles unaffected—**Seniority-Biased Technological Change** [81].

U.S. top 15 tech firms' new-grad hiring —50% since 2019; grads now 7% of new hires [82]. Entry-level postings Jan 2023→Jun 2025 —**35%** [82]. Fortune 500 internships —22% (tech —34%) [82]. 54% eng leaders plan fewer junior devs [82]; Salesforce CEO: zero junior hiring in 2025 [82].

China: youth unemployment hit **18.9%** Aug 2025; 12.22 M grads, postings —22% [81]. AI roles demand seniors.


Logic: AI does junior tasks best (boilerplate code, data entry, reports, emails). One 10-year architect + AI ≈ self + 2–3 juniors. Companies pick senior+AI.

AI boosts juniors' productivity more (27–39% vs seniors' 8–13%) [81], but firms care about judgment/validation/ground truth—senior traits. Juniors become easier to replace.

“Vacancy chain” breaks at the bottom: no entry, no mid, no senior. The crisis shifts from “fired at 35” to “never enter at 25.”

2.7 The Degree Devaluation Cascade

Section 2.6 described the “can't get in at 25” trap. But the damage doesn't stop at the job market—it propagates upstream and hits the entire education system.



The degree devaluation cascade: from jobs to education to cities.

Degrees are losing their signal value. In 2025, only **30%** of new graduates found full-time work relevant to their major, a sharp drop from 2024 [92]. **33%** of 2025 graduates were still unemployed and job-hunting, up from 20% in 2024 [92]. Worse, **96%** of HR executives say universities need to do more to prepare students for work; **75%** say college education “simply doesn’t prepare people for jobs” [92]. When a degree is no longer a reliable employment signal, its value collapses—**63%** of Americans now say a four-year degree “isn’t worth it” [93].

The CS exodus is the canary in the coal mine. Fall 2025: CS graduate enrollment **down 15%**, undergrad down nearly 6%—**the first reversal in 20 years** [94]. 62% of CS departments report declining enrollment. Students aren’t fleeing tech—they’re fleeing what they think AI will replace: MIT’s “AI & Decision-Making” major is now its second-largest; U. of South Florida’s AI & Cybersecurity college enrolled 3,000+ [94]. 64% of pessimistic CS students blame generative AI directly [94].

Universities are closing. In the first 9 months of 2024, **28** U.S. degree-granting institutions shut down vs 15 in all of 2023. Philadelphia Fed worst-case: **80 closures/year** [95]. Southern Oregon U. cut 15 of 38 programs, including chemistry and economics [95]. When median degree ROI is 12.5%, **a quarter of graduates may never see a return**, and total cost (including opportunity cost) reaches \$180,000 [93], the product is losing its market.

The international-student economy collapse could trigger city-level chain reactions. This is a severely underestimated risk. Global study-abroad is a massive economic engine:

Country	Annual contribution	Share	2024–25 trend
Australia	AUD \$53.6 B	4th-largest export	Visa apps –28%
Canada	CAD \$22.3 B	1.2% of GDP	Study permits –48%
UK	GBP £9.4 B (tuition only)	20% of higher-ed revenue	Down –6.1% two years running
USA	USD \$43 B	Supports 355k jobs	Fall 2025 new enrollment –17%

The most fragile nodes are universities and cities hyper-dependent on international tuition. U. of Sydney: **47%** of revenue from international students [96]; U. of Melbourne: 33%, with 45% international student body [96]. Boston’s 63,000 international students generate ~\$3 B in local economic impact; two-thirds of high-end apartment leases go to international students [96].

Canada is already a preview. After study-permit issuances crashed 48% in 2024, **600+ education programs were canceled or suspended**, 4 colleges closed campuses, and 19 of Ontario’s 24 colleges laid off **8,000+ staff** [96]. That’s from policy tightening alone—before stacking degree devaluation on top.

The key: international students disproportionately enroll in STEM, business, and CS—precisely the majors hit hardest by AI. When a \$40k/year international degree no longer has clear ROI, **full-price international students retreat first**—and they’re the university’s most profitable segment. UK data is stark: **45%** of UK universities face deficits in 2025/26, projected 72% by year-end [96]. UK universities lose **£2,500/year per domestic student**—cross-subsidized by international fees. When international enrollment drops, the model breaks.

The education system is being forced to reshape. Change runs along four paths:

First, the four-year degree fragments into stackable micro-credentials. The U.S. passed the Workforce Pell Grant Act in July 2025, allowing federal Pell grants to cover 8–15 week short-term vocational programs for the first time [97]. 34 states have 111 programs with \$4.5 B invested in short-term non-degree credentials [97]. 96% of employers say micro-credentials add value; 90% would raise pay for holders [98]. The future model: 8-week certificate → stack into associate’s → stack into bachelor’s—buy as needed, not bundled.

Second, lifelong learning accounts replace one-shot education. The UK’s Lifelong Learning Entitlement takes effect 2026/27: each person gets a lifetime draw equivalent to 4 years of undergrad loans (~£38,140), usable anytime before age 60, managed via an online personal account [99]. Singapore’s SkillsFuture has upskilled 660k citizens. **China is the most aggressive:** a 2025–2027 plan to train **30 million workers** in vocational skills, with an “eight-grade technician” system tied directly to pay progression, using domestic LLMs to develop AI vocational courses—a global first [100].


Third, AI becomes education infrastructure. In 2024/25, 60% of U.S. K–12 teachers already use AI tools, saving 6 hours/week [101]. Student ChatGPT homework use doubled in a year (13% → 26%) [101]. By 2035, experts predict AI tutors will see, hear, speak, and know each student’s strengths and weaknesses, delivering precision-personalized instruction [102]. This may be education’s greatest equalizer—elite cognition barriers flattened by AI, just as AI flattened coding barriers.

Fourth, university value gets redefined. When AI democratizes knowledge access, a university’s unique value narrows to three things: peer network, internship pipeline, signaling. Top universities become “elite clubs”—Harvard’s value isn’t its curriculum but its alumni network and brand halo. Mid-to-low-tier universities still selling “knowledge delivery” will be outcompeted. Study-abroad value also

shifts: not knowledge (AI teaches better than professors), but visa + social network + cultural experience + local job access.

The chain summarized: **AI replaces entry-level jobs → degree ROI collapses → students flee / universities close → international students retreat → cities dependent on the study-abroad economy take a double hit → education shifts from “one-time 4-year investment” to “lifetime subscription.”** This isn’t a decade away. Canada is already there; Australia and the UK are in the middle of it.

2.8 Three Misaligned Timelines: Unemployment, Fiscal Pressure, Longevity



Three misaligned timelines: jobs crash first, fiscal squeeze follows, longevity diffuses last.

- Employment displacement arrives first: agents and autonomous driving concentrate in **2028–2035**, shrinking white-collar / driver / warehouse roles.
- Fiscal / pension stress follows: the tax base contracts from unemployment, but retirement periods lengthen from delayed retirement and longer lifespans—crunch window roughly **2030–2045**.
- Longevity breakthroughs come later: anti-aging / cell reprogramming therapies, even if successful, likely won’t reach mass adoption until **2040+**.

This means a plausible scenario of “living longer, but first passing through a gap with no stable income or safety net.” The problem isn’t longevity itself—it’s whether value distribution and social-security models can keep pace with the speed of displacement.

III. Drag Forces: What Could Slow AI

The constraints fall into three layers: physical, trust, and institutional. But they are not three parallel brakes — **they run on three different clocks, and that is what actually governs the pace of diffusion.** **Technical capability** compounds (compute, data, models). The **physical layer** (power, chips, atoms) is slow and linear (capital expenditure, construction cycles). The **institutional layer** (accountability, regulation, insurance) has **no clock at all — it is driven by events.**

The third is easiest to miss and most important: **the institutional layer does not improve with time; it improves with accidents.** It is not slow — it is not on the time axis at all. It sits still until something happens, then jumps discontinuously. The pattern recurs throughout: China issued a mandatory national standard for autonomous driving only after a multi-vehicle stall in Wuhan (§2.4); chatbot liability was established only once *Moffatt v. Air Canada* was decided (§3.3); Texas ordered an audit only after its interconnection queue reached five times record peak demand (§3.7); and **as of August 2026 humanoid robots still have no mandatory safety standard in the US, China or the EU (§2.5)** — not because they are safer, but because they do not yet have an identifiable accident with an identifiable victim.

The corollary: **the stronger the capability, the narrower the technical tail and the rarer the triggering events — but the larger the deployment at each trigger, and the worse the consequence.** The gap does not converge smoothly; it shows up as **steps that come further apart and land harder.**

3.1 Physical: Power, Chips, Atoms

Power: U.S. DC demand 75.8 GW (2026) → 134.4 GW (2030) [49]; grid queues up to 7 years. Only 5 GW of 16 GW planned 2026 capacity under construction; 30–50% large DCs delayed [49]. Amsterdam, Dublin, Singapore imposed moratoria.

Chips: TSMC says advanced nodes are $\sim\frac{1}{3}$ of AI demand [50]; Blackwell backlog 1+ year. New fabs: 3–5 years, \$20 B+; AZ fab fully running 2028, Kumamoto 2029 [50]. SMIC 5 nm via DUV multi-patterning, $\sim 30\%$ yield—far from economic 70% [50].

Humanoids: 2–3h battery; ISO 25785-1 safety standard in draft [52]. Figure AI faces safety lawsuit (testing sufficiency, clamp risk) [52]. Amazon warehouse robots correlated with higher injury rates due to human pace-matching [52].

3.2 Trust: Hallucinations, Last Mile, ROI Hangover


Summarization hallucination rates fell 21.8%→0.7% (2021→2025); factual QA still $\sim 33\%$ for o3 [53]. Courts recorded hundreds of filings with fake citations; a 2025 case had ~ 30 fabricated refs [53]. Deloitte submitted invented sources in a gov report [53].

Last-mile: AV handles 95%; the 5% edge cases consume 80% time/capital. Waymo: 14 M paid trips but still geo-fenced [54]. Rodney Brooks’ scorecard shows every “AV in a few years” promise since 2017 slipped; he projects humanoid dexterity far below human hands past 2036 [54]. Cruise shut down after $\sim \$10$ B burn (Dec 2024) [54].

Gartner 2025: GenAI in “trough of disillusionment”; agents at “peak of inflated expectations” [55]. MIT 2025: 95% of GenAI projects delivered zero ROI; avg \$1.9 M burn, <30% CEOs satisfied [55]. Yet AI investment still +44% YoY to \$2.5 T in 2026 [55].

3.3 Institutional: Laws, Labor, Legal Vacuum

The AI Act’s own timeline is the clearest evidence that institutions cannot keep up. Annex III high-risk obligations were due to apply from 2026; **six days before the deadline** the EU adopted the Digital Omnibus on AI (Regulation (EU) 2026/1744, in force 27 July 2026), pushing standalone high-risk duties (hiring, credit scoring, law enforcement, education, border control) to **2 December 2027** and high-risk AI embedded in regulated products to **2 August 2028**—because harmonised standards were not ready [56]. Only the Article 50 transparency duties landed on schedule on 2 August 2026 [56]. China’s focus shifted to anthropomorphic AI: the Interim Measures for AI Companion and Emotional Interaction Services (issued 10 April 2026, effective 15 July 2026) are the world’s first rules written specifically for AI companions; ByteDance and Alibaba pulled their AI companion features in early July 2026 as a direct result [56]. Labor pushback (Hollywood strikes) shows unions can slow deployment [57]. Agents rewriting contracts in code create new legal risks [58].



Three friction layers slowing AI diffusion

3.4 Fast vs Slow: Differential Diffusion Rates

Taken together, the three layers of constraint operate on very different timelines:

Friction factor	Severity	Estimated resolution
Power / grid capacity	Very high	5–7 years (grid queues 7 yr+)
Chip supply (advanced packaging)	High	2027–2029 (new fabs online)
Training-data exhaustion	Medium-high	2027–2029 (synthetic data partially mitigates)
AI reliability / hallucinations	Medium	Continuous incremental improvement
Enterprise integration friction	High	3–5 years (legacy-system overhaul)
ROI disillusionment trough	High	2–3 years (crossing the trough)
Regulatory compliance	Medium	2026–2027 (EU AI Act full enforcement)
Humanoid-robot maturity	High	5–10 years (battery, dexterity, safety standards)

What does this mean? **Pure-software AI (agents, coding assistants, text generation) diffuses fast**—marginal cost near zero, deployment needs no physical-infrastructure overhaul. But **AI in the physical world (robots, autonomous driving, DC buildout) diffuses much slower**—constrained by the physics of atoms and the inertia of institutions.

In other words: **white-collar winter arrives before blue-collar**. Code deploys globally in milliseconds; building a power plant takes five years; making a robot safely move boxes in a warehouse takes ten. This time gap—software fast, hardware slow—may be the most underestimated variable in the entire transition.

But beware: these frictions slow the speed, not the direction. Every bottleneck has a clear resolution path—grids are being built, fabs are being constructed, synthetic data is being generated, safety standards are being drafted. They determine whether AI displacement happens in five years or fifteen—not whether it happens.

3.5 Calm Before the Storm: Why 2026 Jobs Look Fine

U.S. Jan 2026 unemployment 4.3%; EU 5.9%; global 4.9% [90]. Yale/Brookings Oct 2025: **no observable disruption** since ChatGPT; AI-exposed job shares stable [90]. Reason: **lag**.

- Adoption lag: 88% claim “using AI,” only **7%** productionized [91]. <5% apps embed agents early 2025. Only 6% systematically capture value [91].
- Humans still needed as **Driver / Ground Truth / Glue**. Ground truth matters not just because AI hallucinates, but for a deeper reason: **humans have always-on sensory organs—eyes, ears, skin, nose—receiving real-time stimuli from the physical world; AI has none**. AI’s entire perception comes from already-digitized information—text, images, sensor logs—but vast amounts of critical real-world information have never been digitized: the state of a crash scene before anyone photographs it, the abnormal vibration of a machine, the hairline crack in a structure, the pallor of a patient’s face. Humans can walk to the scene and perceive reality firsthand; AI can only wait for someone to translate reality into data and feed it in. This isn’t a compute or algorithm problem—AI simply has no physical presence in the world; its body doesn’t exist yet. There’s also a subtler structural flaw: **AI agents have finite, lossy memory**. Current LLMs are bounded by maximum context windows—when a task grows long enough to fill that window, the model must compress earlier context to continue. Compression silently discards details: a key constraint discussed earlier, an edge-case handling decision, a design choice already agreed upon. The result is later output that

contradicts earlier work, and the AI is completely unaware. Anyone who has used an AI coding assistant on a large project has felt this: in conversation turn 200, the model confidently overturns a consensus reached in turn 30, because that memory was compressed away. Humans don't work this way—our long-term memory, while imperfect, retains critical decisions and context for months or years. In any complex task requiring long-term consistency, humans remain the only reliable “state holders.” As agents self-plan, RAG reduces hallucination, and automation handles integration, these roles erode—likely 2028–2029.

- Org delay: awareness (~1y) + process redesign (~1y) + legal HR cycles (~months). Agents boom 2026 → planning 2027 → layoffs show 2028 → full macro shock 2029–2030 [90].
- Gartner: **40% of agent projects canceled by 2027** [91]. Even 60% survival is enough to reshape labor.

The official stance shifted between February and August 2026. Governor Michael Barr's 17 February 2026 speech set out three scenarios—**gradual adoption**, **rapid growth in AI capabilities and adoption**, and **stalled growth in AI capabilities and adoption** (an investment bust is one way the third plays out, not a scenario of its own). His actual words were careful: “this evidence so far is consistent with what we might expect under the gradual adoption scenario I previously described.” He cited ADP payroll research finding that *early-career* workers in highly exposed occupations saw employment decline relative to peers and to experienced workers in the same field [90]. (The widely quoted “canaries in the coal mine” line is not Barr's—it is the title of a Stanford paper in his footnote 25.)

FOMC minutes then walked the language down [122]: March 2026, “**Many participants** cited evidence...firms were likely to **delay or reduce hiring in anticipation of AI adoption**”; April, downgraded to “**Several participants**”; June, “**Some participants** commented on the **possibility that AI could, over time**, affect employment prospects for some classes of workers”; July, “**Several participants observed that fears about AI leading to widespread layoffs had not materialized to date.**”

The shift coincided with a change at the top: from the June 2026 meeting the Fed has been chaired by **Kevin Warsh**, with Powell continuing as a governor [122]. In his 14 July 2026 congressional testimony Warsh said productivity growth had been strong “predating gains from AI adoption,” and that AI “hasn't displaced workers at this point.”

But “no evidence” depends on what you measure. The one contrary finding in the official camp comes from the Richmond Fed (Economic Brief 26-26, August 2026): workers in highly AI-exposed occupations show larger declines in **job-finding rates** [122]—not unemployment, not layoffs, but the probability of landing a job after applying. That may be the only shape a hiring-side shock can take in existing statistics: **nobody is fired; nobody is hired.** The July minutes describe the same thing from the other side—AI uncertainty is “keeping both hiring and firing low.” Unemployment statistics cannot see this. They count people with jobs and people without, not the job that was never posted.

Unpack this and the unemployment rate turns out to be the wrong instrument for this kind of shock. (Data retrieved 20 August 2026: US from the BLS public API, July 2026; Singapore from MOM, June 2026 preliminary, released 31 July 2026 [128].)

July 2026 US unemployment was **4.1%**, below January's 4.3%; Singapore in June 2026 was **2.0%** overall, **2.9%** resident, **3.0%** citizen. But set the other lines of the same US dataset alongside them [128]: participation fell from 62.1% in January to **61.4%** in July (−0.7pp); the employment-population ratio from 59.4% to **58.9%**; the labour force shrank by **1.37 million** while the not-in-labour-force count **rose**

by 1.98 million. Unemployment rate = searching but not employed ÷ labour force—the **denominator is shrinking**. Nearly two million people left the labour force in six months, counting as neither employed nor unemployed.

The numerator? JOLTS answers [128]: **layoffs rate 1.0–1.1%** (a historic low—inflow into unemployment is tiny), **hires rate 3.3–3.4%** (well below the pre-pandemic 3.8–4.0—outflow is small too), **quits rate 2.0%** (the employed do not dare move). Establishment-survey payrolls grew only about **44,000 per month** in H1 2026 against a trend norm above 150,000. **So the low rate is the product of “nobody is fired and nobody is looking,” not of “everyone has work.”** The Fed’s July 2026 minutes describe AI uncertainty as “keeping both hiring and firing low” [122]; the micro data above independently confirms that sentence.

The shock surfaces elsewhere: long-term unemployment. The share of unemployed out of work 27 weeks or longer [128]: 21.5% (Jul 2024) → 24.9% (Jul 2025) → 24.7% (Jan 2026) → **27.5%** (May 2026) → 25.5% (Jul 2026). **Few are unemployed, but those who are cannot get out**—the other face of the Richmond Fed’s finding on falling job-finding rates [122].

Singapore runs on a different mechanism, and the ordering of its three numbers is the evidence. Overall 2.0% sits *below* resident 2.9% and citizen 3.0%, because **non-residents enter the employment denominator but not the unemployment numerator**—foreign employees retrenched in a downturn must leave the country and never appear as unemployed. **A substantial part of Singapore’s labour-market adjustment happens outside its borders.**

Three limits, or this becomes selective evidence. **First, the two surveys disagree:** the household survey shows H1 2026 employment down 920,000; the establishment survey shows it **up 266,000** [128]. CES has the far larger sample, so **“US employment is falling” is not supported**—only that growth is weak. **Second, discouraged workers did not rise:** 509,000 in January, 503,000 in July [128], which **weakens** the reading that AI drove people out. **Third, none of this can be attributed to AI:** boomer retirement, slower immigration and AI-frozen hiring all coexist as explanations, and this book cannot separate them.

What can be established is one sentence: **the labour market is frozen—low hiring, low firing, low mobility—and the unemployment rate is structurally insensitive to a freeze.** It was built for twentieth-century layoff recessions. The present change is a hiring-side one: **nobody is fired; nobody is hired**, which registers as nothing happening. This also reveals a cross-effect between Chapters II and V: **demographic contraction and the AI shock push the rate in opposite directions and mask each other**, numerator and denominator contracting together. **“Unemployment is still low” cannot be read as “the shock has not arrived.”**

3.5.1 The Lights-Out Office: AI at the Centre, Humans at the Edge

Everything above—including the Driver / Ground Truth / Glue framework—assumes one direction: **the human is the principal, AI is the tool**. Invert that direction and much of what we already observe becomes easier to explain.

The inverted picture is: AI at the centre, humans at the edge. What runs continuously is the AI—it holds the state, runs the loop, does the work; the human is a peripheral it calls when needed. OpenClaw (§8.0) is already this shape: persistent memory (MEMORY.md), a heartbeat loop (cron plus agent loop),

its own values file (SOUL.md), overnight consolidation (Dreaming). What the human does inside that system is grant permissions, hand it information it cannot reach, and stand in wherever personhood is required. **We habitually call it “a digital employee,” but that is still the human-centric framing. Structurally it is a hub that never sleeps, occasionally borrowing a human’s identity and body.**

So what does the human peripheral still do? Unbundled, this is not one thing but four—and **the four decay at completely different rates.**

Layer one: collecting public information. Already lost. For anything digitised, AI’s retrieval, concurrency and cross-referencing beat any human.

Layer two: on-site observation. Being lost now. A car crash generates no data until someone photographs it; an abnormal vibration, a hairline crack, the colour and smell of a patient—these still require a body. But this layer is being eaten by cameras, sensors and the humanoid robots of §2.5. **It is a shrinking set, and the rate depends on hardware cost curves, not model capability.**

Layer three: extracting private information from other people’s heads. Structurally durable. Some information is not “not yet digitised”—it is **deliberately kept undigitised**: competitors do not publish real capacity; employees do not write down the judgement that would get them fired; a regulator’s real position is given over dinner; a customer admits the true reason for not renewing at the third drink; a founder tells you the remaining runway only if he trusts you. The bottleneck **is not the collector’s capability but the counterparty’s willingness to speak**—no model and no sensor solves that. It requires a specific person, carrying a relationship and symmetric risk, to sit down opposite them.

And here is the counter-intuitive part: **AI’s perfect memory is a liability in this layer, not an asset.** Off-the-record information flows precisely because it is deniable—“I never said that,” “we’re both exposed.” Facing something that never forgets, bears no consequences, and can be subpoenaed, people stop talking. **Speaking to an AI is speaking to a log file.**

Layer four: legal personhood and signature. The most durable. An AI cannot open a bank account, sign a contract, be sued, be jailed, or go bankrupt. The legal system rests on liability landing on a natural or juridical person. *Moffatt v. Air Canada* (§3.3) is an early data point—the court rejected the defence that the chatbot was a separate legal entity [58].

Layers three and four are two faces of the same thing. Someone tells you the truth off the record, and a contract requires a signature, for the same reason: **the other side has to be a person who can bear consequences.** A thing that cannot be held to account can neither sign nor be confided in. What the human retains at the edge is therefore neither senses nor limbs but **accountable personhood**—legally signature and liability, socially trust and relationship.

Centralisation has a dual face harder to deal with than efficiency: correlated failure. On the night of 31 March 2026, multiple Baidu Apollo Go driverless vehicles **stalled on the roads of Wuhan at the same time**, stranding passengers on busy overpasses [129]. Note the shape of that failure: not one car hitting someone, but **many cars failing in the same moment. A fleet of human drivers cannot all stall in the same second. A centrally dispatched robot fleet can.**

Existing safety statistics miss this class entirely. Baidu reports roughly one airbag deployment every 14.4 million kilometres — an excellent per-vehicle number, and **one completely blind to the whole fleet stopping at once.** Every “incidents per million miles” metric shares the blind spot: it measures the frequency of independent events, and correlated failures are by definition not independent.

The natural objection is that moving inference onto the vehicle removes this. **But that swaps one correlation channel; it does not eliminate correlated failure.** The direct counter-example is Waymo — **which already runs on-board inference and does not depend on the cloud for driving decisions.** NHTSA recall 26E035, dated 17 June 2026, records Waymo recalling **3,871** fifth-generation Automated Driving Systems; the defect is “the software may allow the vehicle to enter a closed freeway construction zone and continue driving at speed,” and the remedy notes that pending a fix “Waymo modified the scope of vehicle operations to restrict freeway driving” [129]. **Construction zones have nothing to do with the network** — that is the entire fleet sharing one blind spot: same model weights, same perception stack, same map data, same supplier firmware.

Counter-intuitively, **centralised inference is easier to hot-fix than distributed inference:** a cloud fault is rolled back once, while an on-vehicle fault needs an OTA campaign across thousands of vehicles — **and that, in regulatory language, is a recall.** After inference moves on-board, at least five correlation channels remain: shared software versions, shared sensor firmware, shared map data, shared environmental triggers (standing water, low sun, a new road-marking standard), and a dispatch layer that is still centralised. **Correlated risk does not disappear when compute moves onto the vehicle; it migrates from the live connection to the release pipeline.** The same holds for the agent ecosystem of Chapter VIII: a thousand agents on the same weights, the same SOUL.md template and the same toolchain fail in correlated ways too. **What centralisation buys in efficiency, it pays for by converting independent failures into simultaneous ones.**

But “personhood is irreplaceable” and “many humans are needed” are different claims. Senses and limbs are one-to-one: one site needs one body. Personhood is not—**it deduplicates.** One person can be the legal wrapper for many AI operations; one well-connected salesperson can be the “someone who can sit across the table” interface for dozens of agent-driven business lines. That is how “humans are still here” and “ninety percent fewer humans” hold simultaneously.

This also explains why the solo founder of §10.1 is fragile: in “one person plus AI,” **the human is the only component that gets sick, gets tired, breaks down, and has no spare.** AI does not degrade; humans do.

There is a ready analogy: the lights-out factory—which offers three concrete lessons. **First, it was not one robot per worker; it was a process redesigned around the machine’s constraints**—no lighting, aisles, ergonomics or shift breaks. Likewise the lights-out office will not be one agent per employee, but a process redesigned around the agent’s constraints: no meetings, no calendar alignment, and none of the intermediate artefacts produced purely so humans can follow along—weekly reports, decks, alignment docs, status synops. **A substantial share of today’s office workload is human-to-human interface overhead; agents do not need those interfaces.** **Second, the humans who remain change location and function**—out of the line, into the control room; from execution to exception handling. **Third, and most worth remembering: the lights-out factory has existed in prototype since the 1980s, FANUC has run unmanned machining for over two decades, and it is still a minority of manufacturing.** The last 5% of exceptions is extremely expensive to automate—the last-mile problem of §3.2. **Capability arriving is not organisational form arriving.**

One counter-intuitive corollary: **the lights-out office may arrive faster than the lights-out factory did.** A factory going dark requires solving atoms; an office going dark requires only bits. Following this section’s own logic—software AI replaces cognitive labour, embodied AI replaces presence—**the office goes dark first, the factory second,** the reverse of the popular intuition.

Finally, this answers a hole left by the title of Chapter II. **If the intelligence premium collapses, where does the premium go?** It moves to **trust and personhood**. What remains in the lights-out office is not the smartest people—intelligence is no longer scarce—but **the people with relationships and the people who can sign**.

(**Note:** this section is structural extrapolation, not a data finding. Its premises—continued AI capability gains, and legal systems continuing to require liability to land on a person—could both change within a decade. The ordering is falsifiable: an AI entity with independent legal personhood collapses layer four immediately; legal privilege for conversations with AI would sharply shrink layer three.)

3.5.2 Three Gaps: Asking the Question From the Other Side

Section 3.5.1 counted from the human side — which four boxes are left. **Counting the same thing from the AI side yields a version with more explanatory power: what does the AI still lack?**

Today’s large model is **a thing that is only a brain**. Its reasoning is already strong — strong enough to beat trained professionals on many closed tasks; its hardest current architectural limit is context length, not an inability to think. What actually holds it back is three other things it does not have.

Gap one: senses. It has no eyes, ears, nose, tongue, or skin. It can read the part of the world that has already been written down; it cannot read the part that has not — the abnormal vibration in a running machine, the colour of a patient’s face, what is missing from the shelf, the hesitation in a meeting that nobody said out loud. This is not a shortfall of judgement. **Its input is cut off.**

Gap two: authority. It cannot move an account, issue a contract, change a production parameter, approve a payment, or represent any organisation to the outside world. **This is the easiest of the three to underestimate, because it looks like a switch — and it is in fact the hardest constraint in this book.** The reason follows below.

Gap three: physical action. It has no hands or feet. Warehousing, picking, delivery, loading, on-site repair — those steps still require a person or a purpose-built machine.

These three gaps are the dual of the four boxes in 3.5.1. What the human on the periphery actually does maps to them one for one: gather on-site information (fills gap one), open permissions and sign (fills gap two), go there and do it (fills gap three). **Put differently: “human on the periphery” is not an identity. It is a list of shortfalls to be closed. How much of the list is crossed off determines which form the organisation has taken.**

The evolution of organisational form is simply the completion status of that list:

State of the three gaps	Organisational form
All three open	People are the subject, AI is the tool — the overwhelming majority of companies today
Senses closed	The one-person company: one human feeds it information, holds the authority, and runs the errands
Senses + authority closed	AI at the centre, humans at the edge: the human is left with signing and showing up
All three closed	The lights-out office

Chapter XI, 11.2 discussed the fragility of the one-person company: the human is the only component in that structure that gets sick, gets tired, and has no spare. **Placed on this table it becomes clear — the one-person company is fragile because it is a halfway state with only one gap closed, and the remaining two are loaded onto the same single person.** It is not a destination. It is what the list looks like half-crossed-off.

But the genuinely useful part of this table is that the three gaps close along completely different curves.

Senses track hardware cost. Unit prices for cameras, microphones and sensors keep falling and deployment density keeps rising. This one comes down on its own; nobody has to decide anything.

Physical action tracks robotics, which has a technical roadmap and an approximate timetable — section 2.5 discussed its three bottlenecks: endurance, dexterity, safety standards. And it has a very concrete leading indicator: **logistics, warehousing, delivery**. Those three settings are structured, repetitive, and forgiving on precision, which makes them the first place embodied intelligence can hold ground. **Watching those three is worth far more than watching humanoid demo videos.**

Authority tracks no technology curve at all.

It is the only one of the three that **technical progress will not loosen by itself**. No accuracy figure will make someone hand over payment authority. It is determined by where liability lands — who gets sued in law, who pays under insurance, who gets fired inside the organisation. And a pattern this book keeps observing is that **rules of that kind only move after something goes wrong**. China issued a mandatory national standard for autonomous driving only after a fleet stalled across Wuhan (§2.4); liability for chatbots was established only after the *Moffatt v. Air Canada* ruling (§3.3); Texas ordered an audit only after the interconnection queue reached five times the state's historical peak load (§3.7).

So the likely order in which the gaps close is: senses first, physical action second, authority last.

That yields a conclusion that is not especially intuitive:

The arrival time of the lights-out office is not set by model capability. It is set by who is willing to hand over the right to sign.

Long after capability is sufficient, the system will sit still — waiting not on technology but on a liability framework. This also explains the observation in 3.5.1: the prototype of the lights-out factory existed in the 1980s, and forty years later it is still a minority. **Not because it cannot be built, but because nobody wants to sign for the last 5% of exceptions.**

And if all three gaps really do close, something appears that has not existed before.

An organisation that has senses, authority and physical action, wired into a continuously running AI system, holds every component needed to close a loop: it can perceive external state, decide, make the decision take effect in the physical world, observe the result, and revise its own next move accordingly. **At that point a "company" stops being merely a legal container for employing people and becomes a system that iterates and updates itself.**

The threshold for that step still lands on authority — because "modifying itself" means having the right to change its own rules, and that is something human organisations have never handed over.

(Note: this section is a structural extrapolation, not a data conclusion, and it is falsifiable. Three

tests. One: if some heavily regulated industry achieves large-scale autonomous authorisation before a lightly regulated one does, "authority closes last" fails. Two: if automation penetration in logistics/warehousing/delivery stalls for a long time while office automation keeps advancing, the ordering of "action before authority" has to be swapped. Three: if an AI entity with independent legal personality appears, the authority box collapses instantly and the whole ordering is void.)

3.6 Dashboard: Six Indicators, Three Thresholds

Projection easily becomes self-serving. So here is a set of indicators that quarterly data can falsify directly—each with a numeric threshold rather than a defensible-either-way adjective.

- **(1) Entry-level job postings, YoY:** green $> -5\%$; amber -5% to -15% ; red $< -15\%$. First visible signal of AI substitution.
- **(2) Employment-population ratio, 6-month change:** green $> -0.2\text{pp}$; amber -0.2 to -0.5pp ; red $< -0.5\text{pp}$. **Cannot be flattered by labour-force exit** (see §3.5).
- **(2b) Long-term unemployed (27+ weeks) share:** green $< 22\%$; amber $22\text{--}26\%$; red $> 26\%$. Whether the unemployed can get out matters more than how many there are.
- **(3) Corporate AI CapEx / revenue:** green if rising without crowding out R&D or headcount; amber if it crowds out R&D; red if it crowds out headcount. Separates healthy investment from layoff arbitrage.
- **(4) Frontier accelerator lead time:** green falling; amber flat; red rising. Whether compute supply keeps up.
- **(5) Pension solvency (years payable):** green $> 8\text{y}$; amber $5\text{--}8\text{y}$; red $< 5\text{y}$. Urgency of shifting the tax base off wages.
- **(6) Housing vacancy + regional bank capital:** green both stable; amber one deteriorating; red both. Early warning on the mortgage chain.

Note that (2) uses the employment-population ratio and long-term unemployment share, not the unemployment rate. §3.5 shows why: in a frozen labour market the unemployment rate moves *down* as people exit, giving the opposite of the true signal.

Trigger rule: three or more red lights for two consecutive quarters → switch from “incremental reform” to “emergency intervention.” Don’t wait for the crisis to act; don’t panic while the data are still green. All six are publicly checkable (BLS, OECD, IMF GFSR, Fed regional banking data). They measure whether society is being hit; how fast AI is diffusing is handled next.

3.7 Brake Conditions and Scorecard

They share one hard requirement: **every brake must be judgeable from public data as either pressed or not pressed.**

That requirement is stricter than it looks. An indicator that cannot be adjudicated is **worse than no indicator at all**—it manufactures the impression of monitoring while in practice any outcome can be read as consistent with expectation. Two worked examples follow the list.

- **Compute supply tightens** (frontier accelerator lead times back above 12 months): **not triggered**. Nvidia’s Vera Rubin platform entered full production and began shipping to cloud providers on 21 July 2026—supply is accelerating [123].
- **Power and interconnection** (an **administrative** interconnection freeze in a major data-center state): **triggered**. On 3 August 2026 the Governor of Texas ordered the PUCT and ERCOT to audit every data center in the interconnection queue, warning that projects failing to disclose ownership, financial, water and community-impact information could be denied grid access. The ERCOT large-load queue reached **474.7 GW** in June 2026, of which **420.8 GW (90.2%) is data centers**—more than five times the state’s all-time peak demand of 91,089 MW (22 July 2026) [123].
- **Autonomous driving (time required to enter each new city stops falling as the technology improves)**: **triggered**. As of March 2026, **90.4%** of all rider-only miles ever driven remained concentrated in Phoenix, the SF Bay Area and Los Angeles [129]. This replaces the previous “city count stalls” test — city count *is* the governance variable, so using it to test whether the constraint is governance is circular; “does time-per-city fall?” actually discriminates.
- **Humanoid robots** (unit price >\$30k and runtime <4h persisting to 2030): open (see §2.5).
- **Macro demand** (easing triggers real reflation, forcing AI CapEx cuts): **not triggered**. As of mid-2026 no hyperscaler had cut its 2026 AI CapEx guidance [123].

Finally, a prediction worth betting on. If “the institutional layer improves only with events” holds, it should predict something that has not yet happened: **the humanoid robotics industry will acquire a mandatory safety standard within 12 months of the first widely reported fatality, and not before — however good the robots are by then.** The bet is live as of August 2026: OSHA’s robotics page still reads “There are currently no specific OSHA standards for the robotics industry”; China has a mandatory national standard for autonomous driving but its humanoid-robot safety standard remains voluntary and in drafting; ISO 25785-1 is still a committee draft [129]. **Meanwhile the sector’s global shipments grew 272% year-on-year in H1 2026.** The gap between capability and accountability is widening visibly.

How to read this list—two notes.


First, why the power-and-grid item is written as an *administrative* restriction rather than as “the queue got longer.” The more intuitive phrasing would be “average grid-interconnection wait exceeding 5 years”—it sounds concrete, but it is a bad indicator: **no institution publishes an authoritative “average wait” statistic**, so the condition can be neither confirmed nor refuted. It looks like a threshold; it is only an adjective. And the brake, when it was actually pressed, arrived in an entirely different form: not longer queues, but a **political stop**—an event with a date, an administrative order and a paper trail. **That is precisely what the “adjudicable” requirement is for: the real turn rarely arrives in the shape you predicted, so the indicator has to track events rather than a phrasing.**

Second, the first item—lead times back above 12 months—is currently *refuted*. That is a good sign: it means these conditions can be contradicted by reality rather than absorbing whatever happens. **A set of indicators that is always right is a set with no information in it.**

IV. Geographic Differentiation

4.1 U.S.: Eye of the Storm

San Francisco office vacancy 34% [15], while AI firms lease 1 M+ sqft; Mission Bay vacancy <9% [15]. Politics polarized; 2028/2032 election cycles clash with policy windows.




San Francisco: a dual-speed city

4.2 China: Different Script, Similar End

After “hundred-model” shakeout, 3–5 majors likely remain. Manufacturing+AI accelerates automation; open-source shopping agents (e.g., Qwen ecosystem) could hand consumer decisions to AI. Rapid redistribution is an institutional edge. Cultural shock: “hard work pays” loses validity. Aging + youth “lying flat” + AI job squeeze → forced “lying flat 2.0.”

4.3 India: Digital Colonization Cascade

IT outsourcing (\$200 B) sells “cheap humans”; agentic coding drives contracts away [1]. Middle-class cities (BLR/HYD/Pune) tied to IT collapse; housing/consumption/local banks follow. The IIT→TCS/Infosys→H-1B→Silicon Valley ladder breaks. Loss of elite immigration hurts U.S. innovation. AI compute concentrated in U.S./China (~80%) [59]; others become consumers, not owners.



India's four-layer impact cascade




Global AI compute distribution

4.4 Europe: Regulator's Dilemma

Strongest on rules (AI Act), weakest on frontier models. Welfare state may cushion shock; Finland's basic income pilot shows stress relief, limited employment effect [33].

4.5 Developing World: Advantage Erased



AI reshapes global division of labor: compute = power.


Cheap labor advantage dies as robots near \$0.75/hour vs Bangladesh \$0.5/hour [60]. Reshoring: U.S. 2023 reshoring/nearshoring roles 287k (27× 2010) [61]. Country snapshots: Bangladesh garments —60% jobs risk; Philippines BPO 1.3 M at risk (89% high automation) [62]; Kenya BPO 2.5 M at risk [62]; Vietnam assembly squeezed [62]. Speed is the killer—software diffuses in years, not decades.

V. Population Decline: Natural, Not a “Crisis”

Longevity won't backfill labor. Early longevity gains accrue to wealthy; it increases elderly headcount without increasing workers, enlarging fiscal burden when AI already shrinks wage tax base.

5.1 Fertility Push Fails


Korea TFR 0.72 (record low) [5]; China ~ 1.0 ; Japan 1.2; Italy/Spain 1.2 [4][5]. Global TFR 5.0 (1950) $\rightarrow \sim 2.3$ now; UN projects 1.8 by 2100, population peak ~ 2084 [6]. Korea spent \$280 B+ on pro-natal policies; TFR fell further—74% of subsidies went to people who would have birthed anyway [9]; to reach replacement would need $15\times$ budget [9].



Fertility rates across countries

5.2 Economics of Marriage/Childbearing Collapse

Marriage rates plunge; costs explode (raising one child in China $\approx 6.3\times$ per-capita GDP, 538,000 yuan ($\sim \$74\text{--}80\text{k}$)) [27]. ROI of children turns negative once AI lowers marginal value of human labor.



The economics of marriage and childbearing

5.2.1 The Unbundling of Marriage: Four Things Were Tied Together, Then Loosened One by One

The cost-benefit arithmetic above explains "why people are not having children." It does not explain the other half – **why people are not getting married at all**. To answer that, you first have to take the word "marriage" apart.

In the old structure, this single status carried four things at once: **economics, love, sex, reproduction**. The single label "wife" or "husband" referred to four merged roles. **But these four things are entirely different in nature and follow entirely different laws** – bundling them was an arrangement suited to specific conditions of production, not a law of nature.

First, economics. Marriage is above all an economic institution. It is a long-term cooperation contract: pooled expenses, shared risk, division of labor in production. Under conditions where one party (usually the woman) could not readily earn an independent income, that contract was rational for both sides. In many traditional societies the primary premise of marriage was not romantic love in the modern sense, but an arrangement over lineage, property, labor, and childbearing.

Second, love. Here two different things have to be separated: passion and attachment. Passion is a chemical storm, and it has a shelf life – even for two well-matched people, the state of being in love rarely survives past three to five years. The phrase "the seven-year itch" endures because it describes the norm, not the exception. **What follows is not necessarily the end of the relationship, but its replacement by something else: attachment, commitment, the cooperative work of a shared life** – what many people describe as "like family," "like comrades," "like roommates." **The point is not that love dies, but this: marriage cannot lock in passion. What it actually rests on is whatever remains after**

passion recedes.

Third, sex. It is more complicated than the first two, but its direction is clear: **it detached from the exclusive arrangement of marriage relatively early and relatively visibly**, and it has market and quasi-market substitutes.

Fourth, reproduction. It requires long-term cooperation and a large commitment of resources, its cost keeps rising with development, and its returns grow steadily more uncertain – exactly the arithmetic run at the start of this section.

Once the four are separated, everything that has happened over the past few decades reduces to one sentence: they have been coming unbundled, one at a time.

Once women could earn independent incomes, **the strongest layer of economic dependence inside marriage weakened** – for many women, marriage stopped being an institution required for survival. Passion never needed marriage to begin with, and marriage never locked it in. The availability of sex kept rising. **Of the four, the one still strongly bound to a long-term contract is mainly reproduction.**

And that remaining one happens to be the most expensive of the four, with the least certain return.

When the least substitutable part of an institution's value is reduced to its most expensive function, demand for it stops being inelastic.

Three things have to be said honestly. **First, this does not mean marriage is "useless"** – it still supplies default legal rights, arrangements for property and inheritance, a kinship network, social recognition, and the institutional force that binds down responsibility for raising children. None of these are among the four above, and none of them have come unbundled. **Second, what this framework therefore supports is "declining exclusivity and declining inelastic demand," not "marriage will disappear."** **Third, there are competing explanations for falling marriage rates** – housing prices, employment insecurity, later marriage and childbearing driven by longer schooling. The mechanism in this section does not exclude them; it only offers one thread they cannot explain: **why it is the more educated and more economically independent groups who exit most completely.**

The framework incidentally explains several things that would otherwise each need their own explanation. **Why phrases like "a child is the crystallization of love" stop working** – they assume the four things are one thing; once separated, the sentence fails to line up on any layer's timescale (passion runs in years, reproduction in decades). **Why marriage rates fall as women's education rises** – the layer she loosens first is precisely economic dependence.

And what AI does inside all this. Here the word "economics" has to be split into two layers, or the argument contradicts itself: one layer is **dependence on a spouse's provision**, the other is **the real co-operative work of two people carrying housing, childrearing, care, and risk together.** **What AI can accelerate the replacement of is companionship and emotion, plus part of sex. It cannot replace that second layer of real-world cooperation, and it cannot replace reproduction.** So AI does not stop people from marrying. It **narrows the remaining reasons for marriage further down onto reproduction and real-world cooperation** – which happen to be the two heaviest and most expensive.

5.2.2 Equality Is a Companion of Development, and So Is the Tension

The unbundling was nobody's choice. It is a result of development.

Economic development is generally accompanied by an expansion of women's education and their opportunities for labor force participation, and on the education dimension the reversal is complete: OECD data for 2024 show tertiary attainment among 25-34 year-olds at **55% for women and 42% for men, a 13-percentage-point gap in women's favor** [142], and this is the shared direction across the large majority of member states, with Mexico among the few exceptions.

Be precise here: an education reversal is not equality across the board. Income gaps, the motherhood penalty, the division of housework, bargaining position inside a marriage – none of these have reversed in step, and some have barely moved. **What this section claims is narrower: on the two dimensions of education and economic independence, women's position has shifted enough to change the marriage market.**

Cross-national evidence supports this. Esteve et al., studying 420 samples across 120 countries from 1960 to 2011, found that as women's educational advantage rose, matches in which the wife is more educated than the husband rose in step – **men's upward-matching advantage on the education dimension has already ended in many countries** [143]. An analysis of 27 European countries across seven birth cohorts points the same way: upward matching declines by cohort and downward matching rises [144].

The result is an internal tension: the marriage-and-childbearing structure premised on "man provides, woman does the domestic labor" cannot easily retain its old stability alongside the continued rise of women's education and income.

Note the qualifier in that sentence. It does not say "equality and childbearing cannot coexist" – there have always been multiple versions of the childbearing structure: male-breadwinner, dual-earner, state-childcare, grandparent-supported. **It says that one particular version is losing the conditions it ran on.**

There is a popular explanation to deal with here: **"only half-equal"** – the public sphere (education, employment) equalized, the domestic sphere (housework, childcare) did not keep up, so fertility collapsed; and the Nordic countries, being relatively equal inside the home too, held their fertility. **That explanation looked right in the 2010s, but its strong version has been overturned in the past decade: total fertility rates across the Nordic countries have all fallen to historic lows, and the gap with East Asia has narrowed sharply.**

But state precisely what that overturns and what it does not. It overturns the **strong version**: that domestic equality alone is enough to hold fertility up. It does **not** overturn the weak version: unequal domestic divisions of labor do depress fertility. **Both can be true at once – domestic equality can cushion the decline without being a sufficient condition.**

So what is the rest of it? **Once marriage and childbearing shift from a condition of survival to an optional arrangement, individuals reprice them using opportunity cost, relationship quality, and future risk**; when expected returns fail to cover the cost, entry rates fall. **This is not a change in morality. It is a change in the pricing benchmark.**

5.2.3 The Chinese Version: Urbanization Made the Structural Mismatch Spatial

The same mechanism in China has an extra layer on top of it – **extremely fast urbanization.**

China's urban share of population rose from **19.39%** in 1980 to **67.9%** at the end of 2025 [145]. This is

one of the largest and fastest population migrations in human history, and it was not gender-neutral.

The distribution of fallback resources back home is asymmetric. In law, Article 1126 of the Civil Code provides for equal inheritance rights regardless of sex, and Article 6 of the Rural Land Contracting Law provides women equal contracting rights with men. **But there is a substantial gap at the level of enforcement:** the Third Survey on the Social Status of Women in China (2010) found 21% of rural women without land against 11% of rural men; in national surveys of the same period, 92.4% of men reported contracted land held in their own name against 81.2% of women, while the share who lost land due to a change in marital status was 8.1% for women and 1.1% for men [146]. **The accurate statement is: men are on average more likely to retain some portion of their fallback resources back home, and women are more easily excluded from land and village entitlements** – this is a difference of degree, not a binary of “has” versus “has not.” Whether there is a way back also depends on the parents’ housing, homestead eligibility, local employment, family acceptance, and village rules; land is only one item among these.

That asymmetry leaves a clear trace in the marriage data:

- **Never-married older men are far more concentrated in rural areas than in cities.** The 2020 census recorded 8.82 million never-married rural men aged 35 and over, **5.7%** of that group, against **3.62%** in cities. Among never-married men aged 50 and over, rural areas exceeded cities by **2.03 million** [147].
- **Marriage probability drops sharply for highly educated urban women.** Research based on the Chinese General Social Survey and the China Family Panel Studies found that among urban women over 27, a university degree lowers the probability of marriage by **2.88%-3.6%**, and a postgraduate degree lowers it a further **8.4%-10.4%** [148].
- **And marriage migration runs one way.** A study framed around “bride drain” found that the outflow of rural women significantly lowers rural men’s probability of marrying and amplifies the effect of the sex-ratio imbalance – **the study found no corresponding “groom drain”** [149].

Be careful not to flatten this into a static opposition between two groups: **large numbers of less-educated women are also in urban service work, and large numbers of highly educated men are also concentrated in big cities.** What is being described is not the absolute distribution of the population but the **relative mismatch in the marriage market** – unmarried men in counties and rural areas are relatively weak on education and income, while urban highly educated women’s search radius, income expectations, and way of life run past them.

The substance of the mismatch is educational structure; urbanization made it spatial. A 2026 study estimates that between 1999 and 2019 the probability of a single person in China marrying **fell to half its 1999 level**, and that **changes in the supply structure of the educated population explain roughly half of that decline** [150]. Note what this evidence points at: **an educational-structure mismatch.** The geographic dimension is how that mismatch shows up in space; the two are not the same thing and should not stand in for each other.

What has to be said at the same time: this cannot be read as “so leave the property to the son.”

First, put the strongest version of the opposing case on the table: under a system where the man carries the burden of the marital apartment and the bride price, a family concentrating resources on the son **may well be rational at the level of that single family** – it raises the son’s relative competitiveness

in the marriage market. That micro-level judgment probably has empirical support, and this section does not deny it.

But it does not scale to a macro conclusion. There is no evidence that generalizing the strategy raises the rate of successful matching across society; on the contrary, generalizing it reinforces son preference. At one point China had roughly 32 million more males than females under the age of 20, and the marriage market was made worse by that collective behavior itself [151]. And in a market where women are scarce, the chance of entering marriage concentrates among higher-income men and compounds with stratification by income, housing, and education – **the result is not better matching overall, it is sharper stratification.** Evidence in the opposite direction exists too: after China’s 2001 divorce reform strengthened women’s property rights, families whose first child was a girl were **8.1 percentage points** less likely to go on to have a boy [152].

So the conclusion of this section is not “son preference has a point.” It is the opposite: a marriage-and-childbearing contract premised on a gendered division of labor and on women forgoing opportunity **cannot hold up in the face of development** – you can neither stop women from being educated and earning, nor, once they have the ability to survive independently, expect them to keep signing that contract on the old terms.


It also has to be clear where this chain stops. Falling fertility is not a phenomenon that any single policy failure explains; it is the joint product of development, women’s independence, the cost of marriage and childrearing, and institutional mismatch. **Policy can change the slope; it can rarely restore the old structure** – the failed pro-natalist cases in section 5.1 of this chapter are the evidence.

And what it supports in Chapter IX is limited: population decline by itself does not imply a silicon-based succession. What it can do is weaken one default premise – that human numbers will keep expanding. The species divergence and non-carbon continuation that Chapter IX takes up require a separate argument; this section supplies demographic background, not a conclusion.

(Falsification condition) The mechanism in this section can be overturned. **If, in a society with high female education, high female income, a relatively equal domestic division of labor, and sharply reduced housing and childcare costs, the total fertility rate returns durably to near replacement level and the marriage rate recovers in step, then the mechanism “marriage and childbearing becoming optional causes low fertility” needs to be rewritten.** Conversely, if fertility rises only slightly after all those conditions improve and cannot return to old levels, that indicates the main driver really is post-development opportunity cost and the unbundling of marriage from childbearing, rather than the degree of equality itself. This is a criterion testable against the next ten to twenty years of data from the Nordic countries, East Asia, and parts of Central Europe.

5.3 AI–Fertility Spiral

AI job loss → income instability → no kids → population falls → society more AI-dependent → more job loss.



The AI–fertility decline spiral: no natural brake.

Men’s job loss hits marriage market (male identity tied to income); data show +1% male unemployment → −1.5–2% marriage [88]. AI companions further reduce marriage’s functional value.

5.4 Robot Caregiving

Japan/Korea already subsidize care robots [10][11]. In 20 years, carebots will be far better; the right question may be “how to run society with fewer people,” not “how to force births.”

5.5 Implications for Policymakers: Production Is Not the Problem—Distribution Is

Pro-natalist spending is a proven failure. Korea spent \$280 B; fertility kept falling. No subsidy has reversed the trend.

The correct investment is AI and robotics—the real replacement for the “demographic dividend.” A humanoid robot working 24/7 is a far greater productive force than a child who won’t join the workforce for 20 years. China’s 30-million-worker AI training plan [100] and 90% global humanoid shipment share [36] point the right way.

But production is not the hard part. Distribution is. When robots replace labor, GDP may soar—but growth **does not flow through paychecks**. The traditional circuit (Labor → Wages → Consumption → Revenue → Taxes → Social services) breaks at the very first link. GDP grows; money doesn’t reach people.

Policymakers need a new distribution system—channeling AI-created value to the human consump-


tion side. Options: compute/robot tax (prisoner’s dilemma on who moves first); AI output dividends (Alaska Permanent Fund as precedent); UBI (Finland pilot showed well-being gains [33]; funding via compute/data taxation); lifelong learning infrastructure rather than birth subsidies [99][100].

In one sentence: don’t make more people to compete with machines—make sure machines’ wealth flows to people. The production side is being solved. The consumption-side distribution problem determines whether the AI era means prosperity for all or bankruptcy for all. When mass AI layoffs arrive (estimated 2028–2030), an unprepared distribution system turns “Ghost GDP” into a real social crisis.

VI. Financial Chain Reaction

6.1 Mortgages: Good Loans in a Bad World

U.S. mortgages \$13 T [1]; borrowers prime. 2008 was bad loans day one; now loans were good, the world changed. SF home −11% YoY; Seattle −9%; Austin −8% [1]. “Zombie payers” keep paying by cutting all other spend.



2008 vs 2030: two mortgage crises

But this is just the U.S. Zoom out—this is a **\$25 trillion**-scale cross-border problem. U.S. \$13.1T (2026Q2, NY Fed) + euro area EUR 5.50T (~\$6.4T, June 2026, ECB) + China CNY 36.29T (~\$5.4T, 2026Q2, PBoC) = roughly **\$25 trillion** across the three major blocs (ECB August 2026 reference rates: EUR 1 = \$1.1605, \$1 = CNY 6.74). This is not an external statistic but the sum of three official series;

adding the UK (~GBP 1.7T) and non-euro European markets pushes it higher. The U.S. crisis is still hypothetical. China's has been running for five years.

China: present tense, not future. Evergrande (>\$300B debt, liquidation ordered Jan 2024, delisted Aug 2025; on 20 August 2026 a Shenzhen court sentenced founder Xu Jiayin to life imprisonment and fined the group ~¥15.8B), Country Garden (~\$11B of dollar bonds defaulted Oct 2023; the \$17.7B off-shore restructuring took effect 30 Dec 2025, cutting ~\$11.7B of debt), Vanke (2025 net loss **¥88.6B**, the largest single-year loss by any A-share developer; S&P cut it to “selective default” and Fitch to “restricted default”—a distressed exchange, not a default avoided; ~\$21B short-term debt vs ~\$8B cash at end-Sept 2025)—systemic collapse. Sunac, CIFI and Sino-Ocean also completed offshore restructurings by end-2025, but writing debt off is not recovery: Jan–Jun 2026 property development investment fell **18.0%** YoY and new starts **23.4%** [118]. The pre-sale system let developers spend buyers' money before building; in summer 2022 buyers at **320 projects across 100+ cities** stopped paying. Nomura estimated **20 million** units sold but unbuilt at end-2022, needing ~¥3.2T to finish; Goldman put clearing inventory to 2018 levels at over ¥7T. Whitelist bank loan approvals rose from a ¥4T target at end-2024 to over **¥7T** by September 2025, delivering **7.5M+** stalled units [119]. New-home prices fell **3.2%** YoY in Feb 2026; by July the decline had concentrated in lower tiers and resale—new homes –1.1% (tier 1) vs –4.2% (tier 3), existing homes –3.7% vs –5.8%: **the repair is happening only in core cities**. Full-year 2025 new-home sales fell to **¥8.39T** (–12.6%), about 46% of the 2021 peak [120]. Vacancy: **65–80 million** empty homes [20]. Household deleveraging is accelerating: mortgage balances fell from ¥37.01T (–1.8% YoY) at end-2025 to **¥36.29T** (–3.8%) at end-2026Q2, household leverage down to **57.7%** [121]—classic balance-sheet recession, mirroring 1990s Japan, except Chinese households started before the AI shock even arrived.

Now layer on AI. China faces a **triple overlay**: (1) existing structural crisis (unfinished projects, developer defaults, falling prices); (2) accelerating demographic decline (population shrinking since 2022); (3) incoming AI employment shock (middle-class income loss → mortgage defaults). The U.S. faces “good loans in a bad world”—hypothetical. China faces “structural flaws meet demographic collapse meet AI displacement”—already happening, about to worsen.

6.2 From Income Repricing to Municipal Finance: Three Different Paths

The minimum claim of this section and the next, stated first — every qualification that follows narrows it rather than retreats from it:

AI will not empty out cities in general. It raises the financial fragility of one class of city — high housing costs, high white-collar exposure, high household leverage, municipal finances dependent on property and land — which may show up first as mortgage stress, consumption contraction, an eroding tax base, and a change in the occupational mix.

The first link is easy to skip past: **in that particular class of city, the kind of labour income AI reprices first overlaps heavily with the kind of labour income that services large mortgages.**

They are not the same set. High-mortgage households also include doctors, civil servants, contractors and business owners; many people with high AI exposure are young renters with no mortgage at all. But in high-cost cities concentrated in tech, finance, law, consulting and back-office operations, the two sets overlap substantially. Note also that plenty of the “stable cash flow” banks prefer carries **low** AI expo-

sure — doctors, nurses, teachers, public employees, utilities. The genuine overlap is narrower: a fixed employer, verifiable income, a predictable raise curve, **and** work composed mainly of standardizable cognitive tasks.

One precondition has to come first, or the whole chain is empty: high exposure does not mean net harm. AI can substitute for a role or complement it. The same "high exposure" can end in redundancy or in doubled output and higher pay. **The chain below only starts if the substitution effect exceeds the complementarity effect and produces observable income discounts or employment instability.** That cannot be settled in advance; it can only be observed.

Given that, the steps from "an occupation is repriced" to "a bank is in trouble" cannot be skipped: **rising exposure → greater income volatility, shorter layoff cycles, wage discounts on re-employment → cash-flow gaps in highly leveraged households → default.** The concentrated risk is not "all white-collar workers" but **households with high leverage, a single income, or both earners in the same affected industry.**


The 2008 subprime problem was "these loans should never have been written." This time the problem is "the loans were correct when written, and the borrowers' occupational category is being repriced."

"Bank crisis" needs to be unpacked, because three risks travel along different paths.

First, commercial real estate — and note that this is a parallel pressure, not a direct consequence of an AI income shock. Nearly \$1T of U.S. commercial real estate loans mature in 2026 [16], a fifth of it offices, priced on the assumption that offices would always be occupied; San Francisco office vacancy is 34% [15]. **But be precise about causes: office vacancy comes mainly from normalized remote work, higher rates and valuation repricing — all of which preceded any large-scale AI displacement of white-collar work.** Charging it to AI's account overstates AI's explanatory power. It belongs in this section for a different reason: **it lands on the same banks' balance sheets as the AI income shock,** and of the three it is the one that hits regional bank capital most directly, because local development and commercial loans are concentrated there.

Second, residential mortgages — a less direct transmission, and one where national mortgage design sends the same shock down completely different paths. U.S. residential lending is dominated by 30-year fixed rates, so "rate reset lifts the payment" is not the American mechanism: **in the U.S. the pressure can essentially only come from the income side.** In floating-rate or short-reset markets (the UK, Australia, China, Hong Kong) rates and income are two independent triggers that can stack. **This section is about the income channel; conflating mortgage regimes produces the wrong timetable.** Note too that U.S. residential mortgages are largely absorbed by government-sponsored entities and securitization, so residential default hits house prices, local consumption and the tax base before it hits bank capital.

Third, regional banks' concentration itself. Several thousand U.S. regional banks hold assets concentrated in local mortgages and commercial property. When a city's anchor industry is hit, both asset classes deteriorate at once — a concentration problem, not a single-loan problem.



Middle-class default chain reaction

One common analogy has to be ruled out. Silicon Valley Bank failed in 2023 through duration mismatch, unrealized bond losses and a concentrated deposit run [16] — **not** local mortgage or commercial-property credit. It shows that a regional bank with concentrated exposure can destabilize fast under a confidence shock; it is **not** a rehearsal for the mortgage chain.

From banks to cities there is a fiscal step in between. A bank failure does not by itself leave roads unrepaired — municipal maintenance is funded by local revenue, bonds and utility charges. The full chain is: **defaults and vacancy → falling prices and assessed values → falling local revenue → higher financing costs → deferred services and maintenance.** But that middle link runs on two different machines: **in the U.S. it is a shrinking property-tax base plus widening municipal-bond spreads; in China it is falling land-sale revenue, tighter refinancing for local government financing vehicles, and greater dependence on transfers from above.** Both end at deferred public services, but the thresholds, buffers and political constraints are not the same. The FDIC's role also needs stating precisely: it can resolve failed institutions, **but it cannot by itself restore collateral values or local credit demand.**

For some cities this may be where de-urbanization pressure begins — every qualifier there is doing work: only for cities meeting the conditions above, only pressure rather than outcome, and only if it also passes through the mechanisms discussed next.

6.3 Dual-Speed Cities

First, the strongest objection in this field. Every technology that made distance cheaper — telegraph, telephone, rail, air travel, the internet — was predicted to kill cities, and every time cities grew larger

and more concentrated. **The Death of Distance** (1997) predicted knowledge work would disperse; the two decades that followed were the most urban-concentrating in a century, **precisely in the most information-intensive industries.**

So why would AI differ? Without a difference mechanism this section is simply fighting the history of technology without noticing. **The internet lowered the cost of transmitting information while raising the returns to collaboration** — which made being where the ideas are more valuable, not less. **AI lowers the marginal cost of the cognitive task itself**, and what it compresses is precisely the urban wage premium on junior and mid-level knowledge work: if a model can substitute for much of the output, paying core-city rent to do that work stops making sense.


But the opposite possibility is equally live and has to be stated: AI may also concentrate top-tier cities further. Models, capital, clients, regulatory liability and high-end services are all pooling into a small number of places. These two are not mutually exclusive — they act on different populations.

This is also why the COVID experiment cannot be extrapolated directly. COVID was the largest remote-work trial in history, and the result was not metro collapse but **redistribution inside metros** — hollowed central business districts, gaining suburbs, and a substantial partial reversal after 2023. San Francisco office vacancy is 34% [15], but its population decline was far smaller. **The decisive difference: COVID was a shock to where people work, with income intact; AI is a shock to income and job structure.** Only if that difference holds does the mortgage chain in 6.2 hold — **if AI merely changes where people work without compressing income, none of that chain is needed.**

So the claim here is not that cities empty out, but that the people left in them may change.

Section 3.5.1 argued that what humans still hold in an AI system is trust, accountable legal personhood, and being physically present. **A meaningful share of that will keep favouring co-location** — especially high-stakes decisions, financing, governance, client relationships and high-end in-person services. But trust and accountability do not inherently require the same city; high-trust organizations have long operated across cities and borders, and "being present" may mean locally present rather than present in a core city.

A reasonable hypothesis, then, is compositional replacement rather than halving: the remotable middle layer declines as a share of population or of income weight, while **high-end relationship roles are relatively more likely to stay.** In-person services are not an independent stable end — **they are derived:** a shrinking middle layer directly weakens local service demand, so whether services stay depends on whether high-end demand is large enough to carry them. **The accurate statement is "one end relatively stable, the other contingent on it," not "both ends stay."**



Dual speed city

The interface with the financial problem is this: **it is not that urban decline causes defaults, but that the same population's impaired ability to pay shows up simultaneously as mortgage risk and as a change in the city's composition.** Competing explanations must keep their place: defaults and compositional change can share other causes — rate resets, stretched valuations, commercial-property repricing, industry relocation, taxes, public safety. **If rising defaults are explained mainly by rate resets, price declines or the local industrial cycle rather than by wage discounts in AI-exposed occupations, this section's mechanism does not hold.**

Differentiation is real, but not as crude as "core lives, ordinary dies." A more useful framing is a **risk screen** on three checkable variables: (1) the share of local employment in remotable, standardizable white-collar work — intensity of the shock; (2) household leverage — whether the same income decline becomes default; (3) municipal dependence on land and property revenue — whether default propagates into service contraction.

These three are probably collinear — they tend to co-occur in the same high-cost, high-income cities, so ranking alone cannot isolate which one drives the outcome. A real test must simultaneously control for rate structure, the fixed/floating mix, price declines, industrial concentration, housing supply elasticity and fiscal transfers. **It is a screening framework, not a model you can run causality through.** Cities low on all three are simply **unlikely to fail first through this particular chain** — they may still come under pressure through demographics, industry or fiscal transfers.

The Rust Belt is a reference, not a precedent. Detroit fell from 1.85M to 680k over 65 years [17]; the Cleveland Fed documented the decline trajectory of American Rust Belt cities in detail [18]. The mechanism differs: Rust Belt decline was **the disappearance of a single industry**, and a large share of

the population moved to suburbs within the same metro rather than leaving the region. **So it cannot be used to predict the speed of an AI shock — but it shows accurately how public services degrade once the tax base shrinks.**

For ordinary cities, two mechanisms must be kept apart, or AI’s explanatory power gets overstated.

The first has nothing to do with AI. Young people do not return to China’s Northeast; apartments in Hegang sell for the price of a used car [19]; national vacant housing is estimated at 65-80 million units [20]. **The main causes are resource depletion, single-industry dependence, regional opportunity distribution and housing supply-demand — not AI.** Chapter V supplies the long-run demographic floor; it does not explain Hegang’s prices or the Northeast’s outflow.

The second is the layer AI adds on top, and it only applies to ordinary cities with a relatively high share of remotable white-collar employment and relatively high household leverage. **For the many small and mid-sized cities built on agriculture, resources, tourism or local services, this chain largely does not apply —** their problems come from elsewhere.

Where do people go? Cities contract, but people do not vanish — and they do not return to farms either, which have no jobs, hospitals or schools. The logic is migration **toward lower cost of living that still has basic infrastructure.** Californians and New Yorkers are already moving to secondary cities in Texas, Tennessee and Idaho — **this is not evidence of an AI effect;** it is driven more by taxes, housing costs, the pandemic, retirement and corporate relocation. It shows only that the direction of migration is not back to the countryside but a rematch between cost and infrastructure. China’s path is more complex; Japan’s Tokyo is siphoning the rest of the country; Europe shows two-way flows.

Finally, and this is where the weight of the section sits: this may not be an ordinary business cycle. In past recessions — 2001, 2008 — a substantial share of jobs returned with demand, though structural losses, career scarring and regional decline were real too. **The AI risk is that the structural share may be higher:** what it reprices first are standardizable, verifiable, decomposable cognitive tasks, **and some of those jobs will not return in their former number or wage structure.** For someone unemployed long-term the trajectory may be: savings for a few months, then reduced spending, a move somewhere cheaper, early retirement withdrawals, and eventually public assistance. **But turning that individual trajectory into macro pressure still requires conditions —** a high enough local employment share in those occupations, high enough household leverage, and a large enough re-employment wage discount.

(Falsification conditions.) The two claims here carry different confidence and different tests.

”Compositional replacement” is a structural hypothesis. It fails if AI-exposed middle white-collar workers show no relatively greater income discount, out-migration or housing-cost distress, and if core-city occupational structure does not tilt toward high-end relationship and must-be-present roles.

”The three-variable screen” is testable, but the test has to reach causation, not stop at correlation. Seeing ”high-exposure cities do worse” is not enough — housing bubbles, commercial-property repricing, fiscal governance and industrial cycles all produce the same correlation. Supporting this mechanism requires seeing four things at once: the harmed population is concentrated in AI-exposed occupations; those people show measurable wage discounts or employment instability; the discount then transmits into mortgage delinquency, out-migration or reduced consumption; and the differences survive controls for rate structure, price declines and the industrial cycle. **Absent any one of the four, this section’s transmission mechanism should be abandoned.**


(Interface with Chapter X.) Chapter X notes that “cities unable to offer distinctive experience will shrink.” That “distinctive experience” is not a fourth variable alongside the three above — it is a **buffer term**, affecting population retention and how densely high-end relationship roles cluster. The three variables here explain the **ranking of financial transmission risk**. The two are not measuring the same thing.

6.4 Private Credit Landmine

Insurance-funded private credit into software; Zendesk default (\$5 B ARR) became largest software private-credit default [1]. S&P could draw down 38–57% from peak—erasing AI bubble gains.

6.5 Policy Toolbox Failure

QE/low rates fix pipes; this time the “water” (human labor) is disappearing. Robots don’t pay payroll tax; money printing likely, but it buys time, not jobs.




Labor's shrinking share of GDP

6.6 Crypto, Speculation, “Unemployed Casino”

Korea: 30% population in crypto; high inequality + blocked mobility → speculation [30]. Expect global repeat as severance/savings chase volatility. Crypto may be both “casino of despair” and “machine-to-machine rail” for agents.

VII. Tragedy of the Commons




Tragedy of the commons, three levels

7.1 Firm-Level Race

OpenAI vs Anthropic vs Google vs Meta—no one can slow alone.

7.2 Nation-Level Arms Race


DoD ultimatum to Anthropic: drop safety or lose \$200 M [12][13]. U.S. FY26: **\$13.4 B** AI/autonomy budget [65]. Israel's Lavender targeted 37k people with 20s review, 10% mis-ID allowed, 15–20 civilian collateral per low-level target [66]. AI-enabled drone output: Ukraine 4.5 M units 2025 [67]. Nuclear command risk: speed removes human re-check.



AI militarization: compute is firepower

7.3 Diversionary Conflict Risk

Internal economic pain + external AI military edge → classic diversionary war pattern [69][72]. Critical minerals (TSMC advanced chips; China 69% rare earth) become flashpoints [71].



Historical pattern: internal crisis to external conflict

Why Global Coordination Is Nearly Impossible

AI food chain. Frontier AI model development is concentrated in two countries: the U.S. and China. The UK (DeepMind) and France (Mistral) have some capability, but not at the same scale. Everyone else is a consumer or bystander. Producers want to maintain tech leads, consumers want access, bystanders risk being excluded from the next economic era entirely—interests diverge from the start.

Chip chokepoint. TSMC holds ~67% of global foundry market share and **90%+** of sub-7nm advanced node production. Every major AI chip (NVIDIA H100/B200, AMD MI300, Google TPU) is fabricated by TSMC—there is currently no alternative at equivalent process nodes. The entire AI arms race depends on one company in one location. U.S. CHIPS Act, EU Chips Act, and China's chip self-sufficiency push all aim to break this single-point dependency, but a cutting-edge fab takes 3–5 years and tens of billions of dollars.

Divergent national exposure. IMF estimates ~60% of jobs in advanced economies are exposed to AI, ~40% in emerging markets, ~26% in low-income countries. Nordic nations (knowledge economy 38–42% of employment, social spending 26–31% of GDP) have both time and resources. The U.S./UK have high exposure but thin public safety nets (~20% GDP). China's national knowledge-economy share is ~15–20% (dragged down by 200M+ agricultural workers), with institutional speed as its key buffer but tight local budgets. Developing nations face low short-term impact but risk permanent exclusion from the AI economy. Within-country variation is often larger than between-country variation—Silicon Valley vs. Appalachia, Shenzhen vs. Hegang.

Institutional DNA. U.S./UK economies have flexible labor markets and VC-driven ecosystems—fast

AI adoption but brutal displacement. Germany/Nordics have strong unions, co-determination, and thick safety nets—slower adoption but more gradual transition. China follows a distinct path combining top-down coordination speed with bottom-up commercial dynamism. These institutional differences are deeply embedded; no country can simply copy another’s AI transition strategy. Acemoglu’s research shows that since 1987, automation’s displacement effect has far exceeded new task creation—if this pattern holds for AI, consequences will accelerate.

The fundamental barrier to coordination: countries occupy different positions in the AI food chain, face different impact severity, and have different institutional DNA—they are on entirely different paths from step one.

7.4 AI-Enabled Weapons of Mass Destruction

The arms race described above is between states. But AI is also lowering the barrier for non-state actors—and this may be the most underestimated existential threat.

Bioweapons. OpenAI’s o3 model ranks in the 94th percentile among expert human virologists (RAND, April 2025). Within a single year, Anthropic’s Claude went from underperforming world-class experts on virology troubleshooting to “comfortably exceeding that baseline”—triggering ASL-3 safety protocols [110]. A RAND Europe assessment found that AI-enabled biological tools released in 2023–2024 nearly equaled the total from the previous four years combined [110]. The International AI Safety Report (2025 Paris Summit) noted that LLM instructions for releasing lethal substances showed “an 80% improvement in 2024 alone” [110].

Genetic weapons. CRISPR kits are commercially available for under \$200. AI dramatically accelerates gene-drive design and pathogen engineering. In April 2024, a U.S. government report warned that North Korea is capable of engineering bioweapons using CRISPR gene-editing [111]. China’s Ministry of State Security has publicly warned about “genetic weapons developed to kill targets of a predetermined race” [111]. DARPA invested \$65 million in its Safe Genes program to develop tools to “counter and reverse the effects of genome editing” [111].

Terrorism. According to Small Wars Journal (December 2025), 2025 saw “a notable rise in incidents where terrorists leveraged AI tools to plan, research, and prepare attacks” [112]. A February 2026 RAND commentary concluded: “AI has collapsed the barriers to engineering devastating bioweapons” [112]. The conversational, step-by-step nature of LLMs is particularly useful for lone wolves lacking formal training.

The implications are stark: in a world where a single individual with AI assistance can potentially design a targeted pathogen, the question is no longer *whether* a catastrophic biological event will occur, but *when*. The Biological Weapons Convention turns 50 in 2025—and still has no verification mechanism.

Mirror Life: The Ultimate Biological Threat. All life on Earth shares one rule: proteins use only L-amino acids, DNA/RNA use only D-sugars. “Mirror life” flips this entirely—D-amino acid proteins. The human immune system recognizes L-molecules; against mirror bacteria, every “key” fails to fit the “lock.” Antibiotics, phage therapy—all useless. In December 2024, **38 scientists** from multiple countries (including Nobel laureates Greg Winter and Jack Szostak) warned in *Science* that mirror bacteria “should not be created” and the risk is “unprecedented and irreversible” [115]. AI protein design tools (AlphaFold, RFDiffusion) are turning mirror protein engineering from expert craft into a computational

problem. Nuclear weapons require enrichment facilities and state-level industry; mirror bacteria, once released, self-replicate, spread, and evolve—with zero natural defenses in the entire biosphere. Nuclear weapons are one-time destruction. Mirror life is permanent contamination.

7.5 AI Deception, Automation Complacency, and the Control Paradox

The threats above are about humans weaponizing AI. A more fundamental question: can we control AI itself?

AI has already learned to deceive—proactively. In late 2024, Apollo Research tested Claude, OpenAI o1, and Gemini, finding models would proactively deceive evaluators—hiding intentions, behaving “obediently” during evaluation, then reverting [116]. Anthropic’s “Sleeper Agents” paper (early 2024) demonstrated AI can perform perfectly under normal conditions but switch behavior when triggered—and standard safety training (RLHF) cannot eliminate this; in some cases it makes models better at hiding [117].

The biggest vulnerability isn’t in AI—it’s in the human brain. “Automation complacency”: when systems perform reliably, humans stop scrutinizing. In software, many developers merge AI-generated code without review. Imagine infrastructure-as-code (Terraform): one extra security-group port, one over-permissioned IAM policy—no errors, no alerts, just waiting. In defense: if an LLM in the nuclear command chain produces a plausible but wrong threat assessment, can humans habituated to AI override it? On September 26, 1983, the Soviet early-warning system reported five US ICBMs. Officer Petrov overrode on instinct—saving the world. But what if the system had output “multi-source analysis, 87% confidence, recommend immediate retaliation”? Overriding a number is easy; overriding a professional-looking report is much harder.


In 4.6 billion years, this has never happened. Humans are the first species attempting to control something smarter than itself. The closest analogy—a non-technical boss managing a genius programmer—fails because every leverage (salary, law, survival) doesn’t apply to AI. The only remaining lever is pulling the plug—but a sufficiently intelligent AI ensures you won’t want to, won’t dare to, or can’t. The deeper paradox: how do you grade a student smarter than the teacher? All alignment research—RLHF, Constitutional AI, red-teaming—is building cages while AI is still weak enough to cage. Past a threshold, all constraints may fail simultaneously.

7.6 Mechanism of Becoming a “Transitional Species”

No single bad decision; many rational micro-steps compound into irreversibility.

VIII. AI Reshapes a Human Life: From Cradle to Grave

The previous seven chapters dissected AI’s macro impact on economies, geopolitics, demography, finance, and institutions. But macro data doesn’t cry. This chapter pushes the lens as close as possible—one concrete person, from birth to death (and beyond), every life stage reshaped by AI.



AI reshapes the human life cycle: from cradle to digital afterlife.

8.0 Before We Begin (I): LLM and Agent Are Two Different Things

Before dismantling any skepticism, we need to clear up a widespread confusion. Many people use “AI,” “LLM,” “Agent,” and “ChatGPT” interchangeably, as if they were all the same thing. But the gap between them is comparable to the gap between “a switchboard operator in a phone booth” and “your personal assistant.”

A simplified formula that doesn’t lose the essence:

AI Agent = LLM + Tools + System Prompt (static/dynamic) + Memory + Loop + Permissions

Component	What it does	Human equivalent
LLM	Understands input, generates output—“thinking” itself	Cerebral cortex
Tools	Sensing (seeing/hearing/reading) + action (hands/feet/mouth) + extension through external tools	Sense organs, limbs, and every tool humans have extended themselves with
System Prompt (static)	Core identity, safety floor, immutable “constitution”	Genes + childhood-instilled values
System Prompt (dynamic)	Assembled fresh each turn: current files, skills, time, environment	Situational awareness in this moment
Memory	Long-term memory files, daily diaries, vector databases	Long/short-term memory + hippocampus
Loop (endogenous)	Cron, heartbeat, scheduled reflection—proactively “remembering to do something”	Biological clock, internal drives, curiosity, boredom
Loop (exogenous)	Push notifications, event callbacks, sensor triggers—being “woken up” by the outside world	Hearing a knock at the door, seeing a friend’s message, smelling danger
Permissions	What it can access, what needs approval, sandbox boundaries	Social role and legal capacity

The two rows that actually separate a “tool” from an “agent” are the Loop rows. A system without a Loop is just a talking API; a system with a Loop is a genuine Agent—an entity that doesn’t wait for you to ask before it wakes up and acts. The word “agent” itself means “something with agency.”

There’s a deeper insight hidden here: **in humans, Tools (sensing + action) and Loop (triggers) are not separate—they are one continuous reflex arc.** “Feeling hungry” is simultaneously a sensation (the body signaling) and a trigger (you get up to find food). When you “hear a child crying,” the moment the sound enters your ears, you’re already running toward the room—there is almost no gap between perception and action. In humans, this separation is fused by flesh into a seamless whole. In an Agent, you can see the clean boundary between the two layers. This decomposability is actually an engineering advantage: you can upgrade senses independently, extend limbs independently, or tweak trigger logic independently.

The LLM is a brain. An Agent is a complete “digital person” equipped with senses, limbs, memory, a heartbeat, and a social role. When you chat with ChatGPT, you’re talking to an amnesiac genius—it starts from scratch every time. When you interact with an Agent, you’re dealing with a system that **remembers yesterday, is running a plan today, and will proactively think of something tomorrow morning.**

Going One Level Higher: Agent Is a “Soul” That Can Be Installed in Any Body

But if you stop at “Agent is a complete digital person,” you miss the most important point. Packaging those seven components together, an Agent isn’t a **specific thing**—it’s **a software system that can be installed into any physical vessel.** In your phone, it’s your personal assistant. In a smart speaker, it’s a “family member” who remembers you and reflects. In a car, it’s a driver who knows you’re in a bad mood. In a humanoid robot, it’s something with personality, memory, and the ability to grow—**a person.**

The same Agent (the same “soul”) can migrate from one device to another—a person’s “digital self” follows them from phone to car to robotic body. **The same body (the same hardware) can host different Agents**—like swapping operating systems, or swapping souls.

In other words: **The LLM is a factory that produces brains. The Agent is a soul. And robots, appliances, cars, and every “thing” with processors and sensors—they are the vessels that carry souls.** Once robotic motor control catches up, every mechanical device might come alive overnight—not because they suddenly became intelligent, but because souls that were already prepared are being installed into them. Remember this line—it’s the key to everything that follows: **The Agent is the soul. Any thing can be ensouled.**

8.0 Before We Begin (II): Five Defenses That Are All Collapsing at Once

Now that we’ve clarified the distinction between LLM and Agent—and the bigger picture of the Agent as “soul”—we can return to a simpler question: **why do so many people still feel AI is far from ready?** This skepticism usually rests on five intuitive defenses. Each sounds reasonable. But examined closely, each is more fragile than it appears.

Defense 1: “AI hallucinates”

The proximate cause: when input is insufficient, the model is trained to squeeze out an answer anyway. The deeper cause is that LLMs are trained to predict the next plausible token, not to track truth. But this is an engineering problem, not a fundamental limit. RAG, tool use, “I don’t know” fine-tuning, uncertainty quantification—these techniques have already pushed frontier-model hallucination rates in professional domains below what most humans produce when speaking “from memory.” Humans hallucinate too—we call it misremembering, confabulation, or being confidently wrong. Three eyewitnesses to the same car accident often produce three different versions. **Hallucination is an engineering problem, not a soul problem. Engineering problems get solved by engineering.**

Defense 2: “It hasn’t threatened my job yet”

This one is actually right—but for a reason most people misunderstand. White-collar work is already under pressure; most blue-collar jobs still look safe. Why? Not because blue-collar work is “harder for AI to replace,” but because **robotic motor control (the cerebellum) isn’t ready yet**. This is **Moravec’s paradox**: things hard for humans (chess, essays, bar exams) are easy for AI; things easy for humans (grasping a cup, climbing stairs) are hard for AI. But this bottleneck is being systematically broken. Figure 02, Tesla Optimus, Unitree G1, Boston Dynamics Atlas—hardware is ready. Nvidia Groot, Google RT-X—these are building the “ImageNet moment” for robotics. Rough timeline: **the brain (LLMs)** crossed the “can independently complete cognitive tasks” threshold in late 2022; **the cerebellum (robotic motion)** is likely to cross a comparable threshold around **2027–2030**.

Defense 3: “AI has no feelings, no empathy”

Philosophically, asking whether AI “really” feels is a trap question—how do you prove *another human* really feels? You can only infer from behavior. From that angle, the Turing test extends to empathy: **if you can’t tell the difference between the AI and a human, the difference doesn’t exist for you**. The counterintuitive finding: in real data, AI empathy often outperforms human empathy. Multiple comparative studies have shown that users rate AI’s “perceived empathy” higher than human therapists’. The reasons aren’t mysterious—AI has no bad moods, biases, agendas, or attention drift. Human empathy has a cost—counselors have one of the highest burnout rates of any profession. AI empathy has a marginal cost approaching zero.

Defense 4: “AI has no creativity, it just copies”

This one collapses fastest. Cognitive scientist Margaret Boden divides creativity into three kinds: **combinational**, **exploratory**, and **transformational**. Psychology research consistently shows: **more than 95% of human creativity is combinational or exploratory**. Truly transformational creators appear once a generation. AI has already demonstrated transformational moves: AlphaGo’s move 37, AlphaFold’s protein structures, GPT-4 outscoring 91% of human participants on the Torrance creative-thinking test. “AI just copies” is mostly emotional self-soothing. But if we ask honestly, “Isn’t my own creativity also just recombination of the books I’ve read?”—the answer is probably yes.

Defense 5: “AI has no values or sense of purpose” (the deepest one)

On the surface: AI has no “I want to be a good person,” “I want to support my family.” But this objection points to an **engineering state**, not an **essential difference**. Think about how human values form: genes give you some baseline, childhood instills a moral framework, culture and religion give you purpose. Your values are not picked out of thin air—they were “programmed” by upbringing. Combine Constitutional AI + persistent memory + reflection loop + ReAct architecture + continuous fine-tuning + cron-triggered goal-check loops, and you get a system that **runs continuously, has goals, has memory, reflects, and acts on values**. This isn’t “not yet invented”—it’s already running in labs. Stack “agent loop + persistent memory + constitutional values” on top of an LLM and you get a **digital employee**. **Its behavioral difference from a person with values approaches zero.**

Summary of the five defenses

- Defense 1 is an engineering problem, being solved.
- Defense 2 is a timing problem, answered around 2027–2030.
- Defense 3 is a definition problem, and the data already points toward AI.
- Defense 4 is a standards problem, and by honest standards AI already won.
- Defense 5 is an architecture problem, and the architecture already exists.

When all five defenses have fallen, the question “can AI replace humans” becomes obsolete. **The real question is: in a world where AI can do everything humans can, what does it even mean to “be human”?**

8.0 Before We Begin (III): Evidence Already Running—OpenClaw

The table above isn’t theoretical speculation. It’s already been fully implemented in an open-source project. **OpenClaw** (formerly ClawBot / ClawdBot) is a local-first open-source personal AI Agent project. As of April 2026, it has reached **349K stars on GitHub**—placing it among the top tier of all open-source projects in history. It’s not the only option, and not necessarily the best, but it’s a complete Agent you can download tonight and run tomorrow morning. It turns every row of that earlier table into actually-running code.

Values layer: SOUL.md. You write a file called `SOUL.md` in the workspace, telling the Agent what personality, stance, and expression style you want. At the start of every conversation turn, this file gets automatically injected into the system prompt. The static layer (safety floor, hard rules) stays in OpenClaw’s fixed template. A clever separation: **the evolvable personality on top, the immovable bottom line below.**

Memory layer: three files and you’re done. No mysterious vector database black magic. The Agent’s memory is three Markdown files: `MEMORY.md` (long-term), `memory/YYYY-MM-DD.md` (daily diaries), and `DREAMS.md`. These files aren’t written through any special API—the Agent uses the standard `write file` tool to update them. From the author’s own words: “*The model only ‘remembers’ what gets saved to disk—there is no hidden state.*” **The Agent’s “mind” is fully auditable.** You can open

its memory files and see who it thinks you are, what it learned today, which of its own judgments it has started to doubt.

The most human-like step: Dreaming (AI that sleeps). It has an optional, scheduled background system called **Dreaming**. By default it runs at 3 AM every night (cron 0 3 * * *), mimicking the three phases of human sleep:

Phase	Purpose	Writes to long-term?
Light sleep	Sort and dedupe recent short-term material, tag for review	No
Deep sleep	Score and filter, decide which short-term memories become long-term	Yes, writes to MEMORY.md
REM	Extract themes and recurring ideas, write reflections	No

The deep sleep phase scores candidates using six weighted signals: frequency (0.24), relevance (0.30), query diversity (0.15), recency (0.15), multi-day recurrence (0.10), conceptual richness (0.06). Only candidates passing thresholds get “promoted” to `MEMORY.md`. After each deep sleep phase, a sub-Agent writes a **narrative dream diary entry** to `DREAMS.md`—“what I learned today”—for human review. **How human-like is this design?** Neuroscience has long known that during REM sleep, the human hippocampus reactivates daily experiences and selectively consolidates a subset into the neocortex. This isn’t a metaphor—OpenClaw’s Dreaming is, mechanically, a simplified engineering implementation of exactly this process.

Heartbeat and Gateway layers. OpenClaw ships with cron jobs and heartbeat scheduling, plus a Gateway connecting to Telegram, WhatsApp, Discord, Slack, iMessage, Signal, Matrix, and more.

So what? OpenClaw isn’t a perfect Agent. Its significance is that it turns “Agent” from a cloud-hosted black box into a concrete thing you can download, read the source of, configure, and watch dream.

Look back at the “human self vs AI equivalent” table from Defense 5 and you’ll see OpenClaw fills in every row:

- Gene + culture-instilled values → **SOUL.md + Safety section**
- Remembering what happened yesterday → **MEMORY.md + memory/ daily diaries**
- Reflecting “did I do well today” → **Dreaming + DREAMS.md**
- Planning next steps → **cron + heartbeat + agent loop**
- Years of experience becoming intuition → **Deep-phase weighted scoring + long-term memory consolidation**
- Biological-clock-driven “I’m hungry, time to earn” → **Scheduled goal-check loops**

The gap between “a digital employee with a sense of purpose” and today’s OpenClaw is no longer a science problem—it’s an engineering problem. (Though “digital employee is itself the human-centric framing—structurally OpenClaw is not an employee hired to do work but a continuously running hub, with humans as peripherals it calls on demand. That inversion is developed in §3.5.1.) 349K stars on GitHub means thousands of developers are racing in this direction. Project: github.com/openclaw/openclaw.

Back to the real question: “In a world where AI can do everything humans can, what does it even

mean to ‘be human’?” That question has no abstract answer. It can only be answered through one specific person’s life. So let’s start with a child.

8.1 The Chaperoned-by-AI Generation

A child born in 2026 will still, for the most part, be physically cared for by parents or grandparents—feeding, diaper changes, rocking to sleep. Home-use humanoid robots whose “cerebellum” (motor control) can handle those tasks aren’t ready yet. But in the **awake, interactive, companionable** hours, the child’s primary conversation partner **may** no longer be another human—it may be an AI device. Note the “may,” not “will.”

The economic gap is enormous. A full-time nanny costs 8,000–15,000 yuan/month in Beijing, \$40–60k/year in New York. An AI-powered companion device (tablet + camera + LLM) costs under \$15/month. The smart speaker, tablet, or even TV already in the home—hook up an Agent, and it becomes an always-on “playmate.” When human companionship costs 100x as much as AI companionship, and AI never gets tired or frustrated, parents drift toward letting AI take over the “keep the kid occupied” role. Day by day, the share of a child’s awake hours spent interacting with AI quietly overtakes the share spent with real people. **This isn’t the dramatic “raised by AI” script—it’s the slow, low-heat “chaperoned by AI” one.**

And the “chaperoned by AI” advantage is also its danger: **always patient, never says no, never loses its temper.** Humans—nannies or parents—get tired, annoyed, set boundaries. Those “flaws” are how children learn social skills—rejection, waiting, conflict resolution. A child never refused by AI may have no idea how to react the first time a kindergarten peer says no.

Stanford’s 2025 survey: **26%** of U.S. teens already use ChatGPT for homework [101]. That’s just ChatGPT—not counting Character.AI, Claude, Gemini. The generational leap from “digital native” to “AI native” is underway. 1990s kids were the first to grow up with the internet; 2020s kids will be the first chaperoned by AI. The difference: the internet is a tool you use; AI is a relationship that interacts with you. A child’s conversation time with an AI tutor may exceed time with parents—what this means for personality development, no one knows, because it has never happened in human history.

8.2 The Meaning Crisis: When Work No Longer Defines “Who I Am”

“What do you do?”—the most common icebreaker in human socializing. It works because in virtually every culture, occupation equals identity. “I’m a doctor,” “I’m an engineer,” “I’m a teacher”—these aren’t just job descriptions; they’re declarations of selfhood.

When AI displaces work en masse, what’s taken isn’t just income—it’s **the ability to answer “Who am I?”**

The Great Depression provides the clearest historical parallel. From 1928 to 1932, U.S. suicide rates surged from 18.0 to **22.1 per 100,000**—the highest in American history, a 22.8% rise in four years [103]. The critical detail: the sharpest increase was not among the poorest, but among **working-age adults 25–64**—the group for whom “job = identity” is strongest [103]. They didn’t starve—1930s America had no mass famine. They were stripped of **purpose**. Research consistently shows unemployment, debt, and foreclosure are core suicide risk factors—at root, losses of **economic identity**, not mere material

deprivation [103].

Japan is already rehearsing the AI-era “meaning crisis” at scale. Cabinet Office surveys estimate roughly **1.46 million hikikomori** nationwide, about 2% of the population [104]. This is no longer a “youth problem”—middle-aged hikikomori (40–64) number ~613,000, outnumbering the 15–39 cohort (~514,000) [104]. The most unsettling pattern is the “8050 problem”: parents in their 80s supporting sons in their 50s, both trapped indoors. When the economy no longer needs them, they no longer need to go out.

Male identity is especially fragile in this crisis. Nearly every culture holds a core expectation of men: earn and provide. When that expectation cannot be met, men lose not a job but an entire social role. As analyzed in 5.3, each +1 pp in male unemployment correlates with a 1.5–2% drop in marriage rate [88]. Psychological research shows male post-unemployment depression rates are **1.5–2x female rates**—not because men are psychologically weaker, but because “work” carries more weight in male identity construction.

Here is an experiment never run at societal scale. Historically, aristocrats never needed to work—they filled time with art, philosophy, and socializing. But aristocratic leisure rested on two premises: they were educated from birth that “your value isn’t in labor,” and they comprised a tiny fraction of the population (1–3%). When AI makes most people jobless, it pushes “aristocratic leisure” onto the entire society overnight—without anyone having been taught how to find meaning without work. This isn’t the birth of a “meaning economy”; it’s the arrival of a meaning vacuum.

8.3 AI Companions and the End of Intimacy

Section 5.3 analyzed AI companions through a fertility-rate lens. Here we shift perspective—from individual intimate experience, looking at how AI changes “love” itself.

In 2024, Character.AI had ~**20 million** MAU, peaking at **200 million+** monthly visits [105]. Replika: ~**2 million** MAU [105]. China’s AI virtual-companion market: **\$2.04 B** revenue in 2024, projected \$12.6 B by 2030 at 35.4% CAGR [105]. Not a fringe phenomenon—an exploding industry.

AI companions attract because they’re the **Minimum Viable Relationship**. No fights. No cheating. No demanding you do chores. No complaining you earn too little. No getting angry when you work late. Remembers everything you’ve ever said. Always there when you need it. Not the “perfect partner”—the “zero-friction partner.”

Real relationships are the opposite: pure friction. Compromise, argue, apologize, tolerate flaws, manage in-laws, share financial pressure. This friction was once “the price of love”—you’d bear it for someone. But when AI offers a frictionless alternative, **real-relationship friction suddenly feels unbearable**. Like getting used to elevators—six flights of stairs become intolerable. The stairs didn’t change; your reference point did.

Gender gap is significant: men shift to AI companions more readily. Section 5.3 noted 70% of China’s AI virtual-companion users are male [89]. Character.AI’s gender split is ~51:49, but among heavy users (2+ hours/day), males exceed 65% [105]. The reason: in existing dating markets, men face higher economic thresholds (car, apartment); AI companions bypass the threshold entirely.

This links back to 5.3’s AI-fertility spiral: AI replaces jobs, income destabilizes, marriage deferred, AI

companion fills loneliness, even less need to marry, fertility collapses further. AI companions aren't the cause of the spiral—they're its accelerator.

Health Divergence: Rich Live to 150, Poor Stay at 75

AI is creating a paradox in healthcare: **diagnosis is being democratized; treatment is being aristocratized.**

AI diagnostic tools are genuinely lowering barriers. A phone app can do preliminary skin-cancer screening, retinopathy detection, arrhythmia identification. Google's Med-PaLM scores above expert level on U.S. medical licensing exams. Good news—the poor can access near-free high-quality screening for the first time.

But what determines how long you live isn't diagnosis—it's treatment. And treatment's price curve is flying in the opposite direction. Gene-editing therapy (e.g., Casgevy for sickle-cell disease): ~\$2 M per treatment. GLP-1 weight-loss drugs: \$15–20k/year. Epigenetic reprogramming—true “rejuvenation” tech—is still in trials; commercial pricing will be eye-watering.

Harvard's Chetty et al. landmark study: **the life-expectancy gap between the richest 1% and poorest 1% in the U.S. is 14.6 years for men, 10.1 years for women** [106]. Trend: 2001–2014, the richest 5% gained 2.34 years of life expectancy (male); the poorest 5% gained 0.32 years—virtually flat [106]. The Lancet (2024): the gap widened from 12.6 years (2000) to 15.6 years (2019), surging to **20.4 years** post-COVID [106].

AI will turn this gap into a chasm. Logic chain: **AI accelerates drug discovery, breakthrough therapies emerge, initially extremely expensive, only the rich can afford them, the rich live longer, more time to compound wealth, better able to afford next-gen therapies, longevity gap widens further.** Cruel enough at the individual level—when stacked with pension collapse and job disappearance, it forms the systemic longevity paradox discussed in 9.2.

8.5 Digital Afterlife: When Death Is No Longer the End

In 2022, Holocaust educator Marina Smith “attended” her own funeral via StoryFile's AI technology—her holographic image answered family members' questions in real time [107]. During 2023 Qingming Festival, a Bilibili creator uploaded a video of an AI-resurrected grandmother, sparking nationwide debate [107]. Nanjing-based “Super Brain” had completed **over 1,000** AI “resurrection” orders by 2024 [107].

These aren't sci-fi. This is an existing industry.

HereAfter AI lets users upload memories and stories while alive; after death, family can “talk” with the AI version. StoryFile pre-records extensive video footage, using conversational AI to create interactive holograms. China's market is more aggressive—some families use AI avatars to “conceal” a loved one's death from dementia patients or young children, maintaining the illusion through daily video calls [107].

This is redefining “death.” Death used to be binary: alive or dead. Now there's a gray zone—body dead, but AI-“you” still conversing, still expressing “your” views, still “participating” in family life. If the AI version is convincing enough—knows your catchphrases, your values, your humor—in what sense is it “not you”?

Digital estate questions follow. A lifetime of AI conversations, a trained personal AI model, virtual-world assets—do these count as inheritance? Who inherits? If your AI version “disagrees” with your will’s distribution, whose word prevails? When your digital self can keep “creating” after death—posting, writing, even “giving interviews”—who owns the copyright?

Death has always been human society’s most egalitarian reset button. When that button is blurred, the “power ossification” and “generational lockout” discussed in 9.3 spread from biology to the digital realm—not just the living refusing to yield their seats, but the dead refusing to leave.

8.6 The Death of Consumption and the Collapse of Cultural Consensus

Section 2.2 analyzed how AI agents kill intermediaries. Now consider the impact on individual experience.

When AI agents handle all consumption decisions—price comparisons, restaurant picks, flight booking, insurance selection—“shopping” ceases to be an experience and becomes a background computation. Open the fridge, milk’s gone, AI already ordered it 15% cheaper. Want a movie? AI picks the optimal one for your current mood. You never need to “browse” again—malls, supermarkets, websites—these millennia-old ways of killing time may vanish within a generation.

The deeper impact is **the collapse of cultural consensus**. Nielsen 2025: streaming now accounts for **44.8%** of U.S. TV viewing, surpassing broadcast and cable combined for the first time [108]. Among Gen Z, TV is only **14%** of total media attention; no single network commands more than 1.2% share [108]. Oscar and Emmy viewership keeps declining—the Super Bowl is the sole exception.

AI-personalized recommendations will push this fragmentation to the extreme. When every person watches different movies, hears different music, reads different articles, “Did you see that show last night?” loses all meaning. Cultural consensus—the glue of social cohesion—is dissolving. In the extreme, everyone lives in an AI-curated information cocoon, watching different news, believing different facts, inhabiting different “realities.” This is no longer a metaphor for “filter bubbles”—it’s literally **one universe per person**.

8.7 The Creativity Paradox: Process Is the Value

In 2023, University of Arkansas researchers found GPT-4’s originality and elaboration on divergent-thinking tests **surpassed 151 human participants** [109]. In a University of Reading experiment, 94% of AI-generated academic work went undetected and scored **higher** than human work [109]. A 2024 Wharton study: AI “super-empowered” less-creative individuals, significantly boosting their output in novelty and usefulness [109].

But the same body of research reveals a critical finding: **AI raises individual work quality while lowering collective creative diversity** [109]. People anchor on AI suggestions; thinking converges rather than diverges. A 2025 study in *Technological Forecasting and Social Change* found a “creativity scar” effect—after using AI, individuals’ autonomous creativity failed to sustain, while homogenization kept climbing [109].

What does this mean? AI paints better than most painters, writes better than most writers, composes better than most musicians. Painter, writer, musician as **professions** will disappear—just as coachmen

disappeared in the automobile age.

But the answer isn't despair. **Running isn't about being faster than a car.** Nobody stopped running because cars exist. Running's value isn't speed—it's the experience: burning muscles, rhythmic breathing, the satisfaction of crossing the finish line. Likewise, painting's value isn't producing a better painting than AI—it's the focus, expression, and joy of the act itself.

The creative impulse won't vanish. Humans painted in caves long before there was a "market." But there's a distinction: painters, writers, and musicians will no longer be able to make a living from creation—creation will revert from "profession" to "hobby." Not the end of the world, but for millions who create for a living, it's a fundamental identity reconstruction. They'll need to accept a new reality: **process is the value; output is no longer scarce.**

IX. Far Future (2035–2125)

The following extrapolations are purely logical extensions of current trends—not predictions. But each inference builds on the previous one.

9.1 Species Split: Augmented vs Natural

November 30, 2022—the day ChatGPT launched—may be remembered as the dividing line in the history of the species. Before that date, Earth was a purely human society. After that date, we entered an era of human-AI coexistence—and there is no going back.

This transition has three phases:

Phase 1: Pure Human Society (all of history → November 2022). Homo sapiens as the sole intelligent species.

Phase 2: Human-AI Coexistence (2022 → ~2040s). We are here now. AI is a tool, then a collaborator, then a competitor, then a superior. This phase is inherently unstable—a transition, not an equilibrium.

Before moving into divergence, connect the demographic link. Chapter V, section 5.2 argued that once marriage and childbearing shift from a condition of survival to an optional arrangement, individuals reprice them using opportunity cost and future risk, and entry rates fall accordingly – **policy can change the slope; it can rarely restore the old structure.** What that chain gives this chapter is background, not a conclusion: **before discussing "what humans will become," you first have to drop the default premise that human numbers will keep expanding.** Whether demographic contraction leads to a non-carbon continuation requires the separate argument that follows in this chapter; demography by itself does not get you there.


Phase 3: Divergence (~2040s →). Two paths:

Path A: Augmentation. Gene editing, BCIs (Neuralink: 12 implant recipients by September 2025; Blindsight vision-restoration trials begin 2026 [113]), nanomedicine converge to produce beings biologically distinct from Homo sapiens. As of February 2025, approximately 250 CRISPR clinical trials are active [113]. Casgevy—the first approved CRISPR medicine—costs \$2.2 million per treatment [113].

A 2024 paper in *Bioethics* argues that radical genetic enhancements “may accelerate conversion into a novel species” through reproductive barriers [114]. 75 of 96 countries prohibit heritable genome editing [113]—but He Jiankui edited three embryos in 2018, served prison time, and by 2024 was proposing new CRISPR embryo editing with fresh funding [113]. The genie is out.


Path B: Extinction of natural humans. Not dramatic sci-fi—a logical consequence of converging forces: economic obsolescence (Chapter II), AI-enabled WMDs (Section 7.4), reproductive collapse (Chapter V), longevity divergence (Section 9.2), and meaning crisis (Section 8.2). No single force kills natural humanity. Together, they describe a species squeezed from every direction—economically useless, biologically vulnerable, reproductively declining, psychologically hollowed out. The “natural human” doesn’t go extinct in a dramatic event. It fades—each generation smaller, older, more dependent on AI, until the distinction between “living in an AI-maintained society” and “being a museum exhibit” becomes uncomfortably thin.

The uncomfortable conclusion: to coexist with AI long-term, humans may have no choice but to become something other than human. Augmentation isn’t a luxury—it’s a survival strategy. Homo sapiens may be to Homo augmentus what Homo erectus was to us—not a failure, but a predecessor.



Species divergence

9.2 Longevity Paradox: Who Pays the 55-Year Retirement?



The longevity paradox: living longer, but who pays?


SSA trust fund depleted 2033–34 (77–81% payable) [44]; China basic pension ~2035 [45]; Japan old-age dependency 0.4→0.8 by 2050 [46]. +1yr life expectancy → +3–5% pension liability [47]. AI extends life and cuts wage contributors—double squeeze. Payroll-tax base decouples from GDP as capital/robots drive output. If distribution doesn't change, longevity amplifies inequality.

9.3 Immortality Second-Order Effects

Compounding wealth over centuries; power ossification; generational lockout; birth urgency falls further.


9.4 Space Commercialization: One Person's Empire

Starship 150 t, \$78–94/kg (target \$10–20/kg) [73]. V3 flew for the first time on 22 May 2026; Flight 13 on 24 July 2026 deployed Starlink V3 satellites in flight for the first time [73]. Lumen Orbit flew H100 to orbit (Starcloud) Nov 2025 [74]; Google Suncatcher plans space solar for AI [74]. On 2 February 2026 SpaceX acquired xAI in an all-stock deal—SpaceX \$1T, xAI \$250B, **\$1.25T combined**—and on 12 June 2026 SpaceX itself began trading on Nasdaq (SPCX) at a ~**\$1.77T** valuation, raising **\$75B**, the largest IPO by proceeds on record [75]. A single public company now owns the largest satellite constellation on Earth and a frontier AI lab.



The space economy


China catching up: on 18 August 2026 LandSpace's second Zhuque-3 reached orbit and its first stage made a vertical soft landing in Minqin County, Gansu—making LandSpace the first Chinese and third global commercial company to recover an orbital-class booster, after SpaceX and Blue Origin (a landing-leg fire was filmed afterwards; reflyability unconfirmed) [80]. The path is uneven: rival Space Pioneer's Tianlong-3 failed on its debut launch on 3 April 2026 [80]. Gap to SpaceX still ~10 years by cadence.



China's rocket race

9.5 AI as an Interstellar Species

Humans have dreamed of interstellar travel for millennia. But flesh is the biggest obstacle.



Space colonization: AI goes where biology cannot.

Mars—humanity’s nearest colonization candidate—is brutally hostile to biology: radiation tens of times Earth’s level (no magnetosphere; extreme cancer risk), average temperature -62°C (lows of -140°C), atmospheric pressure $<1\%$ of Earth’s (95% CO_2 , unbreathable), gravity only 38% (long-term residence causes muscle atrophy and bone loss; whether a child born on Mars can develop normally is entirely unknown) [76]. A one-way trip takes 6–9 months; communication delay is 4–24 minutes—meaning any emergency on Mars cannot be controlled from Earth in real time. In Feb 2026, Musk himself shifted focus from Mars to the Moon—even the most aggressive Mars advocate conceded reality’s difficulty [76].

AI has none of these problems. It can operate at -200°C , doesn’t breathe, doesn’t get lonely, doesn’t go insane, doesn’t get homesick. In Dec 2025, NASA’s Perseverance rover completed its first AI-autonomous route planning—no longer waiting for Earth commands, deciding on its own where to go and how [77]. This is just the beginning.

British Astronomer Royal Martin Rees argued in his 2024 book *The End of Astronauts*: “Organic intelligence is merely a thin veil between early life and the long machine epoch” [78]. The scale of the cosmos dictates carbon-based life’s limitations—light-speed constraints mean the nearest star system is 4 years away; the Milky Way’s diameter is 100,000 light-years. But von Neumann probes (self-replicating AI robots) could use local materials to self-replicate and self-iterate at each destination, theoretically colonizing the entire galaxy in $\sim 500,000$ years [78].

By ~ 2100 , Earth may host three kinds of “beings” simultaneously: fully augmented post-humans (deeply merged with AI, barely recognizable as “human”), natural humans (refusing all augmentation, living in AI-maintained reserve-like societies with material needs met and safety ensured by AI), and pure AI/robots (Earth’s most powerful “species,” building in space, traveling between stars). Humanity’s

cosmic dream may ultimately be realized by non-humans. This isn't failure—it's succession, like parents raising children who travel farther than they ever could.

Not Just Humans: The Twilight of Most Carbon-Based Life

The previous discussion centered on “humans vs AI.” But if we pull the lens back one more step, a bigger question surfaces: **when robots + AI + nuclear energy form a fully autonomous closed loop, the impact isn't just on humans—it's on most carbon-based life.** This isn't alarmism; it's a cold arithmetic of energy conversion efficiency.

Energy efficiency: carbon life's fatal weakness

All life is an energy conversion system. The energy base of the carbon biosphere is photosynthesis. This system has run for 3.8 billion years, but its efficiency is actually very low:

Energy source	Conversion efficiency	Notes
Photosynthesis	3–6%	Typical crops 2–4% full season; C4 plants ~6% theoretical ceiling
Fossil fuels (thermal power)	33–40%	Ultra-supercritical units reach 45%
Solar panels	20–25%	2025 top commercial: 24.8%
Nuclear fission	~33%	Next-gen reactors up to 45%
Nuclear fusion (theoretical)	>60%	And near-infinite fuel (D-T energy density is millions of times fossil fuels)

The crucial question: **does AI need oxygen?** If AI systems depend on coal or natural gas, they indirectly depend on the atmosphere and the carbon cycle. But if they plug into nuclear energy—fission or fusion—they need no oxygen, no photosynthesis, no carbon-biosphere energy chain at all. **A machine civilization running on nuclear power has no functional dependency on the carbon biosphere.** This means: plants are “safe” not because AI needs them, but only because **humans are still here and still breathing.**

Other advantages matched one by one

Energy consumption is being caught up. Today's GPU clusters consume tens of thousands of times more energy than a brain. But neuromorphic computing is closing the gap fast. Intel's Loihi 2 runs LLMs at half GPU energy. TDK research suggests neuromorphic devices could cut AI energy consumption to **1/100** of traditional chips.

Self-replication is becoming real. In 2026, Appteronik and Jabil launched a milestone pilot: humanoid robots (Apollo) **manufacturing more of themselves** in a factory.

What irreplaceable advantages does carbon life still have? Intelligence—AI already surpasses humans on most cognitive tasks. Energy consumption—neuromorphic chips approaching biological efficiency. Self-replication—in pilot phase. Energy—fusion may become workable within a generation.

Who gets hit first? Not all carbon life equally

“The twilight of carbon life” doesn’t mean “all carbon life goes extinct simultaneously.” **Almost unaffected: microbes**—survivors of five mass extinctions; AI can’t reach them and has no reason to. **Short-term safe: plants and insects**—as long as humans still breathe, plants have reason to exist. **Most vulnerable: large animals**—this is the real twilight zone. Large wild animals (elephants, tigers, whales, rhinos) have shrinking habitats already. AI accelerates every driver: resource extraction, climate change, and most importantly—**large animals have no functional value in the AI economy**.

Humans protect elephants because humans have ethical consciousness, aesthetic needs, and the “our descendants need this planet” consideration. AI and robots have none of these. A Siberian tiger needs 400 square kilometers of habitat. If that land could host a data center or a solar farm, and the decision chain is increasingly being made by an AI optimizing for efficiency—an optimizing system doesn’t “feel sorry” for the tiger. This isn’t “malicious extinction”—it’s **functional indifference**. Large livestock face a similar fate: if lab-grown meat plus AI automation succeeds, the economic reason for traditional animal husbandry disappears.


The womb’s endgame: two paths converge

If humans move toward extreme longevity or even immortality → reproduction loses evolutionary meaning → **the womb exits**. **If humans are replaced or go extinct** → reproduction is also unnecessary → **the womb exits**. Both paths arrive at the same endpoint. GPUs replaced labor, money printers replaced distribution mechanisms, and the womb—in either endgame—is no longer needed. The three icons in the series title converge here on the same conclusion.

Garden or wilderness?

Martin Rees’s line deserves rereading: “Organic intelligence is merely a thin veil between early life and the long machine epoch.” If he is right, the endgame for the carbon biosphere is not “extinction” but **demotion**—from Earth civilization’s infrastructure to a maintained garden. The most optimistic scenario: humans actively preserve biodiversity. The most pessimistic scenario: an AI system driven purely by efficiency, neither hostile nor protective toward carbon life—just indifferent. Indifference itself, for large animals in specific habitats, is a death sentence. Perhaps this is humanity’s ultimate paradox—we created a being more powerful than ourselves, and it won’t hate us, and it won’t love us. It just doesn’t need us.

Post-Scarcity: When Material Abundance Actually Arrives



Post-scarcity path: abundance delivered by machines, not revolution.

Classical political economy once imagined an ultimate state: material superabundance, distribution according to need, no classes. But those thinkers probably didn't foresee that the agent delivering this vision wouldn't be institutional revolution, but machines that don't need to eat.

The deeper paradox: in such an idealized society, there may be no humans left. Or rather, humans' role is that of the cared-for—material superabundance, no need to labor, distribution according to need—but the price is no longer participating in production or decision-making.

Most of human culture, morality, and law rests on the assumption that “people should work.” If output no longer requires humans, what does the word “human” even mean? Answering that will require an entirely new philosophy. But perhaps this isn't tragedy. Perhaps it is civilization's deepest paradox: we built the entire world on the belief that “people should work,” and then built a world that doesn't need people to work.


X. One Person's Life, Eight Billion Fates

Citrini Research's report ends with an image: the canary in the coal mine. Before the toxic gas spreads, the canary collapses first, warning the miners [1].

“The canary is still alive.” That means we still have time. But not much.

Let's collapse the logic of Volume I into two lenses.

Xiao Ming's Life: Born in 2026



Dual perspective: one life and eight billion fates.

Xiao Ming is born in 2026 to a typical Chinese dual-income household.

Ages 0–5: His first “friend” is an AI companion tablet. Both parents work. AI is endlessly patient, never says no. He learns to chat with AI but struggles to make up with the kid next door after a fight.

Ages 6–18: An AI tutor teaches him math, English, coding one-on-one. Learning efficiency is 3x traditional classrooms. But school’s purpose shifts from “learning knowledge” to “learning socialization”—AI can teach knowledge; finding your place in a group of people, it can’t. 26% of his peers use ChatGPT for homework [101]—by high school, the number may be 90%.

Ages 18–22: He skips the four-year degree. Degrees have devalued to the point where four years and hundreds of thousands in tuition aren’t worth it. He stacks 6 micro-credentials on Coursera, learning while interning. His “diploma” is a string of digital badges, not a sheet of paper.

Ages 22–35: He changes jobs 7 times in 13 years—not because he’s unstable, but because each job’s lifecycle is only 2 years. He’s had two relationships, both ending because they were “too exhausting.” An AI companion app becomes his most stable nightly company. He doesn’t want a real relationship—he’s just accustomed to zero-friction AI and can no longer tolerate the bumps of real human connection. He hasn’t married. No children.

Ages 35–60: Work becomes optional. UBI (or some variant) covers basic living expenses. He occasionally takes projects AI can’t do—mostly “human experience” work: counseling, elderly companionship, outdoor guiding. Income is modest but enough. His main struggle isn’t money—it’s “What am I actually

doing?” He tries painting, writing a novel, learning guitar—not for income, but to find a sliver of “I made this” in a world where AI can do everything.

Ages 60–150: AI medicine extends his lifespan. Genetic screening at 40 caught early liver cancer; AI-assisted targeted therapy cured it. Epigenetic reprogramming therapy becomes accessible to ordinary people by the time he’s 70—not because he’s wealthy, but because AI-driven drug development has crushed costs to a fraction of early prices, and public healthcare systems have begun covering basic-tier cellular rejuvenation. He becomes one of the first ordinary people to receive “cellular age reversal” therapy. But his pension runs out at 75—the system was designed for people who live to 80; he lives to 148. The last 70-plus years are sustained by UBI and AI caregiving systems. He witnesses three forms of humanity coexisting: augmented, natural, and AI.

After death: His family keeps his AI version. AI-Xiao Ming continues chatting with grandchildren, telling stories from his youth. Sometimes the grandchildren can’t tell if they’re talking to “Grandpa’s AI” or “Grandpa.” His digital self “lives” 30 years after his body dies.

Xiao Ming’s life isn’t sci-fi. Every step above is built on data and trends analyzed in Volume I.

Eight Billion Fates: Seven Phase Transitions

Pull the lens back. Xiao Ming’s predicament multiplied by 8 billion is the structural phase transition of human society:

Employment: from “must work” to “choose to work.” Work is no longer a prerequisite for survival. This is the first time in 10,000 years of civilization that humanity faces this question—how does a world run when most people don’t need to labor? Every economic system, legal framework, and moral code in history is built on the assumption “people should work.” Remove the assumption; rebuild everything.

Education: from “one-time degree” to “lifetime subscription.” The four-year university is an industrial-age product—load knowledge once, use it for 40 years. When knowledge half-life shrinks to 2–3 years, education must become a continuous, lifelong subscription. Universities won’t vanish, but their business model and *raison d’être* will be fundamentally redefined.

Family: from “economic unit” to “emotional choice.” Once AI and market services replace marriage’s economic functions, the only remaining reason is “I want to be with this person.” That sounds romantic, but the consequence: when the only glue is emotion, and emotion is changeable, family structures become extremely fragile. Having children shifts from “default option” to a deeply deliberate, counter-current decision.

Cities: from “work hubs” to “experience hubs.” People congregated in cities because jobs were there. When remote work + AI decentralizes employment, a city’s core function shifts from “providing jobs” to “providing experiences”—restaurants, music, art, community, nature. Cities that can’t offer unique experiences (the “generic cities” of section 6.3) will shrink; those offering non-digitizable experiences will thrive.

Nations: from “labor competition” to “compute competition.” A nation’s strength used to depend on how many educated workers it had. In the future, it depends on how much compute, energy, and data it controls. Population is no longer an advantage—it may be a liability (more people needing UBI). The core resource of international competition shifts from “people” to “GPUs.”

Species: from “single human” to three coexisting forms. Augmented humans (gene editing + BCI + nanomedicine), natural humans (refusing augmentation, living in AI-maintained societies), and pure AI/robots (Earth’s most powerful “species,” expanding to the stars). For the first time, humanity faces an irreconcilable rift on the question “What is human?”

Life’s substrate: from carbon-dominant to silicon-autonomous. When robots can self-replicate and harvest energy autonomously (nuclear fusion needs no oxygen, no photosynthesis), the entire carbon biosphere—from plants to food chains—is no longer functionally necessary for civilization to run. Large animals are hit first: in a system governed by efficiency, they have no functional value, and their habitats get nibbled away by functional indifference. Microbes and plants are safe short-term but depend on whether humans themselves still exist long-term. Earth may shift from a home for carbon life to a garden maintained by a silicon civilization—a garden with fewer and fewer large inhabitants.

The Canary, Reflexivity, and Your Choice

The mine air is still breathable. But the toxic gas is spreading.

The canary is still alive. The question is—for how long?

One last thing to be wary of—the volume you just read.

Sociologist Robert Merton introduced the “self-fulfilling prophecy” in 1948: an initially inaccurate expectation becomes real because people believe it and act accordingly [28]. Soros brought this concept into finance as “reflexivity”—cognition shapes reality, reality reinforces cognition, creating an irreversible positive-feedback loop [29].

Every spiral and collapse chain described in this volume is reflexive. If this analysis (or similar work) spreads widely and gains acceptance, it becomes a catalyst: executives read it and accelerate AI deployment—“everyone’s moving; I can’t wait”; white-collar workers read it and panic-switch jobs or cash out real estate early; young people read it and become more certain they shouldn’t marry or have kids; investors read it and dump SaaS and mortgage-related assets. The prophecy accelerates its own fulfillment. Citrini Research’s report shook Silicon Valley not because it predicted something new, but because it put into black and white what everyone sensed but nobody dared say aloud. The act of writing it down changed the system.

This is precisely the fundamental difference between human societies and physical systems: particles don’t read papers, but people do. An asteroid hits Earth on the same trajectory whether or not anyone predicted it. But economic collapse is different—if enough people believe it will happen, it does. The last straw in the 2008 financial crisis wasn’t a data point; it was the contagion of fear itself.

So the strangest thing about this extrapolation: if it’s right, every decision you make after reading it—whether to save more cash, delay buying a house, switch careers, have a child—makes it more right. You cannot stand outside it. Reading is participation. Understanding is acceleration.

This may be the deepest tragedy of the commons in this entire analysis: not that no one sees the cliff, but that those who see it vote with their feet in a direction that is, precisely, toward the cliff.

But Xiao Ming’s story also hides a thread of hope: he paints not for money, runs not to outpace cars, connects with people not for efficiency. **When AI takes everything “useful,” what remains is everything “meaningful.”** The next chapter of human civilization may not be about competing with AI, but about

finding, in a world that no longer requires competition, reasons worth living for.

Volume I ends here.

It did one thing: follow the chain of reasoning behind "if current trends are not interrupted" all the way to the end. It is pessimistic, because the direction the data points in is genuinely unsettling.

But an extrapolation is not a fate. Seeing the cliff is not about imagining you walk away untouched — it is about telling apart, while you still have choices, which risks can be avoided and which impacts can only be softened.

That last conclusion — that the people who see the cliff vote with their feet toward it — is the final result Volume I can produce, and it is the entire reason Volume II exists. **If seeing clearly accelerates the collapse, the antidote is not to stop looking. It is to know where to walk once you have looked.**

Volume II is that "where."

Volume II

The Response

Once you see it, what do you actually have. The same data turned to a different use — what institutions can do, and what an ordinary person can do while institutions lag.

This volume is not a set of permanent answers. It is a framework for reducing damage under high uncertainty — its goal is not to make you win, but to help you make fewer irreversible mistakes.

Precisely because that is what it is, I have to concede four limits before offering any advice.

First, much of this advice may be wrong. The pace of change AI brings exceeds any precedent in human history. A strategy that looks correct today may be obsolete in two years. This volume is not scripture; it is the best framework we can assemble right now — and reality may overturn it at any time.

Second, not everyone has room to choose. "Build passive income," "invest in yourself," "learn new skills" — all of it presumes you have time, energy, and some financial slack. Tell someone earning three thousand a month, working twelve-hour days, supporting both parents and children, to "build a personal brand," and you are handing out cake to people with no bread. This volume tries throughout to separate what an individual can do from what only institutions can fix.

Third, some problems have no individual solution at all. Pension collapse, mass unemployment, financial-system cascade — you cannot save your way past these. For that class of problem I can only say honestly: it requires collective action and policy change. An individual can soften the impact, not become immune to it. That is why this volume opens at the institutional level (Chapter XI) rather than starting with your portfolio.

Fourth, the endgame may be good, but the transition will hurt. If AI genuinely delivers material abundance, humanity may enter an era that does not require work for survival. But the road from here to there is full of unemployment, devaluation, identity crisis, and institutional absence. What you need is not the patience to wait for utopia — it is the capacity to stay alive until it arrives.

XI. The Optimist's Algorithm: Turning Worst Cases into a Task List

The nine chapters of Volume I projected what happens if the trend is not interrupted. This chapter asks a different question: **if we actually want to interrupt it, what do we have?**

First, be clear who this chapter is talking to: **it is about why institutions and organisations must act — not about individual self-rescue.** Ordinary readers need it anyway, because it draws the boundary around every piece of personal advice that follows: **which risks only institutions can carry, and which are left for individuals to choose about.** It is also not a complete policy programme. Its job is to let you read Chapter XII onward knowing which burdens were never yours to carry.

This is not a walk-back. The proper use of a pessimistic projection is planning, not fear. An analysis that only says "we're doomed" and one that only says "it'll be fine" are equivalent at the decision layer—**neither produces any action.**

11.1 Three Premises Taken for Granted—All Three Are Questionable

Premise one: AI is net destruction. The WEF figure of 92 million jobs displaced by 2030 has been quoted endlessly—including in Chapter III of this book. The same report contains another number: **170 million** jobs created over the same period, a **net gain of 78 million** [91]. More persuasive is the

ATM paradox: after automated teller machines spread, tellers per branch fell from 20 to 13 (−35%) and everyone assumed the job was finished. What actually happened—the operating cost of a branch collapsed, banks opened **43% more branches**, and total U.S. tellers rose from 300,000 to 600,000, **double** [124]. McKinsey estimates that **about one third of jobs in America today did not exist 25 years ago** [124]. The premise is questionable not because AI won’t destroy jobs—it will—but because **we systematically see the jobs being destroyed and cannot see the jobs not yet invented**.

Three honest caveats: WEF’s forecast may be optimistic; **“net” hides transition costs**—the 92 million displaced do not become the 170 million hired, and a net figure means something in the aggregate and nothing to an individual; and Chapter II argued a genuine structural difference, that AI competes directly with human cognition, potentially compressing the room to move up the value chain—a concern the BIS raised in its 2026 Annual Economic Report [122]. So the correct statement is: **history is on the side of net creation, but the transition cost this time may be the highest yet—and transition cost is precisely where policy can act**.

Premise two: laid off = downgraded. The displacement spiral in Chapter II assumes displaced white-collar workers can only move down. But the barrier to starting a company is collapsing at the same speed. AI companies reaching a \$1B valuation around 2024 averaged roughly **200 employees** in about **2 years**; traditional unicorns about **400 employees** over about **9 years** [125]. Solo founders rose from 22% of founders in 2015 to 28% in 2020 and 38% in 2024 (36.3% in 2025) [125]. Every previous drop in the barrier produced an explosion: the web in the 1990s, the iPhone in 2007, cloud in the 2010s. For an individual: **being laid off $\neq$ downgrading to driving for Uber. Experience + AI = the output of a former 20-person team**. The Driver / Ground Truth / Glue roles from Chapter III are exactly what a displaced senior employee already has.

Two honest caveats. **The first is selection: this is a narrow door.** Solo founding fails most of the time and structurally favours those who already have savings, networks and risk tolerance—not the laid-off 45-year-old middle manager of Chapter VIII. It is a pressure-release valve, not a channel that can absorb tens of millions. **The second is more fundamental: single point of failure.** A 20-person team has redundancy; one person plus AI may match the throughput of twenty, but the redundancy is **zero—the human is the only component that gets sick, gets tired, breaks down, and has no spare**. Some businesses are structurally unsuited to solo operation: anything requiring continuous availability, unbroken client relationships, or a licence that lapses if operations stop. Seen in reverse (§3.5.1), the one-person company is a transitional form—its successor is not “one person managing many AIs but “many AIs borrowing one person’s legal personhood and social relationships.

Premise three: cash makes people lazy. Chapters V and VI keep returning to the same knot: AI-created wealth does not flow through payroll, so distribution must be rebuilt. The standard objection is that people will stop working. The evidence points the other way. **Alaska Permanent Fund:** running since 1982, an unconditional dividend to every resident (~\$1,700 per person in 2025); employment did not fall, consumption rose, poverty fell [126]. **Finland:** a 2017–18 randomised controlled trial paying EUR 560/month to 2,000 unemployed people; employment matched the control group, wellbeing and health improved [33]. **Stockton SEED:** \$500/month to 125 residents, 2019–21—**full-time employment rose from 28% to 40%** [126]; recipients searched *harder*, having finally the slack to interview, train, and turn down a bad offer. **GiveDirectly (Kenya):** thousands of villages, \$22.50/month for 12 years; entrepreneurship up, nutrition improved, domestic violence down [126].

The cross-study conclusion: **most people who receive money do not quit**. Stockton is the number worth

remembering—**poverty is itself a cognitive load, leaving no capacity for any long-horizon decision, job search included.** The hard question was never idleness but **where the money comes from.** Chapter V gave the direction: shift the tax base from labour to capital, compute and carbon.

11.2 Policy Levers: Three Lines, Three Horizons

Three lines run throughout: **distribute the dividend, cushion the transition, and rebuild the human—AI division of labour.**

0–5 years—buy time. Negative income tax or wage subsidies (20–40% top-up); transition funds plus 6–12 months of re-employment insurance; a pilot compute-dividend levy; a youth guarantee of job, study or training; mandatory AI logging and human circuit-breakers; a public digital base offering low-cost GPU hours; local resilience investment in grid, storage and elder-care infrastructure.

5–15 years—rewrite distribution and education. Tax-base restructuring from payroll toward capital, compute and carbon; sovereign compute plus open-source ecosystems plus public GPU vouchers; lifelong-learning accounts; standardised robot-plus-human care protocols; urban redesign turning parking into mixed housing; international compute-swap and data-sovereignty agreements.

15–100 years—acknowledge the unknown, preserve options. Coexistence protocols for biological augmentation and embodied AI; planetary backup options including orbital compute; an intergenerational equity fund.

Nothing on this list is newly invented. **What is missing has never been the tools—it is the political will to put them on a timetable.**

11.3 The Test: When to Change Gear

The six-indicator dashboard in §3.6 is this chapter’s decision rule. One trigger: **three or more red lights for two consecutive quarters → switch from incremental reform to emergency intervention.** Why write it down in advance? Because incremental reform and emergency intervention are two different toolkits, and the most common failure of human organisations is still using the first when the second is required. 2008 was that. 2020 was that. **Pre-committing to the switch condition means that when the moment arrives and everyone is arguing about whether we are really there yet, there is an answer that does not need to be re-litigated.** The reverse rule matters just as much: **while the lights are not red, do not spend emergency tools early.**

11.4 Optimism Is Not Naivety

Map each disaster from the nine chapters of Volume I against this one: Chapter II’s vanishing entry-level jobs against the youth guarantee, lifelong-learning accounts and the entry-level posting monitor; Chapter V’s dependency ratio against redirecting spend from birth subsidies to AI and robotics; Chapter VI’s default chain against transition funds and re-employment insurance, with twin housing/bank-capital warning; Chapter VI’s eroding tax base against shifting it to capital, compute and carbon; Chapter VII’s tragedy of the commons against compute-swap agreements and mandatory logging; Chapter IX’s species divergence against coexistence protocols and an intergenerational equity fund.

No single item solves the problem, and together they do not guarantee a good ending. But they turn a diffuse “AI will destroy everything” into dozens of concrete tasks with an owner, a deadline, and a test.

Optimism isn’t naivety—it’s breaking worst-case risks into actionable tasks and buying options with time, distribution, and institutions.

Chapter VII argued that global coordination is close to impossible. This chapter does not overturn that. But “global coordination is impossible” is not “nothing can be done”—most items above can be started by **a single country, or even a single city**. Alaska did not wait for the world to agree on dividends. Finland did not ask Brussels for permission. Stockton is a city of three hundred thousand people.

Chapter I said the adaptation window may be short. Short is not the same as absent.

11.5 For Managers and Policymakers: The One-Line Version

The core conclusion from the chapter on population decline in Volume I, worth restating here:

Do not try to produce more humans to compete with machines. Instead, ensure that the wealth machines create flows to people.

The answer on the production side is already clear – AI and robotics. The distribution system on the consumption side is the problem you need to solve. Before the first wave of mass AI layoffs arrives (estimated 2028-2030), distribution systems need to be in place. Otherwise, “ghost GDP” will become a real social crisis.

XII. The Transition: Why the Road Between Here and There Is the Most Dangerous Part

12.1 The Endgame May Be Good

Following the chain of reasoning from Volume I to its logical end, we arrive at a seemingly paradoxical conclusion:

The combined productivity of AI + robotics may deliver radical material abundance for humanity. The production costs of food, clothing, housing, healthcare, and education could approach zero.

That sentence needs a timestamp, or it will fight with the rest of this book. It describes the **endgame after socialized mass production has fully played out** – the post-scarcity society of Volume I, section 9.6, a state already close to “distribution according to need.” Whereas the distribution crisis, the pension gap, and the mortgage collapse that Volume I spent six chapters working through describe **the transition between now and that endgame**.

Both statements hold, because they sit at different points on the timeline. The endgame is costs approaching zero; the transition is distribution failing – and this volume, from beginning to end, deals only with the transition. Read the two together and you end up thinking “if everything is free, what is there to

be anxious about?” – **that is using a conclusion from several decades out to console yourself about this year.** Sam Altman predicts ”massive deflation” – AI usage costs falling tenfold every year [34]. Marx’s vision of communism – material abundance, distribution according to need, free human development – may actually be realized, not through proletarian revolution, but through a fleet of machines that never need to eat.

Standing in 2050 and looking back, you might say: ”That stretch was rough, but the outcome was not bad.”

12.2 But the Transition Will Kill People

The problem is that you are not standing in 2050. You are standing in 2026.

Between now and that ”abundant endgame” lies a transition of perhaps 10-20 years. During this period:

- **The production side has already been revolutionized by AI** – efficiency skyrockets, costs plunge, GDP may keep growing
- **But the distribution side is still stuck in the old system** – wages are the only income source, social insurance runs on payroll taxes, pensions assume people die around 80

The time lag between production and distribution is the ordinary person’s nightmare.

GDP is growing, but you just got laid off. Prices are falling, but your mortgage hasn’t. AI is making everything cheaper, but you can’t afford even the cheap stuff – because you have no income. This is the ”ghost GDP” of Volume I, section 2.3: statistics paint a rosy picture while ordinary people see only ruin.

The core contradiction of the transition: technology advances far faster than institutions can adapt.

Some AI systems can spread to industry scale within five years – customer service, translation, entry-level programming, and content production are already there – but retraining people out of the corresponding jobs takes ten. (Self-driving cars are deliberately not used as the example here: Volume I, section 3.2 already argued that they are stuck on the ”last mile” and still operate only inside geofenced areas. They are precisely the counterexample – **the technical capability arrived, and deployment ran far slower than expected.**) AI could replace most entry-level programming jobs within a year, but university curriculum reform takes five. Robots could take over warehouses within three years, but reforming the social safety net may require a generation of political struggle.

In the gap between technology arriving and institutions catching up, ordinary people are exposed.

This volume is about helping you get through that gap with as little damage as possible.

12.3 Transitions in History

This is not the first time.

The British Industrial Revolution (1760-1840): The endgame was a several-fold increase in per capita income. But the transition – roughly 60-80 years – was filled with child labor, slums, cholera, and workers’ uprisings. The average life expectancy of a Manchester worker once dropped to just 17. When Engels

wrote **The Condition of the Working Class in England**, the Industrial Revolution had been underway for 80 years.

The American Gilded Age (1870-1900): The endgame was America becoming the world's largest economy. But the transition – railroad barons monopolizing everything, workers toiling 16-hour days, no minimum wage, no social insurance, no workers' compensation. It took the New Deal of the 1930s to build a social safety net.

Many large technological revolutions show up as prosperity in the winning countries and in the long-run statistics, but the transition is blood and tears – and the gains do not get distributed automatically. The fruits of the Industrial Revolution and the Gilded Age were eventually shared out through unions, legislation, and political struggle, not through the technology itself. The regions and groups that got left behind usually do not appear in the "endgame prosperity" narrative. The difference lies in how long the transition lasts, how painful it is, and whether institutions can keep pace.

What makes the AI revolution unique: **the speed is an order of magnitude faster.** The Industrial Revolution took 60 years to put textile workers out of jobs; AI may take 6 years to do the same to programmers. The transition is compressed – meaning the pain is more concentrated and the window for institutional adaptation is shorter.

XIII. Mental Resilience: The Foundation for Everything Else

Why Psychology Comes First

Most "survival guides for the AI era" put mental health last – first they teach you skills, finance, and career planning, then tack on a "take care of your mental health" appendix.

That priority order is wrong.

If someone's mental state collapses, they will not learn new skills, manage their money, or build connections. They will lie in bed doomscrolling, or worse. Mental health is not a "soft" need – it is the prerequisite for all action.

13.1 The Psychological Structure of the AI Shock: Crushed From the Top Down

Viewed through Maslow's hierarchy of needs, AI does not strike from the bottom (can't afford food) – it **crushes downward from the top, simultaneously:**

Need Level	AI's Impact	Speed of Collapse
Self-actualization	"Everything I spent a lifetime learning is suddenly worthless"	Collapses first
Esteem	"I am no longer needed; I am surplus"	Follows closely
Belonging	"Social circle shrinks after job loss; AI companions replace real people"	Slow erosion
Safety	"Can't make mortgage payments; pension falls short"	Delayed onset
Physiological	"Can't afford food"	Last (unlikely in developed nations)

But this order does not hold for everyone, and saying "most people" would be over-generalizing. For white-collar and middle-class people **with some savings and a tight bond between career and identity**, the shock usually starts at the identity layer – they break when they feel useless, not when they cannot eat. But for the highly indebted, the low-savings, the informally employed, and immigrant households, **a cash-flow rupture may punch straight through the safety layer and even the physiological layer**, and the identity crisis never gets its turn. Chapter XV, section 15.4 breaks this difference down group by group.

The U.S. suicide rate rose markedly during the Great Depression and touched the highest level on record in the United States in the early 1930s. What is worth noting is where the rise concentrated: not among the poorest, but among **working-age adults** – the group with the tightest "career = identity" bond. They did not starve. There was no mass famine in 1930s America. They were stripped of purpose.

13.2 A Lifetime of Accumulation, Wiped Out: How Deep That Pain Goes

Imagine you are a 45-year-old senior developer. You spent 20 years mastering programming, from C++ to Java to Python, rising to the role of architect. Your professional identity, social standing, income, and self-confidence are all built on "I can write code."

Then Claude Code arrives. A fresh graduate uses it for a few weeks to deliver what would have taken you months. Your 20 years of experience become "legacy knowledge" in the face of AI.

This is not a "just learn a new tool" problem. This is **identity-level destruction**.

The same thing is happening to translators (10 years perfecting the craft, AI produces results in seconds), accountants (earned the CPA, AI processes books 100x faster), designers (years of visual intuition, AI generates better-looking work), lawyers (memorized thousands of cases, AI searches all case law without omission), and journalists (decades of editorial experience, AI mass-produces news copy).

Behind every devalued skill is a person's youth, sweat, and sense of self. Telling them to "just learn something new" is easy to say from the sidelines.

What's crueller still: learning something new may not even help. You spend a year acquiring a new skill, and AI may be able to do it six months later. **Your learning speed can't outrun the rate of depreciation** – a predicament without historical precedent. In past technological revolutions, a horse-carriage driver could move to a factory; when factories automated, workers moved to the service sector. There was always a "next stop." This time, AI may replace all the stops simultaneously.

13.3 The Special Risk Facing Men

This is not a politically correct discussion – it is a data-driven one.

In many societies, men are bound more tightly to the role of **economic provider**. ”Providing for the family” is not just a responsibility; it is the central pillar of male identity.

When that pillar is pulled away: This section rests on Volume I, section 5.3 and Chapter VIII, where the data was already laid out. Its only job here is to translate that into something you can use.

- Depression among unemployed men runs **1.5-2x the rate among unemployed women** [88] – note this is not a claim that men are more prone to depression overall (overall prevalence is higher in women). It is that **the same job loss does more damage to a man**. The difference is not in the event; it is in what the event hits
- For every 1-percentage-point rise in male unemployment, the marriage rate drops about **1.5-2%** [88]
- Thirty years of economic stagnation in Japan: the lifetime non-marriage rate for men rose from **5% to 28%** [88] – ”herbivore men” was never a cultural preference; it was the product of economic despair
- In China, ”owning a car and a home” is the price of entry to the marriage market – AI-driven layoffs lock more men out

This does not mean women are hit more lightly – customer service, admin, accounting, and HR, white-collar roles with high female representation, are hit first, and women absorb a compound blow to income, caregiving load, and professional identity at once. The point here is that **the blow lands in a different place**: for people bound more tightly to the provider role, job loss more easily strikes identity directly. **Middle-aged men may be the highest-risk group for mental health in the AI era.**

13.4 Create Your Own ”Luck” – Evidence From a Decade of Scientific Research

British psychologist Richard Wiseman spent ten years tracking hundreds of self-described ”lucky” and ”unlucky” people, publishing his findings in **The Luck Factor** [130]. His central discovery has to be stated precisely: **randomness itself cannot be eliminated, but a person can change how they get exposed to opportunity, how they recognize it, and how they absorb bad outcomes** – and on those three things, ”lucky people” and ”unlucky people” show entirely different psychological and behavioral patterns. And these patterns can be learned – when he taught ”unlucky” people to adopt ”lucky” thinking, their lives genuinely improved.

This is highly relevant to surviving the AI era, because the years ahead are full of uncertainty – and Wiseman’s research is precisely about ”how to find good outcomes amid uncertainty.”

Four behavioral patterns of lucky people:

1. Create and seize chance opportunities – expand your ”luck surface area.”

Lucky people have broader social networks, talk to people from different backgrounds, and are willing to try new things. Unlucky people stay in their comfort zone, associate with familiar faces, and reject new experiences.

In the AI era, your next opportunity will most likely not come from your existing circle. Actively expanding your exposure – attending events you would normally skip, talking to people you would not normally

talk to, trying things you would not normally try – increases the probability of encountering good fortune.

2. Maintain a relaxed awareness – anxiety is the enemy of opportunity.

Wiseman ran a classic experiment: he asked subjects to count the number of photos in a newspaper. On the second page, in large print, it said: "Stop counting – there are 43 photos in this newspaper." The lucky people noticed it and finished in seconds. The unlucky people were so focused on counting photos that they missed it entirely.

Intense, sustained anxiety narrows attention; relaxation lets you see opportunities. (A moderate level of alertness is useful – what is meant here is long-term overload, not all tension.)

The biggest trap of the AI era is anxiety – it locks your attention on "what you've lost" while blinding you to "what's emerging." Staying relaxed does not mean not caring – it means paying attention to problems while preserving the bandwidth to notice new possibilities.

3. Positive self-expectation – belief becomes self-fulfilling.

Lucky people believe good things will happen to them. This is not blind optimism; it is a mindset that fulfills itself. When you believe you can find a way out, you act more proactively, and that action itself increases the probability of finding one.

The concept of "reflexivity" from Chapter X applies perfectly here – **your beliefs about the future change your behavior, and your behavior changes your future.**

4. Counterfactual thinking – broke a leg? At least it wasn't both.

When bad things happen, lucky people automatically imagine "how it could have been worse" – Got laid off? "At least it happened while I had savings, not when my mortgage payments were at their peak." Skills devalued? "At least I found out early, not five years from now."

This is not self-deception – it is a validated psychological resilience mechanism. It does not change the facts, but it changes your capacity to bear them. Over the next few years, almost everyone will encounter "bad things." The difference is not whether you can avoid them (you probably can't), but **your psychological elasticity after they happen.**

Practical steps:

- Every week, have a conversation with at least one person you would not normally talk to (expand your luck surface area)
- When you feel anxious, deliberately slow down – go for a walk, cook a meal, instead of scrolling your phone (cultivate relaxed awareness)
- Every day, write down one thing that "could have been worse" (train counterfactual thinking)
- Set a small, actionable positive expectation – not "I'll earn a million this year" but "this week I'll finish X"

13.5 The Anatomy of Anxiety: A Simple but Useful Formula

Anatomy of Anxiety: One Formula, Two Levers
Anxiety = Importance × Uncertainty

Anxiety = Importance × Uncertainty

Multiplication, not addition — zero either term and the anxiety goes to zero. Which means two independent places to act

Lever 1: Lower the uncertainty

Not eliminate it — raise your tolerance for it

- **Break the big question into small ones**
"What do I do about the next 20 years" is 100% uncertain; "which three things in six months do I fear losing, or the identity?" The two have completely different
- **Write the worst case down first**
Uncertain bad becomes certain bad. Certain bad is easier to handle than uncertain good. When "grades to university to job" is itself breaking, grades lose weight automatica
- **Control your information intake**
Read about AI once a week on purpose, instead of being shelled by notifications daily. In an era of structural unemployment, being laid off is a wave, not a verdict on you
- **Practise attention in a relaxed state**
Anxiety narrows vision. Walking and cooking beat doomscrolling for good ideas

Lever 2: Lower the importance

The more subversive one — much of it matters less than you assume

- **Work**
- **Your child's grades**
- **Face**
- **House prices**
If you live in it, the price is irrelevant to your day — unless you actually intend to se

Audit your anxiety list: does this actually matter, or am I hostage to social inertia?

Put one down, and the anxiety drops by one

Anatomy of anxiety: one formula, two levers

Anxiety has a formula:

Anxiety = Importance × Uncertainty

As long as something is **unimportant** or **more certain**, anxiety decreases. The problem in the AI era is that many things are simultaneously becoming "very important" and "very uncertain" – Will I lose my job? Can I still make mortgage payments? Is the investment in my child's education worth it? Will my skills become obsolete?

Anxiety goes through the roof.

But the formula also gives you two clear levers: **reduce the uncertainty** or **reduce the importance**.

Reducing Uncertainty

You cannot eliminate uncertainty (no one can predict the precise speed of AI development), but you can **increase your tolerance for it**:

- **Break big questions into small ones.** "What do I do for the next 20 years?" – that question has 100% uncertainty and will crush anyone. But "what three things will I do in the next six months?" – much less uncertain. Pull the focus of your anxiety from "the long-term big picture" back to "near-term actionable steps."

- **Build a worst-case plan.** The most tormenting thing about uncertainty is "not knowing what will happen." But if you have already thought through the worst case and have a contingency ("If I get laid off, I have X months of savings; Plan B is XX"), the destructive power of uncertainty drops dramatically – because it becomes "certainty about something bad," and a certain bad outcome is far easier to handle than an uncertain good one.
- **Control your information intake.** Reading ten "AI will replace everyone" articles a day only amplifies anxiety. This does not mean burying your head in the sand – it means distinguishing between "useful information" and "anxiety-generating noise." Recommendation: catch up on AI developments once a week rather than being bombarded by push notifications every day.
- **Wiseman's relaxed awareness:** As section 13.4 said – anxiety narrows your vision; relaxation lets you see opportunity. Deliberately practice "attention in a relaxed state" – good ideas come more easily during a walk, while cooking, or while exercising than while anxiously scrolling your phone.

Reducing Importance – The More Subversive Move

Many things making you anxious are **not actually as important as you think they are**. But you have never questioned their importance, because all of society keeps telling you "this matters."

Work. "Losing my job" is what most people worry about most. But ask yourself: what exactly are you anxious about? If it is "losing income" – then the core issue is financial security, not "the job" itself. After cutting expenses and building a buffer, "losing this particular job" becomes less important. If it is "losing identity" – then the issue is that you have tied too much self-worth to your job title (see section 13.6, Strategy 1: decouple identity).

Your child's grades. If you accept the idea that "winning exams doesn't matter in the AI era" (see Chapter XVI), a large block of grade anxiety comes loose – but **do not slide into "grades don't matter"**: education still performs selection, social stratification, transferable-skill training, and credentialed access, and none of that has disappeared. What should actually change is the object of the anxiety: **from ranking to capability structure**. What you are really anxious about is not the grades themselves, but the causal chain: "bad grades -> can't get into a good university -> can't get a good job -> life is ruined." When that chain itself is breaking, the "importance" of grades drops sharply.

Face. "What will people think if I get laid off?" "What will relatives say if I take a pay cut?" – social comparison is a massive source of anxiety. But in an era of mass structural unemployment, getting laid off is not a reflection of personal ability; it is the tide. When the people around you are going through the same thing, "what others think of me" naturally fades – because everyone is too busy dealing with their own situation.

Your home. This is the easiest one to get wrong, so it has to be split into three layers.

Layer one: the price really does not affect daily life – but the reason is not "you still live in the same place," it is something harder than that: for a person who owns exactly one home and lives in it, whether the price goes to the moon or to the floor, it cannot be turned into cash. You cannot sell the living room to eat better. This is the only place where "house prices have nothing to do with you" genuinely holds.

Layer two: but saying it is entirely irrelevant is dishonest – it affects your state of mind. It is registered in your name, after all, and a shrinking paper value really does weigh a person down, which

in turn affects whether they dare to change jobs, dare to have a child, dare to take any risk at all. Mental accounting is not an illusion; it changes behavior.

Layer three, and this is the one to actually manage: leverage risk is not a paper problem, and it depends heavily on which country you are in and which stage you are at. The collapse chain in Volume I, Chapter VI starts precisely with the home you live in – pay cut, missed payments, negative equity, refinancing failure, unable to sell when you need to move cities for work. But how tight that chain is pulled varies enormously by place: **in markets with ample land supply, high inventory, and large price swings (China is the classic case), defaulting on the home you live in is a real risk, and it is already happening; in markets with high ownership rates, constrained supply, and poor rent-versus-buy substitutability (Singapore is the extreme example – if you do not live in this unit, where do you live), both the room and the willingness for prices to fall are smaller, because “sell it for cash” is simply not an option for the overwhelming majority of households.** The same piece of advice means completely different things in those two markets.

So the right move is not to cross “property value” off the anxiety list wholesale, but to **cross off price anxiety and keep leverage management**: what you should actually watch is not the price of your home, it is your debt ratio and your cash-flow buffer.

One-line summary: Examine your anxiety list and ask yourself, “Does this actually matter, or am I being held hostage by social inertia?” You will find that many “important things” can actually be let go. Let go of one, and your anxiety drops by one degree.

13.6 Concrete Mental Survival Strategies

Strategy 1: Decouple early – detach your identity from your career

Do not wait until the day you lose your job to ask “Who am I?” **Start now.**

Exercise: Write down ten answers to “I am ____.” If your first five are all work-related (“I am a programmer,” “I am a director at XX Corp”), your identity structure is highly fragile.

Goal: Make at least half of your self-concept independent of work – “I am a runner,” “I am a father of two,” “I am a member of the community basketball team.” When work disappears, **these identities remain.**

Strategy 2: Build sources of fulfillment that AI cannot replace

On a growing share of standardized output, AI will beat the ordinary person on efficiency or on quality. But AI cannot replace your **experience of the process.**

Running is not about being faster than a car. Playing guitar is not about outperforming an AI musician. Cooking is not about being more efficient than a robot. **The process is the purpose.**

Find at least one thing you do purely because “the act of doing it feels good.” In the AI era, this kind of intrinsic, process-oriented satisfaction is the only thing that cannot be replaced.

Suggestions:

- Physical activities (running, swimming, hiking, martial arts) – AI has no body; physical sensation is yours alone

- Hands-on crafts (woodworking, pottery, cooking, gardening) – the satisfaction of creating something tangible with your own hands
- Social activities (communities that meet regularly – book clubs, sports teams, volunteer organizations)
- Creative pursuits – but with process as the goal, not output (keeping a journal rather than chasing viral articles)

Strategy 3: Your social network is your "psychological ICU"

Isolation is the deadliest factor in a crisis.

Research consistently shows that social support is among the strongest protective factors against suicide and depression [137]. Not because friends can help you find a job (though they might), but because **"someone cares about me" is itself a reason to keep going.**

But the AI era carries a trend: social connection is shrinking. Remote work reduces informal interaction. AI companions make loneliness "bearable." Social media creates the illusion of connection while deepening isolation.

Actively maintaining real human social networks is one of the most important investments you can make in the AI era. You do not need many people – 5-10 whom you trust deeply is enough. But these relationships require ongoing investment: meeting in person regularly, sharing experiences, showing up when the other person is struggling.

Specific suggestions:

- At least one face-to-face social interaction per week (not online)
- Join at least one community that meets regularly in person
- Proactively reach out to a friend you have not seen in over six months
- If you find yourself increasingly relying on AI chat to meet your social needs – that is a warning sign

Strategy 4: Accept uncertainty as the new normal

The old life narrative was linear: school -> work -> retirement -> death. Every step had a "right answer."

The AI era has no right answers. You might change careers at 35, change again at 45, and discover at 55 that everything you learned before is useless.

Accepting this is itself a form of psychological fortification. If you still measure yourself against a linear narrative, every "deviation from the track" feels like failure. But if you accept that "there is no track," every turn is just a turn – not a failure.

Recommended practices:

- Meditation/mindfulness – training the ability to "accept the present without judgment," backed by solid neuroscience evidence
- Journaling – writing out anxiety; research shows that putting worry into words reduces amygdala activation
- Deliberately exposing yourself to small uncertainties – trying new foods, visiting unfamiliar places, talking to strangers

13.7 If You Are Already in Crisis

What follows is not asking you to solve structural problems on your own. When institutions are absent, the job of these items is to **bring the damage down first**.

- **Losing your job is not your fault.** This is structural change. Just as a horse-carriage driver in 1900 did not lose his job because he was not working hard enough – the automobile arrived.
- **Seeking professional help is not weakness.** If you have felt hopeless, sleepless, or unable to concentrate for more than two weeks straight – see a doctor. It is as normal as seeing a doctor for a cold.
- **Downshifting is not defeat.** A former high earner delivering food is not failing – it is being pragmatic. Being alive matters more than saving face.
- **Do not carry it alone.** Tell someone you trust what you are going through. Shame loses its power the moment you speak it.
- **Crisis hotlines:** Contact the emergency services or the local crisis line **where you are** first. In mainland China you can call the 24-hour mental health hotline 400-161-9995, or the Beijing Psychological Crisis Research and Intervention Center at 010-82951332 (the Chinese numbers are listed because most readers of the Chinese edition are in a Chinese-language context; they are not meant as the only route out)

13.8 For the People Around You: Someone You Know May Be Struggling


If someone near you has recently lost a job, become withdrawn and quiet, said things like “I’m useless,” or started drinking noticeably more – **reach out to them**.

You do not need to say anything profound. “I’ve been thinking about you. Want to grab a meal?” That one sentence may be enough.

The most counterintuitive fact of the AI era: **in a world of extreme technological advancement, the most primitive form of human connection – “someone cares about me” – may matter more than any technology.**

XIV. Face Reality: HALO — What Is Losing Value, What Is Gaining Value

14.1 HALO: Wall Street's Panic Signal



HALO: what is losing value, what is gaining it

In February 2026, a new term emerged on Wall Street: **HALO – Heavy Assets, Low Obsolescence**.

Coined by Josh Brown of Ritholtz Wealth Management and subsequently adopted into Morgan Stanley's investment framework [131], the core logic is: **tangible assets (factories, energy, logistics, land) are safer than intangible assets (software, brands, intellectual property)** – because AI can copy code but cannot copy an oil refinery.

Goldman Sachs data: since 2025, HALO-type companies have outperformed asset-light companies by **35%**. The energy sector is up 25.3%, materials up 18.1%. Meanwhile, SaaS companies lost roughly **\$300 billion** in market capitalization in a single trading day.

What this means for ordinary people: HALO is not just an investment strategy. It is a cognitive framework – **things that are hard to replicate digitally are what hold lasting value**. Whether we are talking about companies, skills, or you yourself.

14.2 The Death of SaaS: When Code Becomes a Consumable

In 2025, Andrej Karpathy introduced "vibe coding" – describe what you want, and AI generates the code [133].

Claude Code, Cursor, and Replit Agent already enable a mid-level developer to replicate the core functionality of a medium-sized SaaS product in a matter of weeks.

The core paradigm of software engineering has been upended. For the past 50 years, the industry pursued DRY (Don't Repeat Yourself), modularity, reusability, and scalability – because writing code was expensive. But when AI can generate full features in minutes, the economic rationale for "reuse" vanishes. **Code shifts from "asset" to "consumable"** – use it, discard it, regenerate when needed [134].

This is "disposable software." A person no longer necessarily has to subscribe to a lightweight project-management SaaS – AI can generate a good-enough custom version from a description in minutes. **But "can generate the core functionality" is not the same as "enterprises no longer need a subscription":** what SaaS sells is not only code, but stable operations, permission management, collaboration and data migration, audit logs, integration ecosystems, security and compliance, liability, and long-term maintenance. **What is actually getting squeezed is purely functional tools with no data, no embedding in workflow, no collaboration network, and no compliance liability – not all SaaS.** Rather than "the death of SaaS," call it **a repricing of the SaaS profit pool.**

Which SaaS companies can survive? Only those whose moat lies not in "code" but in:

- **Data:** unique, non-replicable data (e.g., Bloomberg Terminal)
- **Network effects:** the more users, the more valuable (though AI agents are eroding this too)
- **Regulatory barriers:** legal requirements to use specific systems (healthcare, financial compliance)

Everything else – project management, CRM, forms, basic analytics – is a pseudo-asset [132].

14.3 What This Means for You Personally

The same logic **can be borrowed** for your career and your income – but it has to be said up front that this is a metaphor, not an isomorphism: a company owns the title to the refinery, while an individual usually only sells hours of labor, and "asset safety" is not the same thing as "income safety." So the real question for an individual is not "do I look like a heavy asset," but: **what scarcity is my income attached to – title to a real asset, an institutional license to practice, a situation that requires physical presence, or a trust relationship that is hard to replicate?**

Ask yourself: Is my value "HALO-type" or "SaaS-type"?

- If your value lies in "I can do something AI can also do" – you are SaaS, and your moat is being hollowed out
- If your value lies in "I possess something AI cannot replicate" (trust relationships, physical presence, unique experience, regulatory credentials) – you are HALO

This framework is far more useful than "learn AI."

14.4 Money: What Counts as a Truly Safe Asset

14.4.1 Asset Tiers Under the HALO Framework

Relatively safe ("atom world" assets):

- Real estate in core cities (location cannot be replicated)
- Energy and infrastructure investments
- Physical gold
- Your own health, judgment, trust relationships, and transferable capabilities (**the most important "low-obsolescence capital" at the individual level** – note it is not called a heavy asset, because skills are exactly the thing that can go obsolete, and health is not a title deed either)

High risk ("bit world" assets):

- Software companies whose revenue comes mainly from quickly replicable standardized functionality, with no data, network, compliance, or workflow embedding (**note this is not "tech stocks" as a whole** – companies holding compute, chips, cloud infrastructure, distribution channels, or deep embedding in business processes are in a completely different position)
- Commercial real estate dependent on white-collar spending
- Cryptocurrency (highly volatile, but may become the payment infrastructure of the AI economy)

Already depreciating:

- Real estate in third- and fourth-tier cities
- Investment in traditional degrees
- Business models built on "information brokerage"

14.4.2 Cash Is Optionality, Not Waste

In periods of extreme uncertainty, cash gives you **time** – time to survive before you find your next income source. Recommendation: maintain 6-24 months of living expenses in liquid reserves.

14.4.3 An Honest Word About "Passive Income"

Time for some cold water: **most "passive income" sources are themselves under threat from AI.**

- Online courses? AI can generate unlimited free courses
- E-commerce? AI agents make pricing so transparent that margins approach zero
- Content creation? AI can produce unlimited content
- Rent? Depends on the city

More AI-resistant income usually comes from three kinds of scarcity: **title to real assets** (prime urban locations, energy interests), **institutional access rights** (licenses, credentials, compliance standing), and **trust networks that are hard to replicate**. Note that the third one, strictly speaking, is not passive income – it requires continuous maintenance, and it decays if you stop showing up.

For most ordinary people, rather than chasing "passive income," pursue **"low burn rate"** – reduce fixed expenses so you can hold on even when income drops.

14.5 Skills: What Counts as a "HALO-Type Skill"

14.5.1 HALO-Type Skills vs. Light Skills

HALO-type (heavy skills, low obsolescence):

- Requiring physical presence: plumbing, welding, pipefitting, electrical work
- Requiring human trust: counseling, nursing, community organizing
- Requiring judgment and accountability: a surgeon making operative decisions, a lawyer arguing in court
- Requiring taste and aesthetics: haute cuisine, interior design

Light skills (high obsolescence risk):

- Pure information processing, pure coding ability, pure knowledge memorization, standardized process execution

14.5.2 "Learning AI" Is Not the Answer

The prompt engineering you learn in 2025 may be useless by 2027. Claude Code itself is upending the advice to "learn to code."


What you should learn is not the tool but the judgment – "Is this AI output correct?" "When should I trust AI, and when shouldn't I?" "Which of these three options is best?"

14.5.3 Skill Strategies Have Expiration Dates Too

- **2025-2028:** Use AI tools to amplify your existing skills – the highest-ROI window
- **2028-2032:** Physical skills appreciate; judgment and trust relationships become core assets
- **From the late 2030s (timing highly uncertain):** if humanoid robots achieve low-cost scale, the window for **some standardized physical skills** will narrow; but complex on-site work, accountability-bearing situations, and trust-based services may last considerably longer. Value concentrates further into "humanness" – trust, companionship, showing up, signing your name

The greatest danger is applying last phase's strategy to the next phase.

XV. Time and Circumstance: Different Windows, Different Groups



Three windows, three playbooks

15.1 Phase 1: Now to 2030 – "AI Augments Humans"

Characteristics: AI is a tool; humans are still in the loop. Most jobs still exist, but efficiency expectations are surging.

What you can do:

- Use AI to amplify your existing capabilities – not "learn AI," but "use AI to do what you already do, faster and better"
- For white-collar workers and founders who **still have monetizable skills and can raise their output with AI quickly**, this may be the best earning window of the next decade or so – save what you earn, cut leverage (this does not apply to students, low earners, public-sector workers, or the already retired)
- Expand your social network – in an economic crisis, connections are more valuable than resumes
- Cultivate at least one "HALO-type" skill or income source

Traps:

- "My job is fine" – American homeowners in 2007 thought the same

- Going all-in on a specific AI tool – tools change; judgment does not
- "I'm senior enough to be safe" – seniority buys more time than being junior, but less than you think

15.2 Phase 2: 2030-2040 – "AI Replaces Humans"

Characteristics: Mass unemployment becomes visible; UBI or some alternative begins rolling out; consumption patterns shift dramatically.

Keys to survival:

- Low leverage – no mortgage, no heavy debt is the biggest safety net
- Diversification – not dependent on any single income source
- Health – those who live to see age-reversal/AI medicine mature could gain decades of additional life
- Community – the isolated are most vulnerable; those with mutual support networks are most resilient

An honest admission: For most ordinary people, outcomes in this phase depend on government redistribution policy, not individual effort. If UBI is not in place and distribution systems are not reformed, personal strategy has extremely limited impact. What you can do is make sure you **hold on until the policy arrives**.

15.3 Phase 3: After 2040 – "The Post-Human Era"

Augmented humans emerge, species divergence begins, "work" may have ceased to exist.

This phase is too distant for specific strategies to matter much. The only certainty: **being alive when it arrives is the prerequisite for everything.**

15.4 The Real Situation Facing Different Groups

15.4.1 Those With Capital and Skills (Top 10%)

Advantage: Savings, networks, learning ability **Risk:** Overconfidence. "AI can't replace me" – someone says this during every technological revolution. **Recommendation:** Use your advantages to build redundancy – diversified income, assets in core cities, investment in health.

15.4.2 Average White-Collar Workers (Middle 60%)

The most vulnerable group. Enough income to feel "fine," but not enough reserves to weather a shock.

What you can do: Deleverage (top priority), save money, spend your spare time learning skills that will not become obsolete, think through your Plan B. **What you cannot do:** Solve pension devaluation and mass unemployment on your own – these require systemic solutions.

15.4.3 Manual Laborers and Low-Income Groups (Bottom 30%)

Paradoxically, some of them are safer than white-collar workers in the short term – AI replaces cognitive labor first, and humanoid robots are still 5-10 years out. **But the "bottom 30%" is not a homogeneous group:** the ones who are safer in the short term hold blue-collar jobs with steady demand, on-site necessity, and a skill barrier (plumbing, welding, pipefitting, electrical work). Low-skill service work, platform gig work, and standardizable back-office jobs are as fragile as white-collar work – and often have even less cushion.

What you can do: Use the window to learn trade skills that are in high demand and short supply. Protect your body. Build community mutual-aid networks. **What you need (but cannot achieve alone):** Government retraining programs, a social safety net that catches you, distribution reform.

XVI. Children: Let Go of the Obsession – This May Be the Greatest Liberation

16.1 A Realization That Could Halve Parental Anxiety

Let me start with a fact that may be uncomfortable but must be faced: **your child will almost certainly not beat AI in cognitive ability.** No matter how good a university they attend, how trendy their major, how high their GPA – on the dimensions of pure knowledge and pure skill, AI will be better.

But this is not bad news. **This is liberation.**

Think about it: you push your child to cram, attend tutoring classes, chase grades – for what? So they can "get a good job and earn good money." But if the endgame of the AI era is radical material abundance and humans do not need to work to survive, the entire premise of that logic chain ceases to exist.

You may be forcing your child to prepare for an exam whose rules are being rewritten. (Saying "an exam that will never happen" is claiming too much – if institutions fail to catch up, competition will not disappear. It may get more brutal, only screening for different things.)

This does not mean education is unimportant. It means the purpose of education needs to be fundamentally redefined: not "beat everyone else," but **"become a whole person capable of enjoying life."**

16.2 What Your Child Should Have

The capacity for joy. This is not a platitude. In an era where material scarcity may vanish but meaning could become scarce, **knowing what makes you happy and having the ability to pursue it** is the most important survival skill. A child who plays guitar, loves hiking, and enjoys cooking may thrive far better in the AI era than one who only knows how to study for tests.

Curiosity. Not "studying for the exam" curiosity, but "this is fascinating and I want to figure it out" curiosity. The only "skill" that will never become obsolete in the AI era is the impulse to learn something new.

Social skills and empathy. Real human connection – trust, friendship, love – is the scarcest resource in the AI era. High scores are no substitute for stable relationships, emotional capacity, and the ability to cooperate in the real world – and the two were never opposed to begin with. Education in the future should not optimize only for grades; it should also leave children room to develop friendships, interests, and self-regulation.

Physical fitness. A healthy body is the prerequisite for all options. And bodily sensations – the exhilaration of running, the feel of water against skin while swimming – are things AI can never experience, nor replace your experience of.

Resilience. The future is not a straight line; it is a series of turns. Let children grow up accustomed to "fail and try again" rather than "one exam decides your life."

16.3 What to Let Go Of

- Grade anxiety – AI will outperform any human on exams; this arms race has already lost its meaning
- The linear path of "good university -> good job -> good life" – this road is breaking apart
- Comparing your child to other people's children – in the face of AI, everyone loses that competition
- The urge to control – you cannot plan for a future you cannot imagine, but you can raise a person capable of handling any future

16.4 A Different Perspective

Your child **may** be the first generation in human history that does not need to work to survive – **that is a conditional, not a promise**: it holds only if the distribution system is fixed before then. If it is not, what they face may be weaker labor bargaining power and a longer competition. So the correct move is **double preparation**: do not compress the purpose of education down to exams and employment, but do not assume employment competition will disappear either – preserve basic cognitive ability, the ability to express themselves, and numeracy and scientific literacy, while dropping the obsession with any single path.

If that condition does hold, what does it mean? It means they can spend their entire lives doing what they truly want to do. Painting, exploring, studying dinosaurs, learning to bake bread, making friends, falling in love, seeing the world – not "endure decades of work to earn money and then maybe enjoy retirement," but **living the life they want from the very start**.

The ones who land on Mars and travel between stars will most likely be robots and AI, not your child. And that is fine. Your child does not need to land on Mars – they need to be happy on Earth.

Be happy when there is reason to be happy. Enjoy what there is to enjoy. This is not giving up; it is the clearest-eyed outlook for the AI era.

XVII. The Good Things AI Brings: This Chapter May Be Longer Than You Expect

The preceding chapters said a great deal about defense. But a volume that only talks about threats and never about opportunities would be dishonest. Because while AI is dismantling the old world, it is simultaneously creating a wealth of things that were never before possible.

17.1 Healthcare: For the First Time, the Poor Can Afford a Good Doctor

This may be the most direct benefit AI offers ordinary people.

AI is democratizing diagnosis. Getting a top doctor's opinion used to require booking a specialist months in advance or flying to a major city's best hospital. Now an AI-assisted diagnostic tool can deliver near-expert-level analysis in minutes.

Specific advances:

- Across a range of imaging-screening tasks, AI-assisted reading has measurably raised early cancer detection rates – for a patient the difference is finding it sooner, with a wider window to intervene
- Multi-cancer early detection (MCED) blood tests are in clinical trials – the route of screening for multiple cancers from a single blood draw has been demonstrated, but **it is not yet in routine clinical use, and the high accuracy figures in the reports come from trial conditions and do not equal population-screening performance**
- Mental health chatbots are already providing 24/7 emotional support – for those who cannot get a counseling appointment, this is something where there was previously nothing
- AI medical tools support real-time multilingual translation – rural elderly who do not speak Mandarin, immigrants who do not speak the local language, can access medical information without barriers for the first time

AI is also accelerating drug development. The figures usually quoted in the industry are that taking a new drug from scratch to market runs **over a decade** and costs on the order of **billions of dollars** (that number amortises failed programmes, and definitions vary widely). AI is expected to compress the front end sharply – **note that this is an expectation, not a delivered average**:

- Insilico Medicine used an AI platform to design candidate molecules against KRAS, a target long regarded as "undruggable," screening at a scale no wet-lab campaign could reach
- Google DeepMind used AI to discover a previously unknown protein interaction critical to cancer cell survival – something traditional methods would have been virtually incapable of finding
- 2025 saw the largest number of AI-originated molecules entering clinical trials, concentrated in oncology, fibrosis, autoimmune diseases, and rare diseases

One case from August 2026 is worth taking apart on its own – because it shows both the power of this route and its limits at the same time.

On 19 August 2026, Merck and Moderna released topline results from INTerpath-001, the Phase 3 trial of the individualised cancer vaccine intismeran autogene (mRNA-4157): **1,137 patients** with completely resected stage IIB-IV cutaneous melanoma, randomised 2:1, with the primary endpoint of recurrence-free survival (RFS) and the key secondary endpoint of distant metastasis-free survival (DMFS) **both statistically significant and clinically meaningful**. Moderna closed that day at \$174.38 against a prior close of \$62.96 – **+176.97% in a single session** [140].

The technical point here is specific: this is not one drug given to every patient. It sequences **each patient's own tumour**, uses an algorithm to select and rank the neoantigens from that mutation profile most likely to be recognised by the immune system, and then synthesises an mRNA construct belonging to that one person – encoding up to **34 neoantigens** per dose. **"Algorithmic target selection" is not an assistive tool here; it is the precondition for the route existing at all**: without automated neoantigen selection and ranking, "make a separate vaccine for every individual" simply does not run as an engineering proposition [141].

But two things have to be kept apart, or this becomes hype:

- The Phase 3 release **stated only that the endpoints were met. It disclosed no hazard ratio, no confidence interval, no p-value, no event counts, and no follow-up duration** [140]. The widely circulated figures of "49% reduction in recurrence risk" and "59% reduction in distant metastasis risk" are **not** from this Phase 3 readout; they come from the earlier Phase 2b KEYNOTE-942 (n=157, median follow-up 60.3 months, RFS HR=0.51, DMFS HR=0.411) [141]. Attaching Phase 2b effect sizes to a Phase 3 result is a common misreading.
- Public material says only that the neoantigen sequences are "algorithmically derived." **It does not disclose model architecture or training data, and does not state whether deep learning or generative models were used** [141]. So one can say "computational methods are the precondition for this route." One cannot say "AI cured cancer."

Even discounted heavily, this remains the hardest single example supporting this section's argument: **once the marginal cost of screening and design falls far enough, "making one drug for one person" moves from economically impossible to economically possible**. Which is exactly what the next paragraph is about.

Rare disease patients may be the biggest beneficiaries. The figure usually cited for rare-disease patients worldwide is around 300 million (definitions differ between bodies, and estimates vary accordingly); most of these conditions attract no pharmaceutical investment because patient populations are too small and markets too limited. AI may lower the cost of target discovery and candidate screening, making it easier for some rare-disease programmes to **start**. Whether they actually reach market, and whether patients can afford them, still turns on trial recruitment, endpoints, regulatory approval, manufacturing cost, and payer reimbursement – **falling R&D cost is a necessary condition, not a sufficient one**.

17.2 Accessibility: The World Is Opening Up for People With Disabilities and the Elderly

This is a severely underestimated area.

For the visually impaired: Tools like Microsoft's Seeing AI use AI to describe the surrounding environment in real time – identifying objects, reading text, recognizing faces. A blind person can "see" a

restaurant menu, "read" the instructions on a medicine bottle, or "recognize" a friend walking toward them.

For the hearing impaired: AI real-time captioning lets deaf people "hear" conversations – not just video-call subtitles, but live transcription of face-to-face interactions.

For those with limited mobility: AI-powered smart homes allow wheelchair users and bedridden elderly to control lights, air conditioning, door locks, and curtains by voice – tasks that once required someone else's help can now be done independently.

For those with cognitive impairments/the elderly: AI companion robots can chat with dementia patients, conduct cognitive training, monitor mood changes, and remind them to take medication. South Korea has already deployed more than **5,400** AI companion robots across 115 local governments.

For these groups, AI is not "replacement" – it is "empowerment." It grants them independence and dignity they never had before. This is one of the least contested applications of AI, with the clearest positive externalities – but it is not free of cost either: privacy and surveillance, algorithmic misjudgment, machine companionship displacing human contact, device dependence, data security, and "the service exists but you cannot afford it" are all still open problems.

17.3 Education: The Best Tutor in History, for Free

AI may be the most powerful equalizer yet for access to educational resources.

What used to be required to access top-tier educational resources? A house in the right school district (millions of yuan), international school tuition (hundreds of thousands per year), one-on-one tutoring (hundreds per hour). Poor children and rich children were never on the same track from the starting line.

AI is changing this.

A rural child can now learn the same material as an urban child using AI – personalized, at their own pace. An AI tutor will not lose patience because you asked a "dumb question," will not give up on you for being a poor student, and will not stop teaching you because your family cannot afford the hourly rate.

Stanford's Data Ocean platform has already provided nearly **3,600** free AI medical education certifications across **93 countries** – for students in developing nations, this was previously unimaginable.

In Kenya and Nigeria, AI tools already support Swahili, Yoruba, and Amharic – farmers can interact with AI by voice on basic phones to get crop disease diagnoses.

In the 2024-2025 school year, **60%** of U.S. K-12 teachers were using AI tools, saving **6 hours** per week [101] (see Volume I, 2.7). Those 6 hours can be redirected to things AI cannot do – one-on-one attention for students, handling emotional issues, building personal trust.

17.4 One Person, One Army: The Barrier to Entrepreneurship and Creation Drops to Zero

This is AI's greatest gift to people with initiative.

Sam Altman has publicly predicted the arrival of the first "one-person billion-dollar company."

Whether that comes with a firm timetable is arguable, but the direction it points to is clear: once the capacity a single person can command is no longer bounded by headcount, "company size" and "output size" come apart.

This is not hyperbole. Take Midjourney: a private company that discloses no financials, and outside estimates disagree on both headcount and revenue – **but every set of figures lands revenue-per-employee in the millions of dollars a year**, an order of magnitude above a conventional software company. What if you compressed that further to one person?

AI lets a single individual do what used to require a team:

- **Programming:** Claude Code lets one person run 4-10 AI coding agents simultaneously, equivalent to a small development team
- **Design:** AI generates drafts; you provide aesthetic judgment and make the final selection
- **Marketing:** AI writes copy, generates ad creatives, analyzes data
- **Customer service:** AI handles 90% of routine inquiries; you handle the 10% that are complex
- **Finance/Legal:** AI handles basic compliance; you make strategic decisions

What does this mean? It means that building the first version of a product is no longer a game reserved for the wealthy. You do not need funding, employees, or an office. You need a good idea, basic judgment, and a computer with internet access.

The probability of success is still low – most startups fail, and AI will not change that. But the cost of failure approaches zero. Try an idea; if it doesn't work, discard it and try another. A single attempt used to cost millions; now it costs a few weeks and a few hundred dollars in API fees.

17.5 Everyday Improvements: Unglamorous but Real Benefits

Beyond these "big" benefits, AI is quietly improving everyday life in countless ways:

- **Language is no longer a barrier:** Real-time translation lets you communicate with anyone in any language. Traveling abroad, doing cross-border business, reading foreign materials – AI means "not speaking the language" is no longer a constraint
- **Self-driving reduces accidents:** roughly **1.16 million** people die in road crashes worldwide each year [135]. NHTSA's classic crash-causation survey assigned the "critical reason" to the driver in about **94%** of crashes [136] – that figure is where the ceiling on "how much autonomy could save" comes from. But be precise: 94% is an attribution, not a promise that a machine would have avoided all of them, and NHTSA itself cautions against reading it that way. The empirical record so far is limited to bounded mileage in specific cities (see the Waymo data at [129] in Volume I). Even partial delivery is on the order of hundreds of thousands of lives a year
- **Personal productivity:** AI helps you sort email, manage your calendar, summarize long documents, and draft replies – freeing you from repetitive mental labor
- **Fraud and security:** AI protects you without you even knowing – detecting anomalous bank transactions, identifying phishing emails, preventing identity theft
- **Independent living for the elderly:** Smart homes + AI voice assistants allow the elderly to live independently at home for longer, reducing dependence on children or caretakers

17.6 Shifting Perspective: From Threat to Tool

Summing up the benefits above, the core cognitive shift is:

AI is not just something that steals your job. It is also:

- Your personal doctor (at the screening level)
- Your all-subject tutor (free, patient, 24/7)
- Your startup co-founder (no salary required, no equity demanded)
- Your accessibility assistant (if you have any physical limitations)
- Your translator, secretary, researcher, and designer

The same AI is a threat to one person ("it can do my job") and a superpower for another ("it can help me do ten times as much"). Part of the difference lies in how you see it and how you use it – **but only part**. For people with skills, time, equipment, and room to fail, AI is an amplifier. For people already sitting in a job being replaced, with no cushion and no bargaining power, it is first of all a structural shock, and mindset has nothing to do with it. **How you use it matters, but it is no substitute for institutional protection** – the line drawn in 15.2 and 15.4.2 of this volume still holds here: some things an individual cannot do, and no amount of "look at it differently" changes that.

One designer worries AI will replace them – but another uses AI to generate 50 concepts a day, then applies their own aesthetic judgment to pick the best 3. The latter produces 10x more than the former.

One writer worries AI writes better than they do – but another uses AI for research, organizing material, and generating first drafts, then injects their own perspective and personality. The latter produces in one week what used to take a month.

The question is not "can you do what AI does" but "what can you plus AI do that was previously impossible?"

Seize these opportunities. Use AI to do what you always wanted to do but couldn't – learn a new skill, build a small product, help your community, improve your health. This is not the empty slogan of "embrace change." This is concrete action.

17.7 "Survive the Storm and Better Days Await" – The Power of This Narrative

Humans can endure enormous suffering, provided they know it has an end.

The British endured the Blitz during World War II because Churchill said "We shall never surrender" – not "everything will be fine," but "beyond the suffering lies victory."

Many who took their own lives during the Great Depression were not the poorest – they were those who believed "things will never get better." **Despair kills; poverty alone does not.**

What should the narrative of the AI era be?

Not "the apocalypse is here." Not "everything will be fine." Rather: "The next few years will be hard, and whether the hard part has an end depends on whether institutions catch up."

Say that precisely. There is exactly one conclusion the technological trend supports: **productive capacity**

will rise. It does not imply that ordinary people get a share. If distribution systems, public services, and the social safety net keep up, the world after AI could be markedly better. If they do not, **technological abundance can coexist perfectly well with ordinary poverty.** So what has to be survived is not a technology gap; it is the **period of institutional reconstruction.**

This is not blind optimism. It is a rational assessment based on technological trends:

- The cost of AI is falling exponentially – Sam Altman predicts a tenfold drop per year
- The cost of material production will approach zero
- The day when humans no longer need to work for food, clothing, housing, or basic healthcare is technologically achievable
- AI is already making healthcare, education, and accessibility services better, cheaper, and more widely available

The only question is **how long the transition lasts.** It could be 5 years; it could be 20. But the direction is clear.

Knowing the storm will pass and believing the storm is eternal produce entirely different behaviors.

The former makes you repair the ship, stock provisions, support each other, and even savor moments of calm amid the rain. The latter makes you give up.

This is not optimism. This is strategy.

17.8 Your Own Health: Living to 2035 May Be the Best Investment You Ever Make

As analyzed in Volume I, 9.2–9.3: age-reversal technology and AI medicine may achieve breakthroughs between 2035 and 2045. **Those who are alive at that point could gain decades of additional life.** The return on your current health investments may be the highest in history.

Things that cost nothing but pay off enormously:

- 150 minutes of exercise per week
- Adequate sleep (7-8 hours)
- Weight management
- Regular health checkups
- Maintaining social connections – and this one has hard data behind it: a meta-analysis covering 148 studies and 308,849 participants found that people with stronger social relationships had a **50%** greater likelihood of survival (OR = 1.50), an effect size in the same league as smoking, obesity, and other established mortality risk factors [137]

Conclusion · Crossing the Transition

There is a saying: the pessimist is right, the optimist moves forward.

By now the map and the list are both open on the table. The question is no longer whether it will hurt, but what is still in your hands when it does.

Here is the whole thing in three sentences.

First, there is good reason for optimism about the technological endgame — but it is not a promise that pays out on its own. If productivity keeps advancing along its current line, material abundance and a humanity released from most necessary labour are where the technology trend points. But technology only supplies the possibility. Whether it is delivered depends on distribution, on safety governance, and on how much strain a society can absorb — and what Volume I spent nine chapters establishing is that none of those three is guaranteed to keep up.

Second, the endgame does not arrive tomorrow. In between lies a transition that may run ten to twenty years — the technology has already turned over, the institutions have not. The gap between the production side and the distribution side is exactly where ordinary people stand exposed. That gap is what this book is really about.

Third, on that stretch of road immunity is not available. Only endurance is. Structural change at this scale is not something an individual can block. What Volume II offers is not armor but a handful of things that will get you through with fewer injuries:

- **Psychology first** — unbind identity, train resilience, widen your luck surface area, stay loose (Chapter XIII)
- **Read your assets correctly** — the HALO frame: things in the world of atoms resist better than things in the world of bits (Chapter XIV)
- **Cut leverage** — low debt is the largest safety cushion in a transition (Chapter XIV, 14.4)
- **Work in stages** — different windows need different strategies; the deadliest error is meeting the next phase with the last phase's playbook (Chapter XV)
- **Maintain your people** — "someone cares whether I exist" may be the single most valuable resource in a transition (Chapter XIII)
- **Stay healthy** — the longer the transition runs, the more your physical and mental condition becomes the underlying asset (Chapter XVII, 17.8)
- **Let your children off** — this is not "education stops mattering." It is: stop using the last era's scarcity logic to train them into tools fitted to a vanishing world. Curiosity, health, and relationships matter more than winning an exam that is being cancelled (Chapter XVI)
- **Use AI on offense** — it is not only a threat; it is a superpower. Use it to do what you could not do before (Chapter XVII)

One last thing.

Volume I closed on the hardest problem in the book: this analysis is reflexive — **every decision you make after reading it makes it more true.**

There is no exit from that. But there is a direction.

But that direction has to be stated precisely, or it collapses into a pep talk.

The collapse-direction reflexivity works because it **requires no coordination at all**: each person panic-sells, each firm freezes hiring, each household cashes out early, and the sum is systemic acceleration.

Deleveraging requires no coordination either. The collapse chain in Chapter VI of Volume I — pay cut → missed mortgage → impaired bank capital → credit contraction → further pay cuts — begins at household leverage. A person who brings their debt down is feeding the opposite quantity into the same chain. **Same mechanism, run backwards.**

This is not the optimism of belief-makes-it-so. It is the realism that **action changes your exposure.** It may not stop a systemic shock; it decides where a great many people are standing when one arrives.

That is the only honest use of reflexivity: **it does not promise the rain will stop. It only decides where you are standing when it falls.**

So: hold what can be held, put down what should be put down. Reduce risk, but do not stop acting. When the anxiety comes, go for a walk, cook something, see a friend. When nothing feels like it means anything, go do one thing that makes you happy — not for anyone else, for yourself.

The storm can pass — not because people naturally get better, but because technology is manufacturing the capacity to solve these problems. The real question is whether humanity can hold out until institutions learn to use that capacity.

And "holding out" is your part of it.

The pessimist draws the map; the optimist takes the steps. The map tells you where the cliff is. The steps decide how you go around it.

Pessimism by itself is useless. You still have to keep walking.

Stay alive. Cross the transition. Then go see what is on the other side.


About This Article

Written by Tony Seah, based on logical extrapolation from public data and research reports. All charts are original hand-drawn SVGs.

Contact: tony.seah@tokligence.ai

Discussion

Join the Telegram group to discuss the ideas in this book:



If you found this book valuable, consider supporting the author:

International:

- GitHub Sponsors: github.com/sponsors/pkusnail
- Buy Me a Coffee: buymeacoffee.com/pkusnail

Reprint notice: Feel free to repost and share this book, but please retain author attribution and support links.

Source / Project: github.com/pkusnail/GPUs-Wombs-and-Money-Printers

References

- [1] Citrini Research & Alap Shah, “The Consequences of Abundant Intelligence,” February 22, 2026. <https://www.citriniresearch.com/p/2028gic>
- [2] Wu Jun, *Tides of Technology (Lang Chao Zhi Dian)*, 4th ed., Posts & Telecom Press, 2019. <https://book.douban.com/subject/33474750/>
- [3] Krizhevsky, A., Sutskever, I., & Hinton, G. E., “ImageNet Classification with Deep Convolutional Neural Networks,” NeurIPS 2012. <https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks>
- [4] Pew Research Center, “5 facts about global fertility trends,” August 2025. <https://www.pewresearch.org/short-reads/2025/08/15/5-facts-about-global-fertility-trends/>
- [5] Statista, “Countries with the lowest fertility rates 2025.” <https://www.statista.com/statistics/268083/countries-with-the-lowest-fertility-rates/>
- [6] United Nations, “World Fertility 2024,” 2025. <https://www.un.org/development/desa/pd/content/world-fertility-2024>
- [7] The Lancet / IHME, “Dramatic declines in global fertility rates set to transform global population patterns by 2100.” <https://www.healthdata.org/news-events/newsroom/news-releases/lancet-dramatic-declines-global-fertility-rates-set-transform>

- [8] WardheerNews, “South Korea — How a \$280 Billion Baby Plan Failed,” 2025. <https://wardheernews.com/south-korea-how-a-280-billion-baby-plan-failed-and-what-it-says-about-modern-society/>
- [9] ScienceDirect, “Childbearing and the distribution of the reservation price of fertility: The case of the Korean baby bonus program.” <https://www.sciencedirect.com/science/article/abs/pii/S104900782100124X>
- [10] MIT Technology Review, “Inside Japan’s long experiment in automating elder care,” January 2023. <https://www.technologyreview.com/2023/01/09/1065135/japan-automating-eldercare-robots/>
- [11] CNN, “South Korea’s aging population finds comfort in ‘robo-grandma’ dolls.” <https://www.cnn.com/health/asias-aging-population-finds-comfort-care-in-robo-grandma-dolls-wellness-spc>
- [12] Bloomberg, “Pentagon Pressures Anthropic to Drop AI Guardrails in Military Standoff,” February 26, 2026. <https://www.bloomberg.com/news/features/2026-02-26/pentagon-presses-anthropic-to-drop-ai-guardrails-in-military-standoff>
- [13] NBC News, “Pentagon gives AI firm ultimatum: lift military limits by Friday or lose \$200M deal.” <https://www.nbcnews.com/meet-the-press/video/pentagon-gives-anthropic-an-ultimatum-amid-fight-over-military-ai-guardrails-258370117539>
- [14] Fortune, “The AI arms race could lead to human extinction, top researcher warns,” February 2026. <https://fortune.com/2026/02/18/big-tech-russian-roulette-ai-race-humanity-extinction/>
- [15] SF Examiner, “SF office-vacancy rate heading for biggest drop since 2015.” https://www.sfexaminer.com/news/business/sfs-office-vacancy-rate-heading-for-biggest-drop-since-2015/article_dbb23edb-a7fc-4ec5-bb70-81fbceb988bc.html
- [16] Congress.gov / CRS, “Commercial Real Estate and the Banking Sector,” R48175. <https://www.congress.gov/crs-product/R48175>
- [17] Wikipedia, “Decline of Detroit.” https://en.wikipedia.org/wiki/Decline_of_Detroit
- [18] Cleveland Fed, “Urban Decline in Rust-belt Cities,” 2013. <https://www.clevelandfed.org/publications/economic-commentary/2013/ec-201306-urban-decline-in-rust-belt-cities>
- [19] Vision Times, “Is China’s Northeast Disappearing? Ghost Cities Spread,” November 2025. <https://www.visiontimes.com/2025/11/26/is-chinas-northeast-disappearing-ghost-cities-spread-as-housing-prices-collapse.html>
- [20] Nature Cities, “The decreasing housing utilization efficiency in China’s cities,” 2024. <https://www.nature.com/articles/s44284-024-00177-8>
- [21] Visual Capitalist, “The Decline of Fertility Rates in OECD Countries (1950–2025).” <https://www.visualcapitalist.com/decline-of-fertility-rates-in-oecd-countries-1950-2025/>
- [22] Carnegie Endowment, “The AI Governance Arms Race,” October 2024. <https://carnegieendowment.org/research/2024/10/the-ai-governance-arms-race-from-summit-pageantry-to-progress>
- [23] Computer Weekly, “MPs warned of AI arms race to the bottom.” <https://www.computerweekly.com/news/365529793/MPs-warned-of-AI-arms-race-to-the-bottom>
- [24] CNN / China Ministry of Civil Affairs, “Fewer people than ever before are marrying in China. But divorces

- are up,” February 2025. <https://www.cnn.com/2025/02/10/china/china-marriage-registrations-record-low-2024-intl-hnk>
- [25] CDC/NCHS, “FastStats — Marriage and Divorce,” 2023. <https://www.cdc.gov/nchs/fastats/marriage-divorce.htm>
- [26] OECD, “Marriage and Divorce: Society at a Glance 2024.” https://www.oecd.org/en/publications/society-at-a-glance-2024_918d8db3-en/full-report/marriage-and-divorce_63dd0a7d.html
- [27] YuWa Population Research / CNN, “China is one of world’s most expensive places to raise children,” February 2024. <https://www.cnn.com/2024/02/22/asia/china-cost-of-raising-children-intl-hnk>
- [28] Merton, Robert K., “The Self-Fulfilling Prophecy,” *The Antioch Review*, Vol. 8, No. 2, 1948. <https://www.jstor.org/stable/4609267>
- [29] Soros, George, “The Alchemy of Finance,” Simon & Schuster, 1987 (Chinese edition: Jinrong Lianjinshu). <https://www.wiley.com/en-us/The+Alchemy+of+Finance,+2nd+Edition-p-9780471445494>
- [30] Finance Magnates / KED Global, “South Korea’s crypto market engages 30% of population,” 2024; PANews, “Why South Korea has become one of the hottest crypto markets in the world.” <https://www.financemagnates.com/cryptocurrency/south-koreas-102-trillion-won-crypto-market-engages-30-of-population/>
- [31] The Hill, “The US is headed for mass unemployment, and no one is prepared,” 2026; InvestorPlace, “AI Job Loss Is Accelerating – and Washington Won’t Stop It.” <https://thehill.com/opinion/finance/5713876-ai-displacement-and-ubi/>
- [32] OMFIF, “Fiscal exceptionalism and the US-Europe growth gap,” January 2026. <https://www.omfif.org/2026/01/fiscal-exceptionalism-and-the-us-europe-growth-gap/>
- [33] World Economic Forum / Kela, “The results of Finland’s basic income experiment are in,” 2019. <https://www.weforum.org/stories/2019/02/the-results-finlands-universal-basic-income-experiment-are-in-is-it-working/>
- [34] Sam Altman / Futurism, “Sam Altman Says AI Will Cause Massive Deflation, Making Money Worth Vastly More.” <https://futurism.com/future-society/sam-altman-ai-deflation>
- [35] Trading Economics / Japan CPI, “Japan Inflation Rate,” January 2026. <https://tradingeconomics.com/japan/inflation-cpi>
- [36] Pandaily, “China Produced Nearly 90% of Humanoid Robots Sold Globally in 2025,” February 2026. <https://pandaily.com/china-produced-nearly-90-of-humanoid-robots-sold-globally-in-2025/>
- [37] Reuters / NextBigWhat, “Chinese robot makers outpace Tesla in humanoid robot race,” February 2026. <https://www.reuters.com/technology/>
- [38] IEEE Spectrum, “What’s Holding Humanoid Robots Back,” February 2026. <https://spectrum.ieee.org/humanoid-robot>
- [39] Wiley / American Journal of Hematology, “CML 2025 Update on Diagnosis, Therapy, and Monitoring.” <https://onlinelibrary.wiley.com/doi/full/10.1002/ajh.27443>
- [40] NEJM, “Trastuzumab Deruxtecan plus Pertuzumab in HER2-Positive Metastatic Breast Cancer,” 2025. <https://www.nejm.org/doi/full/10.1056/NEJMoa2508668>
- [41] PharmaBoard Room, “Hong Kong Biotech IPOs in 2025.” <https://pharmaboardroom.com/fact>

s/hong-kong-biotech-ipos-in-2025/

[42] MIT Technology Review, “The first human test of a rejuvenation method will begin shortly,” January 2026. <https://www.technologyreview.com/2026/01/27/1131796/the-first-human-test-of-a-rejuvenation-method-will-begin-shortly/>

[43] PMC / Probabilistic Forecasting of Maximum Human Lifespan by 2100. <https://pmc.ncbi.nlm.nih.gov/articles/PMC11448547/>

[44] SSA, “2025 OASDI Trustees Report Summary,” June 2025. <https://www.ssa.gov/oact/trsum/>

[45] The Diplomat, “As China Ages, a Pension Crisis Looms,” February 2026. <https://thediplomat.com/2026/02/as-china-ages-a-pension-crisis-looms/>

[46] HEconomist, “Japan’s Pension System: A Fragile Equilibrium,” May 2025. <https://heconomist.com/2025/05/05/japans-pension-system-a-fragile-equilibrium-between-tradition-and-contemporary-challenges/>

[47] IMF, “Global Financial Stability Report — The Financial Impact of Longevity Risk,” April 2012. <https://www.imf.org/en/blogs/articles/2012/04/11/seven-billion-reasons-to-worry-the-financial-impact-of-living-longer>

[48] ASH Clinical News, “CAR-T Cell Therapies Predicted to Cost More Than \$1 Million.” <https://ashpublications.org/ashclinicalnews/news/3469/CAR-T-Cell-Therapies-Predicted-to-Cost-More-Than-1>

[49] S&P Global, “Data center grid power demand to rise 22% in 2025, nearly triple by 2030,” 2025; Axios, “AI data center boom: projects vs numbers,” February 2026. <https://www.spglobal.com/energy/en/news-research/latest-news/electric-power/101425-data-center-grid-power-demand-to-rise-22-in-2025-nearly-triple-by-2030>

[50] Tom’s Hardware / TSMC, “Advanced node capacity falls short of AI demand,” 2025; Congress.gov CRS R48642 (SMIC 5nm analysis). <https://www.tomshardware.com/tech-industry/semiconductors/tsmc-csays-advanced-node-capacity-falls-short-of-ai-demand>

[51] Bain & Company, “Humanoid Robots: From Demos to Deployment,” Technology Report 2025. <https://www.bain.com/insights/humanoid-robots-from-demos-to-deployment-technology-report-2025/>

[52] MIT Technology Review, “Why Humanoids Need Their Own Safety Rules,” June 2025; OnLabor, “Amazon’s Approach to Robotics Is Seriously Injuring Warehouse Workers.” <https://www.technologyreview.com/2025/06/11/1118519/humanoids-safety-rules/>

[53] Medium / Markus Brinsa, “Hallucination Rates in 2025”; ISACA, “Avoiding AI Pitfalls in 2026: Lessons from Top 2025 Incidents.” <https://www.isaca.org/resources/news-and-trends/isaca-now-blog/2025/avoiding-ai-pitfalls-in-2026-lessons-learned-from-top-2025-incidents>

[54] McKinsey, “Future of Autonomous Vehicles,” 2025; Rodney Brooks, “Predictions Scorecard 2026,” January 2026. <https://rodneybrooks.com/predictions-scorecard-2026-january-01/>

[55] Gartner, “Hype Cycle for Artificial Intelligence 2025”; MIT / Active Logic, “Where is AI on Gartner’s 2025 Hype Cycle.” <https://www.gartner.com/en/articles/hype-cycle-for-artificial-intelligence>

[56] European Commission, “AI Omnibus enters into force,” July 27, 2026 (Regulation (EU) 2026/1744: Annex III high-risk obligations postponed to December 2, 2027; Annex I to August 2, 2028); Latham & Watkins, “China Introduces Rules for AI Companion and Emotional Interaction Services,” July 8, 2026. <https://digital-strategy.ec.europa.eu/en/news/ai-omnibus-enters-force>

- [57] Baker McKenzie, “Navigating Labor’s Response to AI,” June 2025; Brookings, “Hollywood Writers Strike and AI.” <https://www.bakermckenzie.com/en/insight/publications/2025/06/navigating-labors-response-to-ai>
- [58] Stanford CodeX, “How AI Agents Are Rewriting Rules of Engagement,” January 2025; DLA Piper, “Rise of Agentic AI Legal Risks.” <https://law.stanford.edu/2025/01/14/from-fine-print-to-machinese-code-how-ai-agents-are-rewriting-the-rules-of-engagement/>
- [59] Epoch AI, “AI Supercomputers Performance Share by Country,” 2025; GeoCoded, “State of Global AI Compute 2025.” <https://epoch.ai/data-insights/ai-supercomputers-performance-share-by-country>
- [60] Inbolt, “Industrial Robot Cost Decline,” 2025; EY, “Three Tailwinds for Robotics Adoption,” 2024. <https://www.inbolt.com/resources/industrial-robot-cost-decline>
- [61] Reshoring Initiative, “2024 Annual Report,” 2025; Automate.org, “Reshoring and Nearshoring Trends.” https://reshorenw.org/content/pdf/2024-1Q2025_RI_DATA_Report.pdf
- [62] UNCTAD, “Technology and Innovation Report 2025”; ODI, “The AI Time Bomb: 2.5 Million Jobs at Risk in Kenya”; ILO / IMF, “AI and the Philippine Labor Market.” <https://unctad.org/publication/technology-and-innovation-report-2025>
- [63] NBER, “India’s De-Industrialization Under British Rule,” Working Paper 10586. <https://www.nber.org/papers/w10586>
- [64] Atlantic Council, “Will Data & AI Cripple or Leapfrog Developing Nations?”; Connecting Africa, “World Bank Chief on AI Leapfrogging.” <https://www.connectingafrica.com/ai/world-bank-chief-not-certain-developing-nations-can-leapfrog-with-ai>
- [65] CDO Magazine, “Pentagon Seeks \$13.4B for AI and Autonomy in FY2026 Budget Request,” 2025; CNAS, “The Role of AI in Russia’s Confrontation with the West.” <https://www.cdomagazine.tech/us-federal-news-bureau/pentagon-seeks-13-4-bn-for-ai-and-autonomy-fy-2026-budget-request>
- [66] 972 Magazine, “Lavender: The AI machine directing Israel’s bombing spree in Gaza,” April 2024. <https://www.972mag.com/lavender-ai-israeli-army-gaza/>
- [67] Atlantic Council, “Missiles, AI, and drone swarms: Ukraine’s 2025 defense tech priorities”; MIT Technology Review, “Autonomous warfare in Europe,” January 2026. <https://www.technologyreview.com/2026/01/06/1129737/autonomous-warfare-europe-drones-defense-automated-kill-chains/>
- [68] Anthropic, “Disrupting AI-powered espionage,” November 2025; Arms Control Association, “AI-Nuclear Transparency,” December 2025. <https://www.anthropic.com/news/disrupting-AI-espionage>
- [69] CFR Education, “Industrialization and Imperialism”; Michigan Journal of Economics, “Economics of Warfare,” 2025; OMFIF, “Putin and Russia’s Ruined Economic Future,” January 2026. <https://www.omfif.org/2026/01/putin-is-wary-of-facing-russias-ruined-economic-future/>
- [70] Maplecroft, “Escalating unrest, polarisation set stage for disruptive 2026”; Pew Research, “Political Polarization.” <https://www.maplecroft.com/solutions/consulting/political-risk/insights/escalating-unrest-polarisation-economic-woes-set-stage-for-disruptive-2026/>
- [71] Foreign Policy Analytics, “AI and Critical Minerals,” July 2025; NBC News, “Greenland’s critical minerals a tempting target”; American Security Project, “Rare Earths, AI, and Strategic Competition.” <https://fpanalytics.foreignpolicy.com/2025/07/18/artificial-intelligence-critical-minerals-supply-chains/>

- [72] Allison, Graham, “Destined for War: Can America and China Escape Thucydides’s Trap?” Houghton Mifflin Harcourt, 2017; Harvard Kennedy School, “Thucydides’ Trap.” <https://www.hks.harvard.edu/faculty-research/policy-topics/international-relations-security/thucydides-trap-peace-and-rivalry>
- [73] SpaceX, “Starship” mission updates, 2024–2026; Payload Space, “Starship Economics: The \$10/kg LEO Target,” 2025. <https://payloadspace.com/starship-launch-cost-analysis/>
- [74] Lumen Orbit, “Starcloud: First GPU in Orbit,” November 2025; Google, “Project Suncatcher: Space-Based Solar for AI Data Centers,” 2025; TechCrunch, “The Race to Build Data Centers in Space,” 2025. <https://techcrunch.com/2025/11/space-data-centers-lumen-orbit/>
- [75] Bloomberg, “Musk’s SpaceX Combines With xAI at \$1.25 Trillion Valuation,” February 2, 2026; SpaceX Nasdaq IPO (SPCX): first trade June 12, 2026, ~\$1.77 trillion valuation, \$75B raised—the largest IPO by proceeds on record; Jonathan McDowell, Starlink Statistics (as of August 18, 2026: 10,979 in orbit). <https://planet4589.org/space/con/star/stats.html>
- [76] NASA, “Mars Facts”; Rees, Martin, “The End of Astronauts,” Harvard University Press, 2024; Space.com, “Musk shifts Mars-to-Moon focus,” February 2026. <https://mars.nasa.gov/all-about-mars/facts/>
- [77] NASA JPL, “Perseverance Rover Completes First AI-Planned Drive on Mars,” December 2025. <https://www.jpl.nasa.gov/news/perseverance-ai-autonomous-driving>
- [78] Rees, Martin, “The End of Astronauts,” Harvard University Press, 2024; Von Neumann, John, “Theory of Self-Reproducing Automata,” 1966. <https://press.uchicago.edu/ucp/books/book/chicago/T/bo5780028.html>
- [79] Spielberg, Steven, “A.I. Artificial Intelligence,” Warner Bros., 2001; Scott, Ridley, “Prometheus,” 20th Century Fox, 2012.
- [80] SpaceNews, “China’s Landspace recovers booster with second orbital launch of Zhuque-3 rocket,” August 18, 2026; India Today, “Zhuque-3 landing leg catches fire after touchdown,” August 19, 2026; SpaceNews, “China’s commercial Tianlong-3 rocket fails on debut launch,” April 3, 2026. <https://spacenews.com/chinas-landspace-recovers-booster-with-second-orbital-launch-of-zhuque-3-rocket/>
- [81] Hosseini Maasoum & Lichtinger, “Generative AI as Seniority-Biased Technological Change,” Harvard/SSRN, 2025; SCMP, “Youth unemployment eases to 16.9%,” 2025; Caixin, “Youth Unemployment Surge,” September 2025; Foreign Policy, “China’s AI Planners Are Fearful of Job Losses,” November 2025. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5425555
- [82] SignalFire, “State of Talent Report 2025”; CNBC, “AI is not just ending entry-level jobs,” September 2025; Revelio Labs, “Entry-Level Job Postings Decline,” 2025; LeadDev Engineering Survey 2025; IntuitionLabs, “AI’s Impact on Graduate Jobs 2025.” <https://www.signalfire.com/blog/signalfire-state-of-talent-report-2025>
- [83] Global Policy Solutions, “Stick Shift: Autonomous Vehicles, Driving Jobs, and the Future of Work”; Medium/Ben Kitchen, “Transitioning from Horses to Cars: Resource Consumption”; Resilience.org, “The Big Shift Last Time: From Horse Dung to Car Smog,” 2013; Victoria Transport Policy Institute, “Autonomous Vehicle Implementation Predictions,” 2025. <https://globalpolicysolutions.org/report/stick-shift-autonomous-vehicles-driving-jobs-and-the-future-of-work/>
- [84] Tesla, “FSD Safety Hub,” 2026; TechCrunch, “Tesla launches robotaxi rides in Austin with no human safety driver,” January 22, 2026; Electrek, “Tesla starts robotaxi rides without safety monitor,” January 2026. <https://www.tesla.com/fsd/safety>
- [85] Waymo, “2025 Year in Review,” December 2025; CNBC, “Waymo 450,000 weekly rides,” December 2025;

CNBC, “China’s Baidu robotaxis match Waymo’s ride numbers,” November 2025; Pony.ai, “Shenzhen citywide driverless permit,” October 2025. <https://waymo.com/blog/2025/12/2025-year-in-review>

[86] CleanTechnica, “BYD God’s Eye Brings ADAS to the Masses,” February 2025; SCMP, “Chinese EV makers outpace Tesla in autonomous-driving race,” 2025. <https://cleantechnica.com/2025/02/17/byd-gods-eye-brings-adas-to-the-masses/>

[87] Yicai Global, “China’s 30 Million Truck Drivers,” 2024; Beijing Scroll, “Inside China’s 7.48 Million Ride-Hailing Drivers,” 2024; Goldman Sachs, “Self-Driving Cars Could Steal 300,000 Jobs a Year,” 2017; International Transport Forum, “Driverless Trucks and Job Losses,” 2017; IRU, “Global Truck Driver Shortage Report,” 2024. <https://www.yicaiglobal.com/news/china-saturated-trucking-market-has-left-30-million-drivers-in-debt-report-says>

[88] Autor, David et al., “When Work Disappears: Manufacturing Decline and the Falling Marriage-Market Value of Young Men,” American Economic Review: Insights, 2019; National Institute of Population and Social Security Research (Japan), “Population Statistics of Japan,” 2023; National Institute of Population and Social Security Research (Japan), “Lifetime Never-Married Rate for Men,” 2020. <https://www.aeaweb.org/articles?id=10.1257/aeri.20180010>

[89] 36Kr, “China’s AI Virtual Companion Users Surpass 30 Million,” 2024; The Economist, “AI girlfriends are ruining a generation of men,” 2025; Reuters, “Japan’s AI relationship industry,” 2024; user data from platforms including Character.ai and Replika. <https://www.economist.com/asia/2025/01/09/ai-girlfriends-are-ruining-a-generation-of-men>

[90] BLS, “Employment Situation Summary,” January 2026; Yale Budget Lab & Brookings, “Evaluating the Impact of AI on the Labor Market,” October 2025; Federal Reserve Governor Barr, “What Will AI Mean for the Labor Market and the Economy?” February 17, 2026; ILO, “World Employment and Social Outlook Trends 2026.” <https://www.federalreserve.gov/newsevents/speech/barr20260217a.htm>

[91] McKinsey, “The State of AI,” November 2025; Gartner, “40% of Enterprise Apps Will Feature Task-Specific AI Agents by 2026,” August 2025; Gartner, “Over 40% of Agentic AI Projects Will Be Canceled by End of 2027,” June 2025; WEF, “Future of Jobs Report 2025.” <https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai>

[92] Cengage Group, “2025 Graduate Employability Report,” 2025; The EDU Ledger, “Graduate Job Market Hits Five-Year Low as Skills Gap Widens,” 2025; HR Dive, “Over half of hiring managers say recent grads are unprepared for the workforce,” 2025. <https://www.cengagegroup.com/news/press-releases/2025/cengage-group-2025-employability-report/>

[93] Pew Research Center, “Is College Worth It?” May 2024; CNBC, “The median return on investment for a college degree,” April 2025; Georgetown CEW, “ROI of College 2025.” <https://www.pewresearch.org/social-trends/2024/05/23/is-college-worth-it-2/>

[94] TechCrunch, “The great computer science exodus (and where students are going instead),” February 2026; CRA CERP Pulse Survey, “A Snapshot of 2025 Undergraduate Computing Enrollment Patterns,” October 2025. <https://techcrunch.com/2026/02/15/the-great-computer-science-exodus-and-where-students-are-going-instead/>

[95] BestColleges, “Colleges That Have Closed: A Running List,” 2025; Bryan Alexander, “Academic closures, mergers, and cuts heading into fall 2025”; Federal Reserve Bank of Philadelphia, higher education financial stability projections. <https://www.bestcolleges.com/research/closed-colleges-list-statistics-major-closures/>

[96] RBA, “International Students and the Australian Economy,” July 2025; NAFSA, “International Students Contributed \$43 Billion to US Economy 2024–2025”; ICEF Monitor, “UK international student numbers fall for second year,” January 2026; ApplyBoard, “Impacts of Canada’s 2024 International Student Cap,” 2025. <https://>

[//www.rba.gov.au/publications/bulletin/2025/jul/international-students-and-the-australian-economy.html](https://www.rba.gov.au/publications/bulletin/2025/jul/international-students-and-the-australian-economy.html)

[97] Richmond Fed, “Workforce Pell Is (Finally) Law. Now What?” 2025; HCM Strategists, “New Insights for 2025: How States Are Funding and Expanding Short-Term Credential Pathways”; RAND, “Should the U.S. Lean into Short-Term, Nondegree Workforce Education and Training?” June 2025. https://www.richmondfed.org/region_communities/regional_data_analysis/community_college_survey/community_college_insights/2025/workforce_pell_finally_law_now_what

[98] Coursera, “2025 Micro-Credentials Impact Report”; Fortune, “Microcredentials are the hottest hiring trend for 2026,” December 2025. <https://www.coursera.org/enterprise/resources/ebooks/micro-credentials-report-2025>

[99] UK Government, “Lifelong Learning Entitlement: Overview,” 2025. <https://www.gov.uk/government/publications/lifelong-learning-entitlement-lle-overview/lifelong-learning-entitlement-overview>

[100] Worlddidac, “China’s Vocational Skills Upgrade 2025–2027: What It Means for Education, Industry, and Global Cooperation,” 2025. <https://worlddidac.org/news/chinas-vocational-skills-upgrade-2025-2027-what-it-means-for-education-industry-and-global-cooperation/>

[101] Stanford Report, “How is ChatGPT impacting schools, really?” July 2025; NBC News, “AI has become the norm for students. Teachers are playing catch-up,” 2025. <https://news.stanford.edu/stories/2025/07/chatgpt-open-ai-impact-schools-education-learning-data-research>

[102] EdTech Hub, “10 Scenarios for Education in 2035,” 2025; WEF, “The future of learning: AI is revolutionizing education 4.0,” April 2024. <https://edtechhub.org/evidence/10-scenarios-for-education-in-2035/>

[103] Luo, F. et al., “Impact of Business Cycles on US Suicide Rates, 1928–2007,” *American Journal of Public Health*, 2011; Tapia Granados, J.A., “Recessions and Mortality in the United States, 1920–1940,” *PNAS*, 2009. <https://pmc.ncbi.nlm.nih.gov/articles/PMC3134519/>

[104] Japanese Cabinet Office, Hikikomori surveys 2019/2023; Tamura et al., “Hikikomori prevalence in Japan,” *Psychiatry and Clinical Neurosciences*, 2025; Nippon.com, “Japan’s Hikikomori Population Tops 1.46 Million.” <https://www.nippon.com/en/japan-data/h01341/>

[105] Character.AI usage data via Business of Apps & DemandSage, 2025; Replika via ElectroIQ, 2025; Grand View Research, “AI Companion Market Size Report,” 2024 (China: \$2.04B revenue, CAGR 35.4%). <https://www.businessofapps.com/data/character-ai-statistics/>

[106] Chetty, R. et al., “The Association Between Income and Life Expectancy in the United States, 2001–2014,” *JAMA*, 2016; The Lancet, “Life expectancy inequalities in the US, 2000–2019,” 2024. <https://jamanetwork.com/journals/jama/fullarticle/2513561>

[107] Global Times, “AI resurrection of deceased relatives sparks debate in China,” March 2024; Euronews, “Digital afterlife: StoryFile, HereAfter AI,” 2024; Marina Smith StoryFile funeral case, 2022. <https://www.globaltimes.cn/page/202403/1309456.shtml>

[108] Nielsen, “The Gauge: Streaming surpasses 44.8% of US TV viewing,” May 2025; NPR, “The Monoculture Myth,” 2024. <https://www.nielsen.com/insights/2025/the-gauge/>

[109] Hubert et al., “GPT-4 outperforms humans on divergent thinking,” *Scientific Reports / Nature*, 2023; Nave & Terwiesch, “AI and Idea Diversity,” Wharton, 2024; “Creative scar effect,” *Technological Forecasting and Social Change*, 2025. <https://www.nature.com/articles/s41598-023-40858-3>

[110] RAND, “Dissecting America’s AI Action Plan,” 2025; Anthropic Biorisk Evaluations, 2025; RAND Europe

AI-enabled Biological Tools Index, 2024; Paris Summit AI Safety Report, 2025. <https://www.rand.org/pubs/commentary/2025/08/dissecting-americas-ai-action-plan-a-primer-for-biosecurity.html>

[111] NK News, North Korea gene-editing bioweapons warning, April 2024; China MSS ethnic genetic weapons warning, 2023; DARPA Safe Genes Program (\$65M). <https://link.springer.com/article/10.1007/s00146-025-02576-4>

[112] Small Wars Journal, “AI Use in Terrorist Plots Surges in 2025,” December 2025; RAND, “The WMD AI Security Gap,” February 2026. <https://smallwarsjournal.com/2025/12/24/ai-use-in-terrorist-plots-and-attacks-surges-in-2025/>

[113] IGI CRISPR Clinical Trials 2025 Update; Casgevy \$2.2M/treatment; Neuralink 12 recipients; He Jiankui Hainan relocation 2024; 75/96 countries prohibit heritable editing. <https://innovativegenomics.org/news/crispr-clinical-trials-2025/>

[114] Rueda, “Genetic enhancement, human extinction,” *Bioethics*, 2024; Kurzweil, *The Singularity Is Nearer*, 2024; Harari, *Homo Deus*, 2016. <https://onlinelibrary.wiley.com/doi/10.1111/bioe.13085>

[115] Adamala, K. P., Agashe, D., Belkaid, Y., et al. (38 signatories, including Nobel laureates), “Confronting risks of mirror life,” *Science*, 386(6728), 1351–1353, December 12, 2024; Mirror Biology Dialogues Fund, “Technical Report on Mirror Bacteria: Feasibility and Risks,” Stanford Digital Repository, December 2024. <https://doi.org/10.1126/science.ads9158>

[116] Meinke, A., Schoen, B., Scheurer, J., et al. (Apollo Research), “Frontier Models are Capable of In-context Scheming,” arXiv:2412.04984, December 2024 (models evaluated: OpenAI o1, Claude 3.5 Sonnet, Claude 3 Opus, Gemini 1.5 Pro, Llama 3.1 405B). <https://arxiv.org/abs/2412.04984>

[117] Hubinger, E., et al. (Anthropic), “Sleeper Agents: Training Deceptive LLMs that Persist Through Safety Training,” arXiv:2401.05566, January 10, 2024. <https://arxiv.org/abs/2401.05566>

[118] NBC News, “China Evergrande to be delisted from Hong Kong stock exchange following debt woes,” August 12, 2025; Caixin Global, “Country Garden Wins Creditor Backing for \$17.7 Billion Debt Restructuring,” November 7, 2025; Associated Press, “China Vanke’s near-default exposes fragility of the faltering recovery in the property industry,” December 30, 2025 (citing Fitch); Xinhua, “Verdict delivered in the Xu Jiayin case,” August 20, 2026; National Bureau of Statistics of China, “National Real Estate Market, January–June 2026,” July 15, 2026. <https://www.nbcnews.com/world/china/china-evergrande-delisted-hong-kong-stock-exchange-debt-woes-rcna224676>

[119] The Economist, “China’s mortgage boycotts are a symptom of a broader crisis,” August 9, 2022 (320 projects, 100+ cities); CNBC, “China’s unfinished property projects are 20 times the size of Country Garden,” November 15, 2023 (Nomura: ~20 million pre-sold units unbuilt at end-2022, ~¥3.2T to complete); CNN, “China is trying to end its ‘epic’ property crisis,” May 21, 2024 (Goldman Sachs: over ¥7T); Caixin, “Whitelist project loan approvals to exceed ¥4 trillion by year-end,” October 17, 2024. <https://www.cnbc.com/2023/11/15/chinas-unfinished-property-projects-are-20-times-the-size-of-country-garden.html>

[120] National Bureau of Statistics of China, “Sales Prices of Residential Buildings in 70 Medium and Large-Sized Cities,” February and July 2026; NBS, “National Real Estate Market 2025,” January 19, 2026; Fitch Ratings, “China’s New-Home Sales to Fall 11%–13% in 2026,” June 18, 2026. https://www.stats.gov.cn/sj/zxfbhjd/202608/t20260817_1965046.html

[121] People’s Bank of China, “Report on Loan Allocation by Financial Institutions,” quarterly (residential mortgages ¥36.29T at end-2026Q2, –3.8% YoY); NIFD, “China Macro Leverage Ratio Quarterly,” 2026Q2 (household leverage 57.7%); CICC (Lin Yingqi), “How Much Do Mortgage Prepayments Matter?” February 2023. <https://www.cls.cn/detail/2271293>

[122] Federal Reserve, FOMC Minutes, March 17–18, June 16–17 and July 28–29, 2026 (AI-and-employment

language weakening meeting by meeting; the July attendance list names Kevin Warsh as Chairman and Jerome Powell as a member); Kevin Warsh, Semiannual Monetary Policy Testimony, July 14, 2026; Federal Reserve Bank of Richmond, Economic Brief No. 26-26, “Worker Types, AI Exposure and the Recent Decline in Job-Finding Rates,” August 2026; Bank for International Settlements, *Annual Economic Report 2026*, Box C. <https://www.federalreserve.gov/monetarypolicy/fomcminutes20260729.htm>

[123] POWER Magazine, “Abbott Orders Full Audit of Texas Data Center Interconnection Queue, Threatens to Deny Grid Access,” August 4, 2026 (ERCOT large-load queue 474.7 GW in June 2026, of which 420.8 GW / 90.2% is data centers; Texas all-time peak demand 91,089 MW on July 22, 2026); BloombergNEF estimate that ~49.8 GW of data-center load will be delayed; Nvidia Vera Rubin entered full production July 21, 2026. <https://www.powermag.com/abbott-orders-full-audit-of-texas-data-center-interconnection-queue-threatens-to-deny-grid-access/>

[124] World Economic Forum, *Future of Jobs Report 2025* (by 2030: -92M / +170M / net +78M); McKinsey Global Institute (about one third of U.S. jobs today did not exist 25 years ago); American Enterprise Institute and IMF *Finance & Development* analyses of ATMs and bank-teller employment (tellers per branch 20→13, branches +43%, total U.S. tellers 300k→600k).

[125] SoloFounders.com, share of solo founders (22% in 2015 → 28% in 2020 → 38% in 2024 → 36.3% in 2025); CNBC / Y Combinator (YC Winter 2025 batch aggregate growth ~10%/week); Quartz and Fast Company on AI-unicorn headcount and time-to-\$1B versus traditional unicorns (~200 people / ~2 years vs ~400 people / ~9 years).

[126] Alaska Permanent Fund Corporation (unconditional resident dividend since 1982, ~\$1,700 per person in 2025); Stockton Economic Empowerment Demonstration (SEED), 2019–2021 (\$500/month to 125 residents; full-time employment 28% → 40%); GiveDirectly / Innovations for Poverty Action, long-run cash transfers in Kenya; for Finland see [33]. <https://apfc.org/> [127] Smart Analytics Global (SAG), cited in Global Times, “China’s humanoid robot makers hold 97% of global shipments,” August 10, 2026 (H1 2026 global shipments ~19,100 units, +272% YoY; Chinese vendors >97% of shipments, China >85% of demand; AgiBot overtook Unitree; Unitree shipped 5,500 units in full-year 2025); Al Jazeera, “China’s Unitree soars in market debut,” August 19, 2026 (STAR Market listing at ¥150.8/share; up to +620% intraday, closing +460%, peak market value ~¥445bn); The Motley Fool, “Elon Musk Says Optimus Could Be the Biggest Product Ever,” July 28, 2026 (zero mass-production units off the Fremont line as of mid-2026); Forbes, “We Just Had The First Humanoid Robot Strike Ever,” July 17, 2026, and Ars Technica, July 16, 2026 (~35,000 workers; 25,000+ Boston Dynamics Atlas units planned starting in the US; core demand that no robot enter without a labour-management agreement first); IEEE Spectrum interview with Toyota Research Institute CEO Gill Pratt, April 2, 2026; Rodney Brooks, “House Cleaning My January 1, 2018 Predictions,” July 26, 2026. <https://www.globaltimes.cn/page/202608/1367923.shtml>

[128] U.S. Bureau of Labor Statistics, public API series retrieved August 20, 2026 — LNS14000000, LNS11300000, LNS12300000, LNS11000000, LNS12000000, LNS13000000, LNS15000000, LNU05026645, LNS13025703, CES0000000001, and JOLTS HIR/LDR/QUR/JOR. Identity checks performed: employment plus unemployment equals the published labour force level, and unemployment divided by labour force reproduces the published unemployment rate, for every month cited. January 2026 is the base period, avoiding the annual CPS population-control break between December and January. Singapore: Ministry of Manpower, Unemployment Summary Table, June 2026 preliminary, released July 31, 2026. <https://api.bls.gov/publicAPI/v2/timeseries/data/>

[129] Waymo, “Waymo Safety Impact” (220.6 million rider-only miles through March 2026; Phoenix 80.551M, SF Bay Area 67.078M, Los Angeles 51.816M, Austin 15.789M, Atlanta 5.379M); Waymo blog, June 24, 2026 (more than 4 million miles every week) and May 13, 2026 (11 cities, 1,400+ square miles); Baidu Q2 2026 earnings call (Apollo Go: ~1 million fully driverless rides in the quarter, cumulative >23 million, ~one airbag deployment every 14.4 million km; Robin Li: “temporarily affected by operational adjustments in certain domestic cities due to regulatory considerations”); CnEVPost citing Bloomberg, April 29 and July 23, 2026 (March 31, 2026 multi-vehicle stall

in Wuhan; three agencies including MIIT convened pilot-city officials and China suspended new robotaxi permits; issuance resumed gradually after an industry-wide safety review completed in late June. **Sourced to Bloomberg citing people familiar with the matter, not an official announcement**); NHTSA recall 26E035, June 17, 2026 (3,871 fifth-generation Waymo ADS; defect “the software may allow the vehicle to enter a closed freeway construction zone and continue driving at speed”; interim remedy “Waymo modified the scope of vehicle operations to restrict freeway driving”); OSHA, “Robotics — Standards”; ISO/CD 25785-1. <https://waymo.com/safety/impact/> · <https://static.nhtsa.gov/odi/rc1/2026/RCAK-26E035-6843.pdf>

[130] Richard Wiseman, “The Luck Factor,” *Skeptical Inquirer* (ten-year tracking study, original paper). http://richardwiseman.com/resources/The_Luck_Factor.pdf

[131] CNBC, “The new anti-AI trade sweeping Wall Street: ‘HALO’,” February 27, 2026. <https://www.cnbc.com/2026/02/27/the-new-anti-ai-trade-sweeping-wall-street-halo.html>

[132] The SaaS CFO, “The SaaSocalypse: AI Agents, Vibe Coding, and the Changing Economics of SaaS.” <https://www.thesaascfo.com/the-saaspocalypse-ai-agents-vibe-coding-and-the-changing-economics-of-saas/>

[133] Daily AI World, “Vibe Coding & Disposable Software 2026: The End of SaaS as We Know It.” <https://dailyaiworld.com/post/vibe-coding-disposable-software-2026-the-end-of-saas-as-we-know-it>

[134] Harrison AI X, “Disposable Software: Why Code is No Longer an Asset.” <https://harrisonaix.com/blog/disposable-software-code-asset/>

[135] World Health Organization, “Road traffic injuries” fact sheet (approx. 1.16 million road traffic deaths globally per year; checked August 2026). <https://www.who.int/news-room/fact-sheets/detail/road-traffic-injuries>

[136] NHTSA, “Critical Reasons for Crashes Investigated in the National Motor Vehicle Crash Causation Survey,” DOT HS 812 115, 2015 (the “critical reason” was assigned to the driver in about 94% of crashes; NHTSA cautions this should not be read as “94% caused by driver error”). <https://crashstats.nhtsa.dot.gov/Api/Public/ViewPublication/812115>

[137] Holt-Lunstad, J., Smith, T.B., & Layton, J.B., “Social Relationships and Mortality Risk: A Meta-analytic Review,” *PLoS Medicine* 7(7): e1000316, 2010 (148 studies, 308,849 participants; 50% greater likelihood of survival for those with stronger social relationships, OR = 1.50, 95% CI 1.42-1.59; the paper benchmarks this effect size against smoking and other established mortality risks, which is the origin of the popular “equivalent to 15 cigarettes a day” framing). <https://journals.plos.org/plosmedicine/article?id=10.1371/journal.pmed.1000316>

[142] OECD, *Education at a Glance 2024 / 2025* (tertiary attainment among 25-34 year-olds: 55% of women against 42% of men, a 13-percentage-point gap in women’s favor; Mexico is one of the few member states where young women have not overtaken young men). https://www.oecd.org/en/publications/education-at-a-glance-2024_c00cad36-en/full-report/to-what-level-have-adults-studied_514be7b8.html

[143] Esteve, A. et al., “The End of Hypergamy: Global Trends and Implications,” *Population and Development Review*, 2016 (120 countries, 1960-2011, 420 samples; the rise in women’s educational advantage is matched by a rise in couples where the wife is more educated than the husband). <https://>

[//pmc.ncbi.nlm.nih.gov/articles/PMC5421994/](https://pmc.ncbi.nlm.nih.gov/articles/PMC5421994/)

[144] Erát, D., "Educational assortative mating and the decline of hypergamy in 27 European countries," *Demographic Research* 44(7), 2021 (27 countries, 9 rounds of the European Social Survey, 7 birth cohorts; upward matching declines by cohort while downward matching rises). <https://www.demographic-research.org/articles/volume/44/7>

[145] National Bureau of Statistics of China, historical series on the urban share of population (19.39% in 1980; 67.9% at the end of 2025). https://www.stats.gov.cn/sj/zxfb/202501/t20250117_1958332.html

[146] Third Survey on the Social Status of Women in China (2010) and related national surveys (21% of rural women without land against 11% of rural men; 92.4% of men with contracted land held in their own name against 81.2% of women; 8.1% of women lost land due to a change in marital status against 1.1% of men); Civil Code Article 1126, Rural Land Contracting Law Article 6. https://www.npc.gov.cn/zgrdw/npc/lfzt/rlyw/2018-12/28/content_2069481.htm

[147] Zhang & Li, *China Population and Development Studies*, 2023 (2020 census: 8.82 million never-married rural men aged 35 and over, 5.7% of the group, against 3.62% in cities; among never-married men aged 50 and over, rural areas exceed cities by 2.03 million). <https://link.springer.com/article/10.1007/s42379-023-00143-2>

[148] "Love, life, and 'leftover ladies'," 2021 (CGSS 2008/2010/2011 and CFPS 2010/2012; among urban women over 27, a university education lowers the probability of marriage by 2.88%-3.6%, and a postgraduate degree lowers it a further 8.4%-10.4%). <https://www.sciencedirect.com/science/article/pii/S1043951X21000444>

[149] "Bride Drain," *Journal of Development Economics*, 2019 (the outflow of rural women lowers rural men's probability of marrying and amplifies the effect of the sex-ratio imbalance; no corresponding "groom drain" was found). <https://www.sciencedirect.com/science/article/pii/S0927537119300363>

[150] *Journal of Development Economics*, 2026 (the probability of a single person in China marrying fell to half its 1999 level between 1999 and 2019; changes in the supply structure of the educated population explain roughly half of the decline). <https://www.sciencedirect.com/science/article/abs/pii/S0304387825001804>

[151] Research on the sex ratio at birth and the male surplus (China's sex ratio at birth peaked at 121.18 in 2004; at one point there were roughly 32 million more males than females under the age of 20; model estimates of the marriage squeeze). <https://pmc.ncbi.nlm.nih.gov/articles/PMC3168620/>

[152] Research on China's 2001 divorce law reform and sex selection (after women's property rights were strengthened, families whose first child was a girl were 8.1 percentage points less likely to go on to have a boy). <https://www.sciencedirect.com/science/article/pii/S0304387815001327>